

Wrangling and Pipelines for Data Handling

8th - 9th April

Hello!



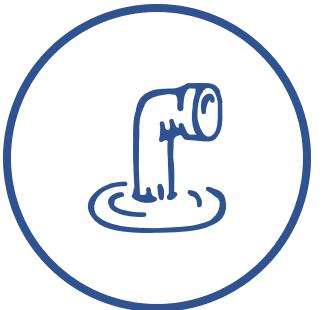
Richard Strange

Environmental Research DTP

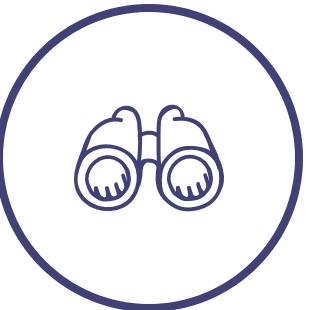
Volcanic Seismology and AI

Data Scientist and Data Architect

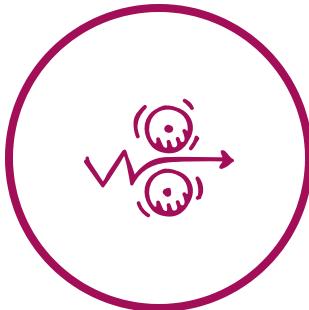
Day 1 Themes



Pipelines



Profiling



Cleaning

Day 1 outline

- 1 What do data pipelines look like?
 - 2 Where do pipelines go wrong?
 - 3 Breakout session: data flow challenges
-
- 4 Building blocks
 - 5 Data Profiling and Cleaning
 - 6 Data wrangling lab

— Lunch! —

Course Objectives

Understand the basics of data pipeline design

Understand the options available in building a pipeline

Understand the likely errors in data, how to spot them and fix them

Understand the advantages of code that pre-emptively reports failures

Understand the purpose and function of new technologies for handling and finding data

Able to begin constructing pipelines

Able to write functional ELT code

Able to interact with online data sources

Course Objectives

Understand the basics of data pipeline design

Understand the options available in building a pipeline

Understand the likely errors in data, how to spot them and fix them

Understand the advantages of code that pre-emptively reports failures

Understand the purpose and function of new technologies for handling and finding data

Able to begin constructing pipelines

Able to write functional ELT code

Able to interact with online data sources

You can make an informed
decision with your data

House rules...

House Rules - I expect you to...



Stop me whenever you want

No such thing as a silly question, and there is no such thing as a badly timed question



Stop underestimating your knowledge

It's impossible to avoid working with data. Just because you might not recognise the terminology doesn't mean you don't already have experience



Stop worrying

This course is a whistle-stop tour of the topic, not a assessed lecture series.

Section 1

Data Pipelines

Modularity



- There are a thousand ways to achieve the same thing
- No one way is correct...

Modularity



- There are a thousand ways to achieve the same thing
- No one way is correct...
- ... but some ways are more correct than others

What is a data pipeline?

"a set of data processing elements connected in series, where the output of one element is the input of the next one [...] some amount of buffer storage is often inserted between elements"

-WIKIPEDIA

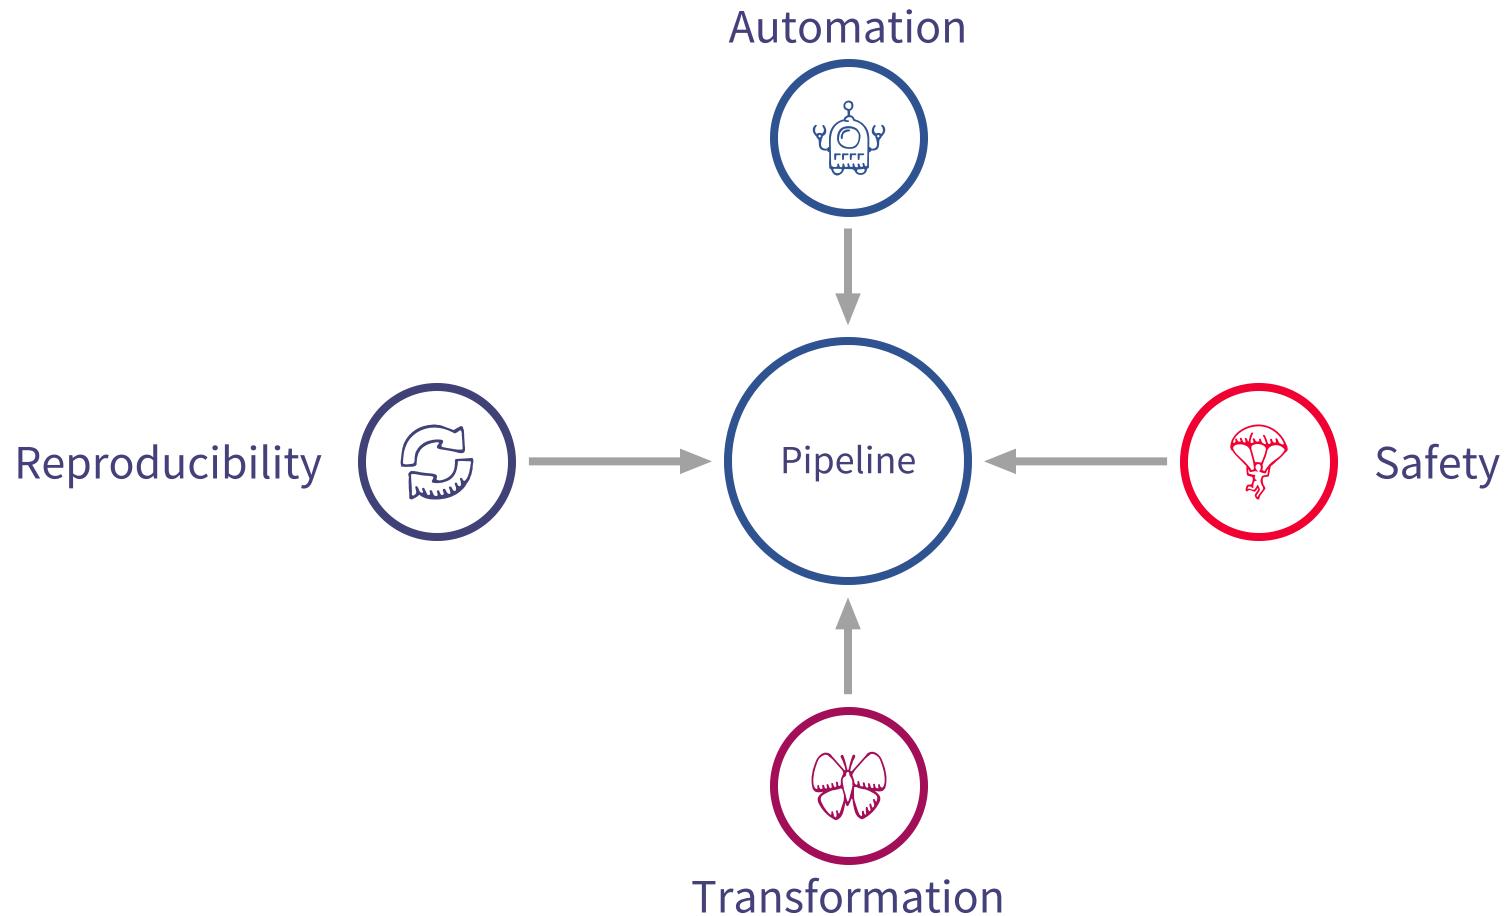
"eliminates many manual steps from the process and enables a smooth, automated flow of data from one station to the next. It starts by defining what, where, and how data is collected. It automates the processes involved in extracting, transforming, combining, validating, and loading data for further analysis and visualization. It provides end-to-end velocity by eliminating errors and combatting bottlenecks or latency."

-ALOOMA

"A data pipeline is a set of actions that extract data (or directly analytics and visualization) from various sources. It is an automated process: take these columns from this database, merge them with these columns from this API, subset rows according to a value, substitute NAs with the median and load them in this other database"

-MEDIUM

What's the point?



Pipeline Reliability

- Data should not arrive duplicated
- Data should not arrive incomplete
- Updated data can be distinguished
- Nothing fails quietly
- Staging catches data, even in partial failure
- Different sources can merge safely

Work
your way
upwards

THE DATA SCIENCE **HIERARCHY OF NEEDS**

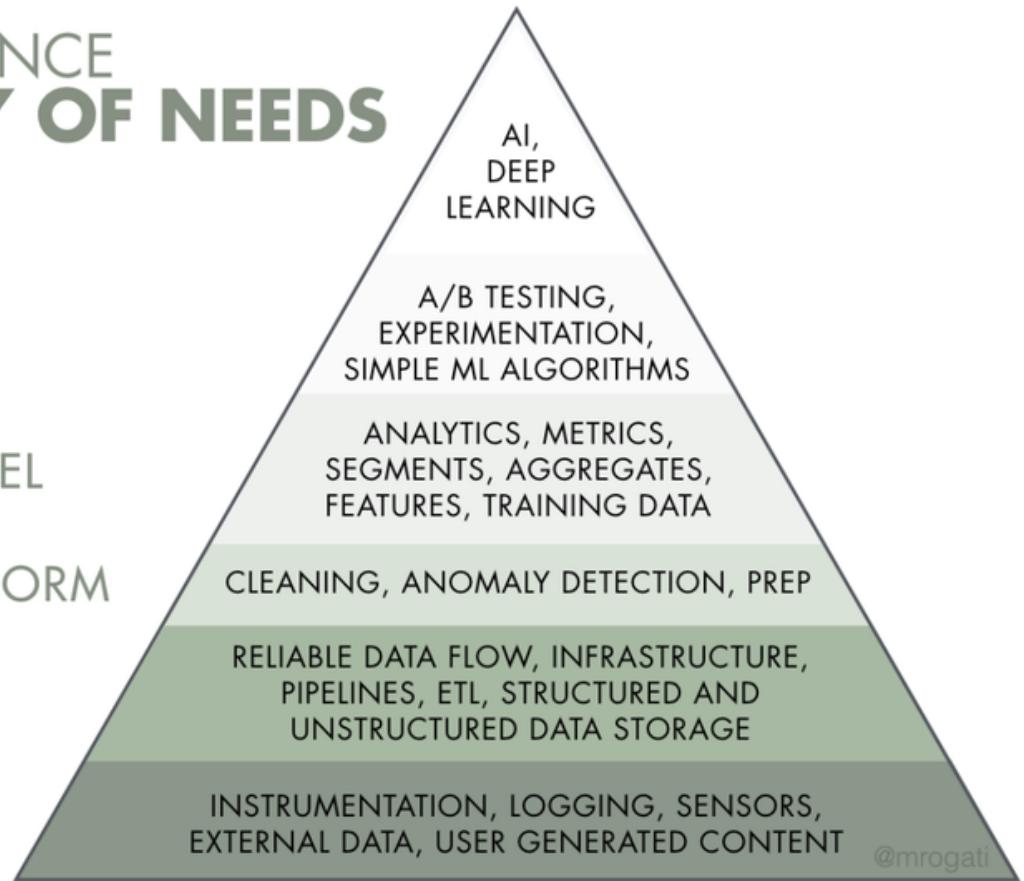
LEARN/OPTIMIZE

AGGREGATE/LABEL

EXPLORE/TRANSFORM

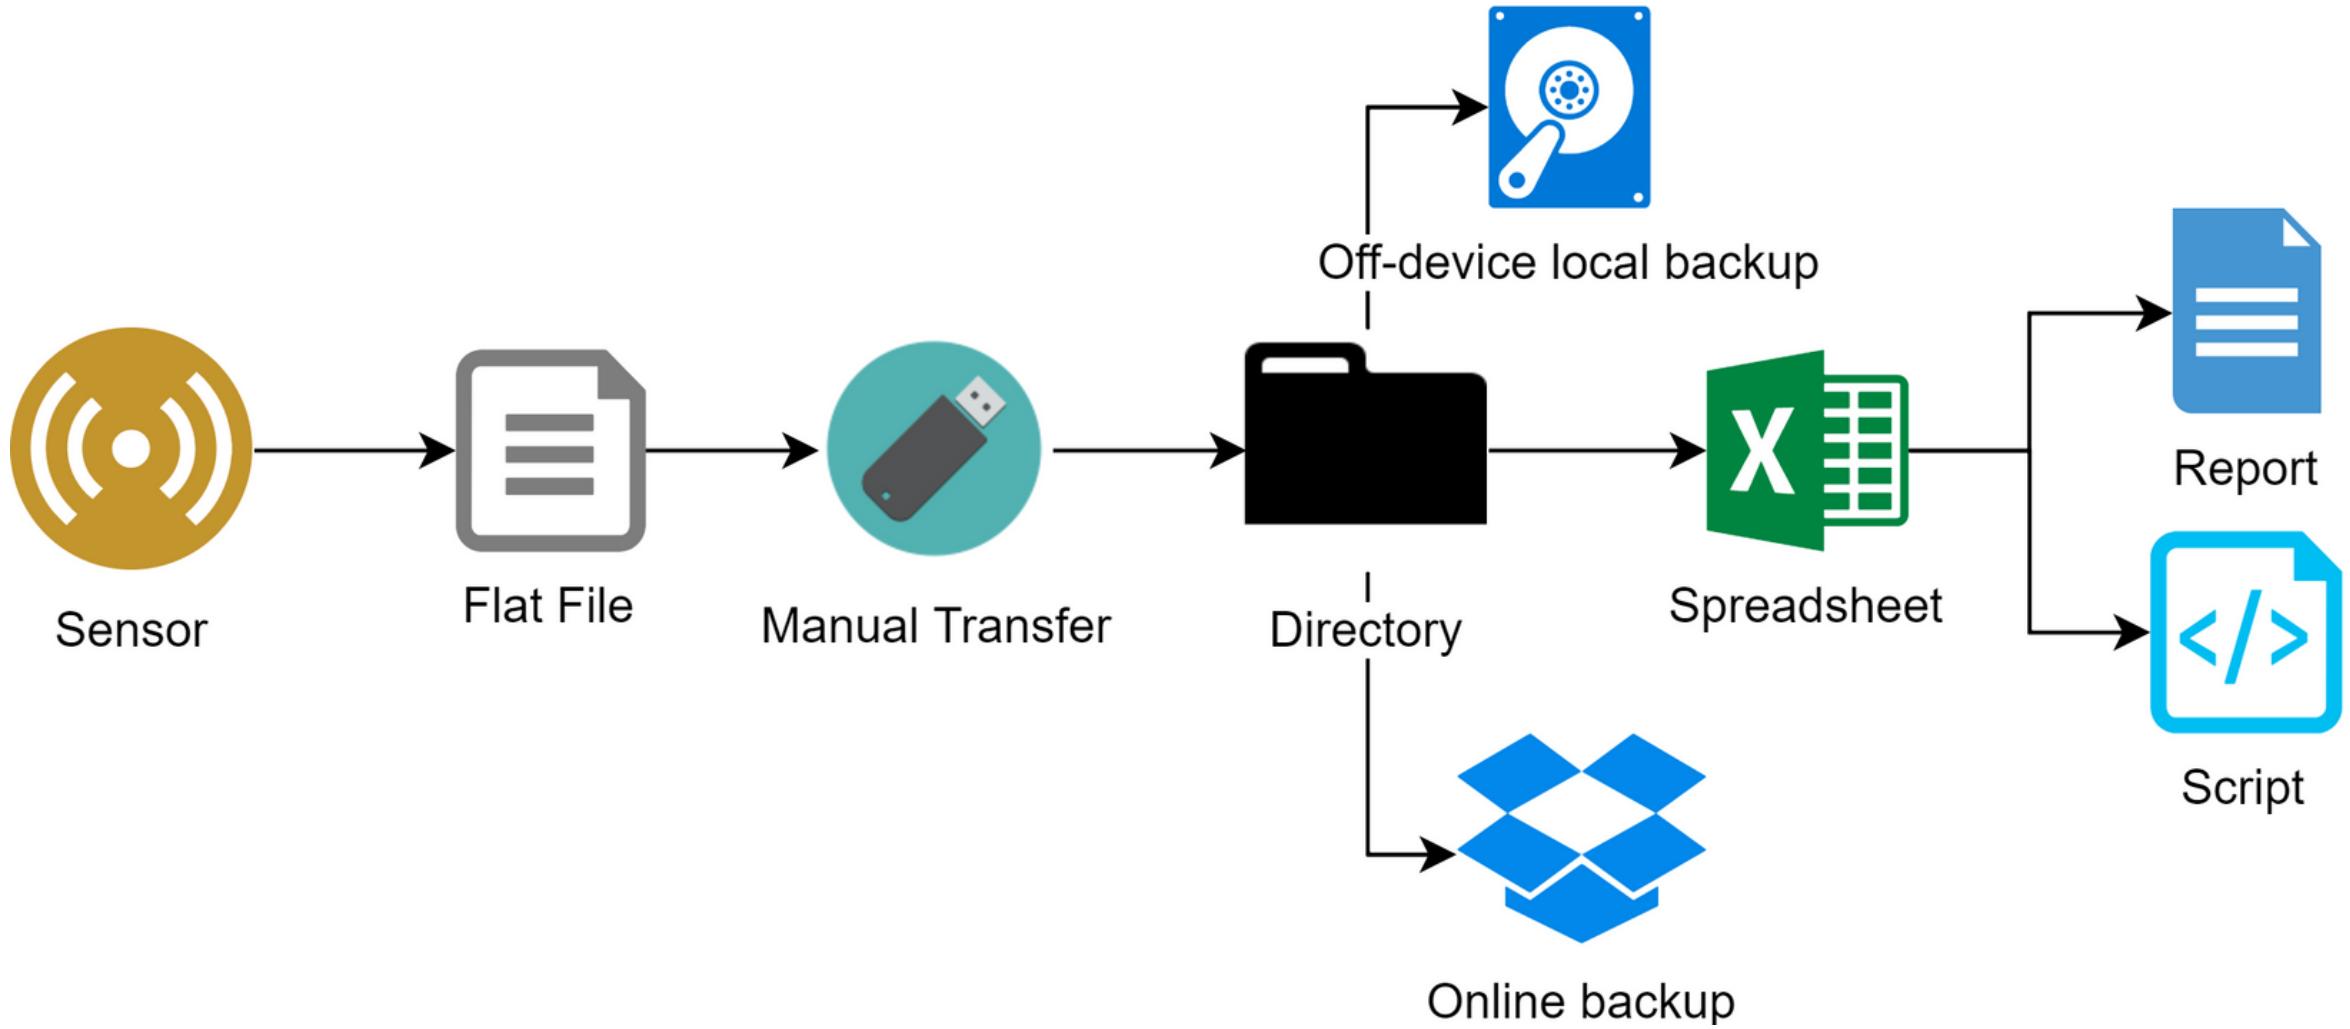
MOVE/STORE

COLLECT

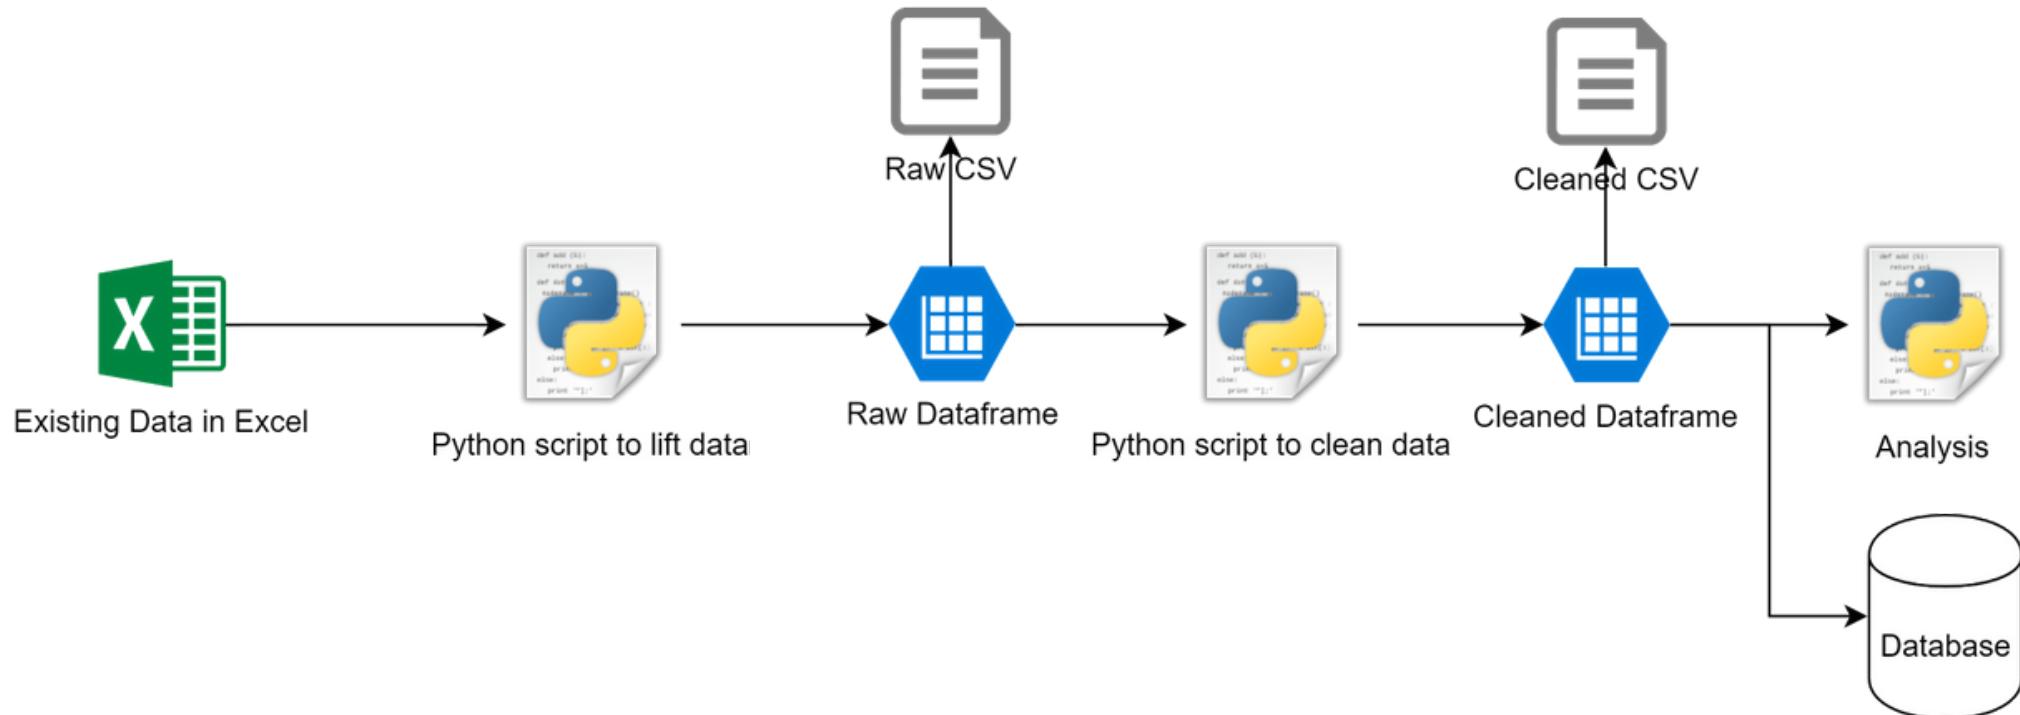


@mrogati

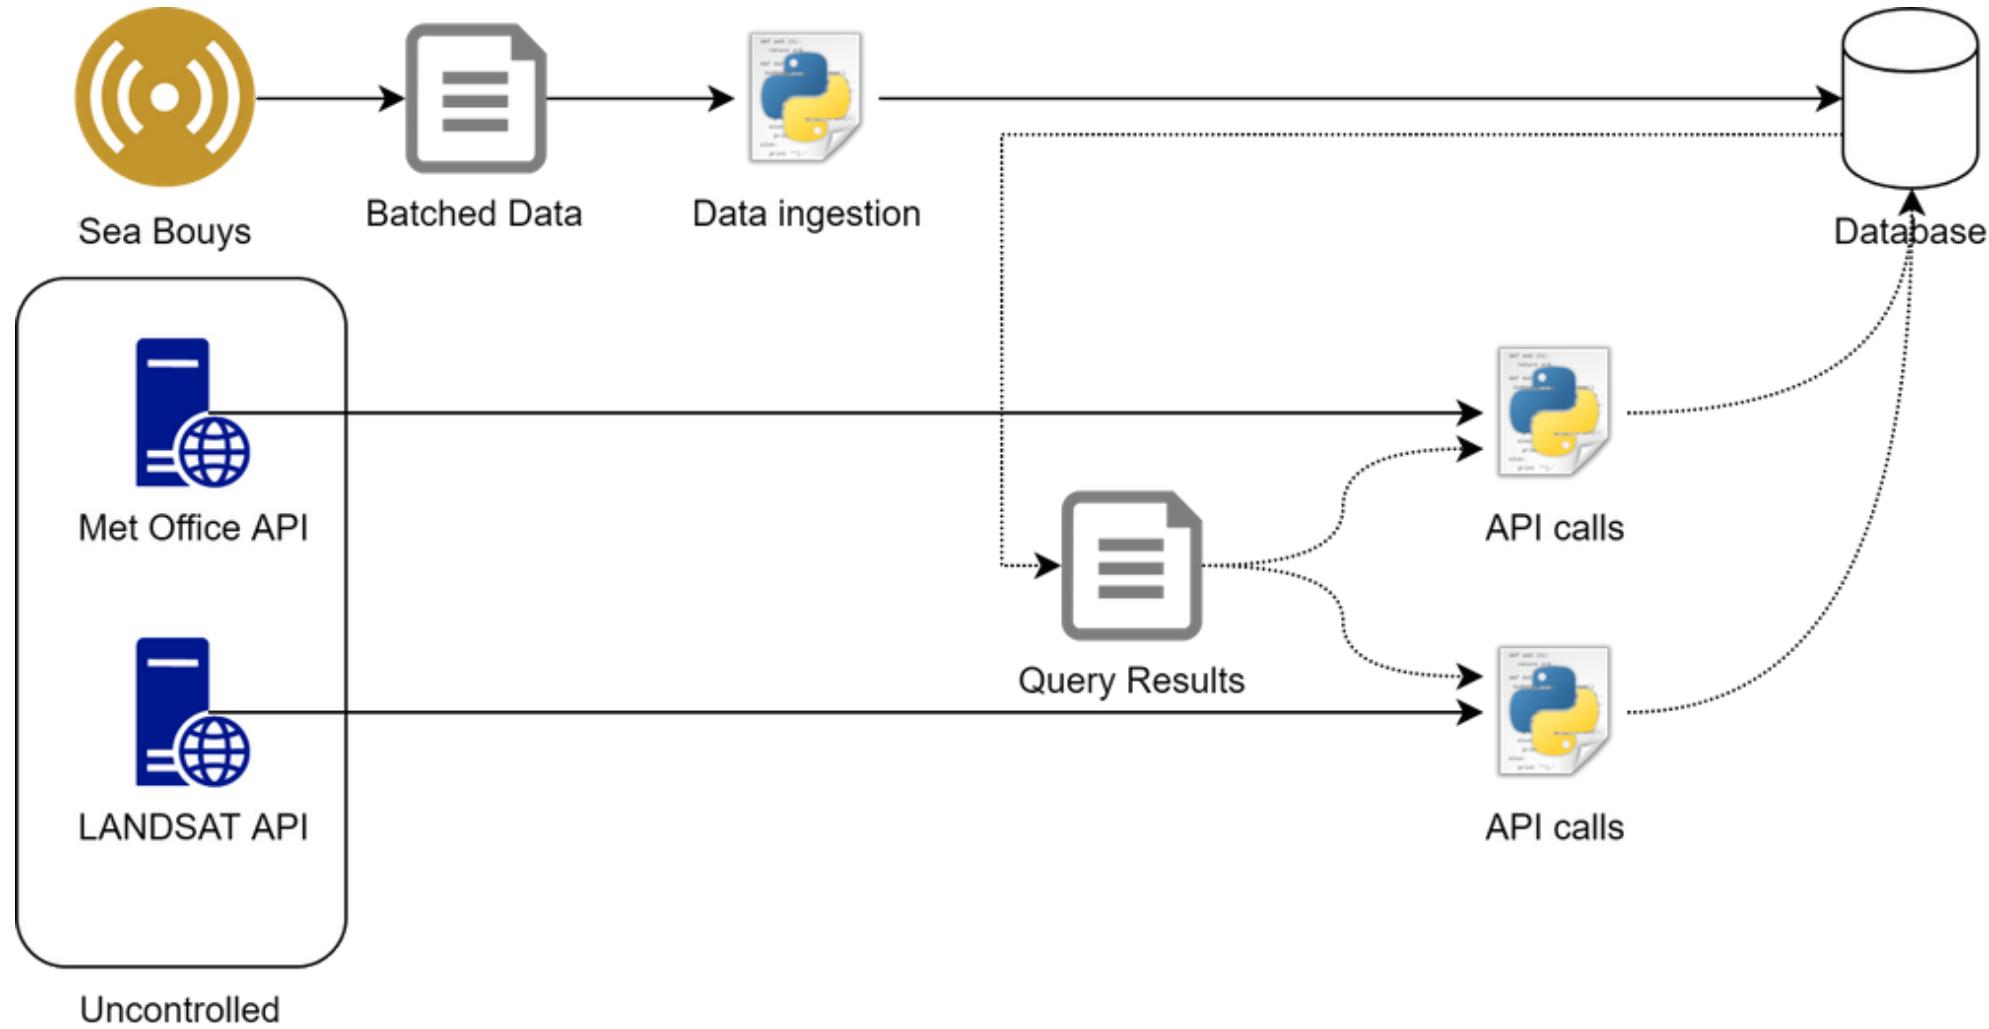
Basic Data flow



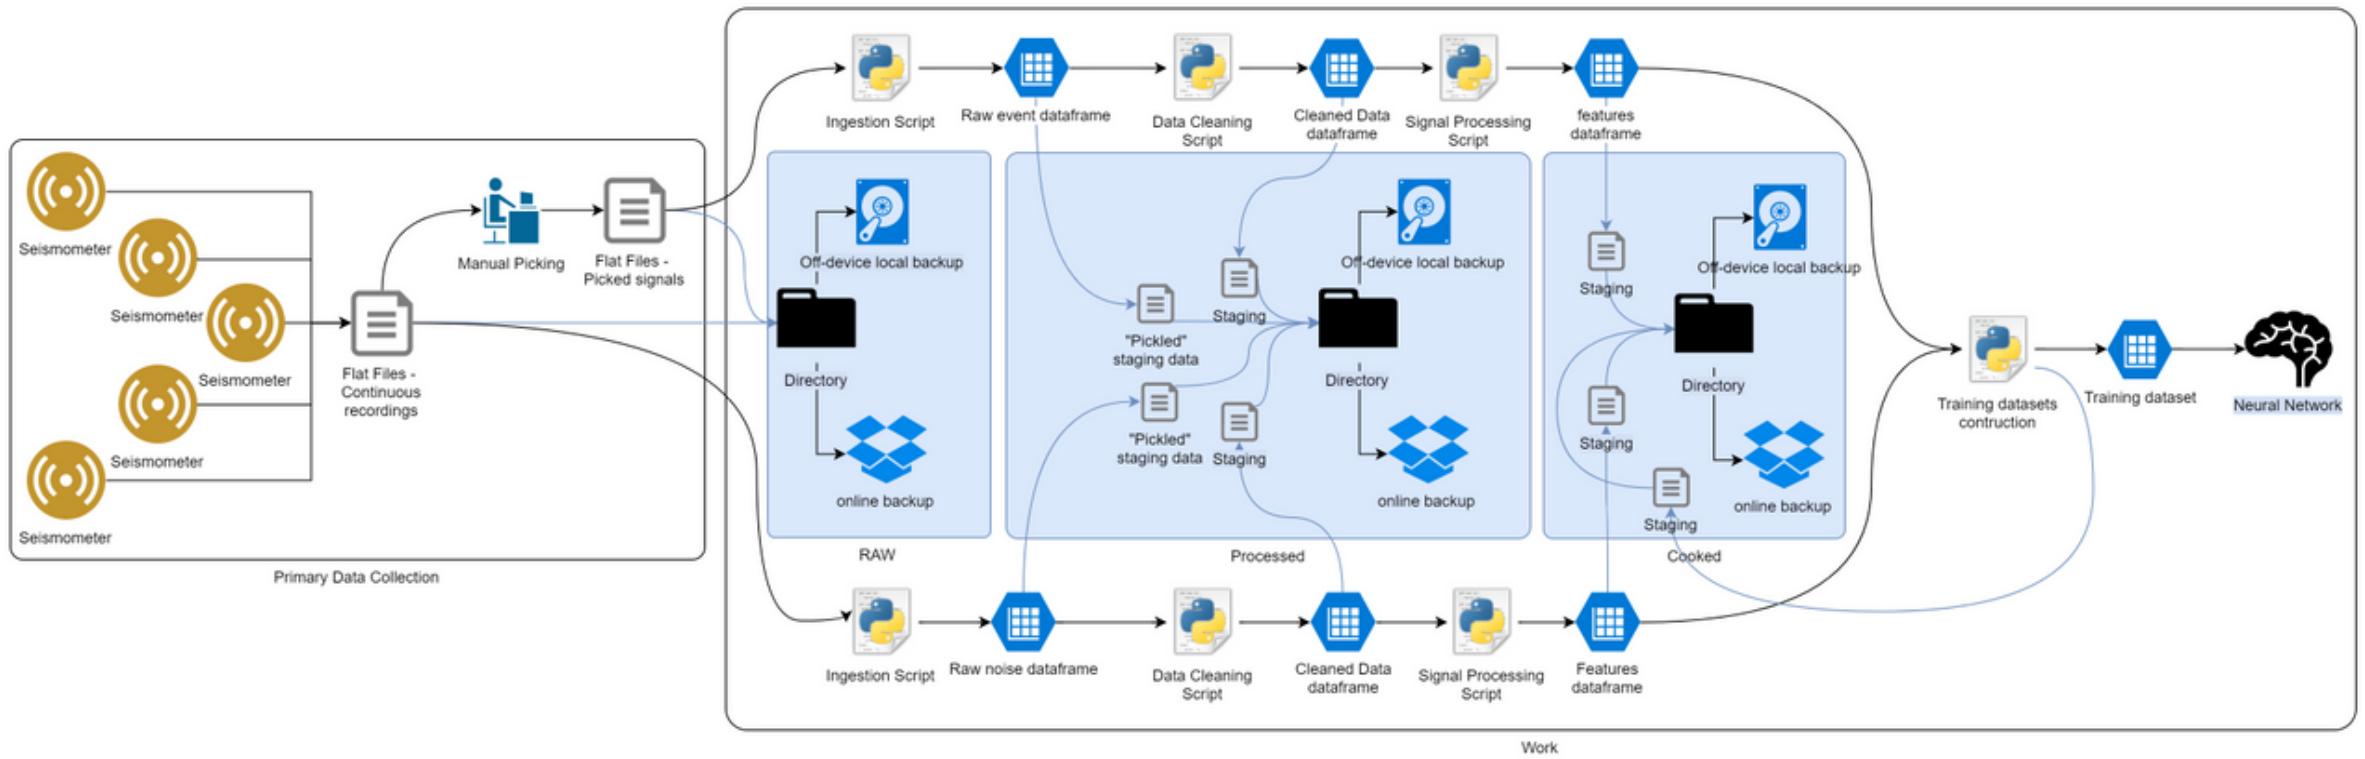
Excel Workflow



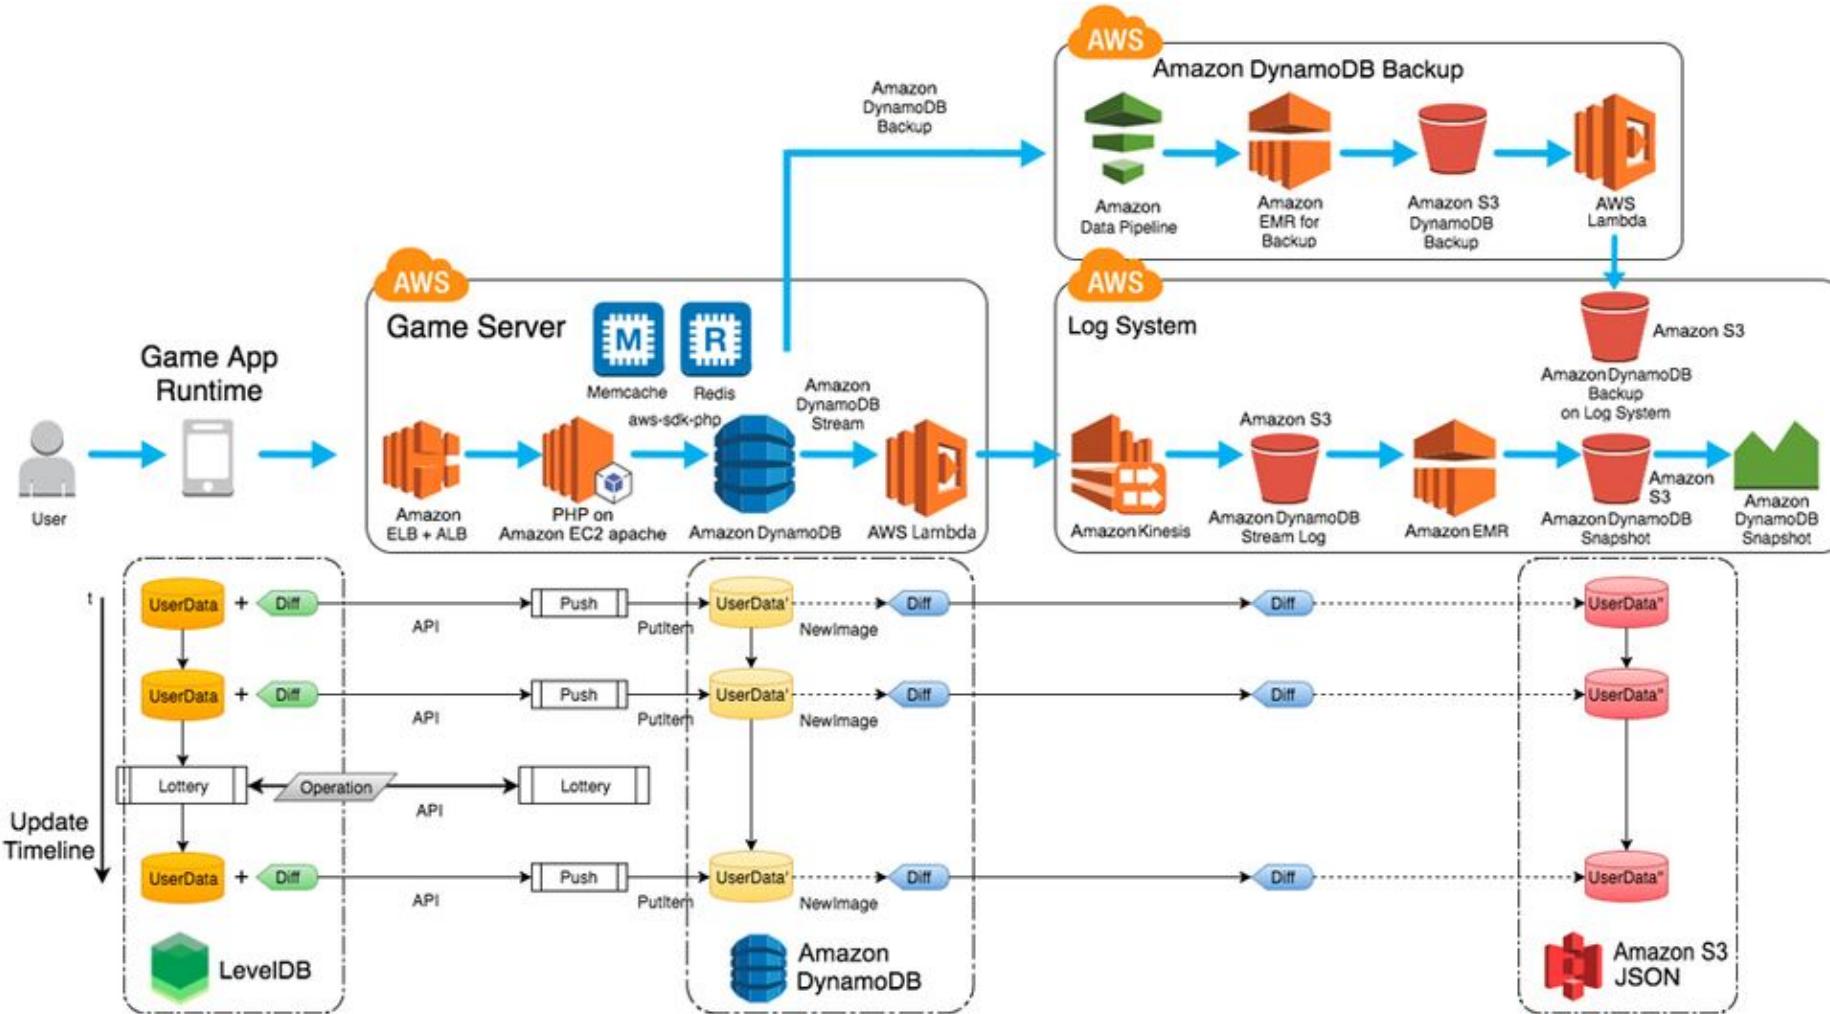
API Workflow



Working seismology pipeline

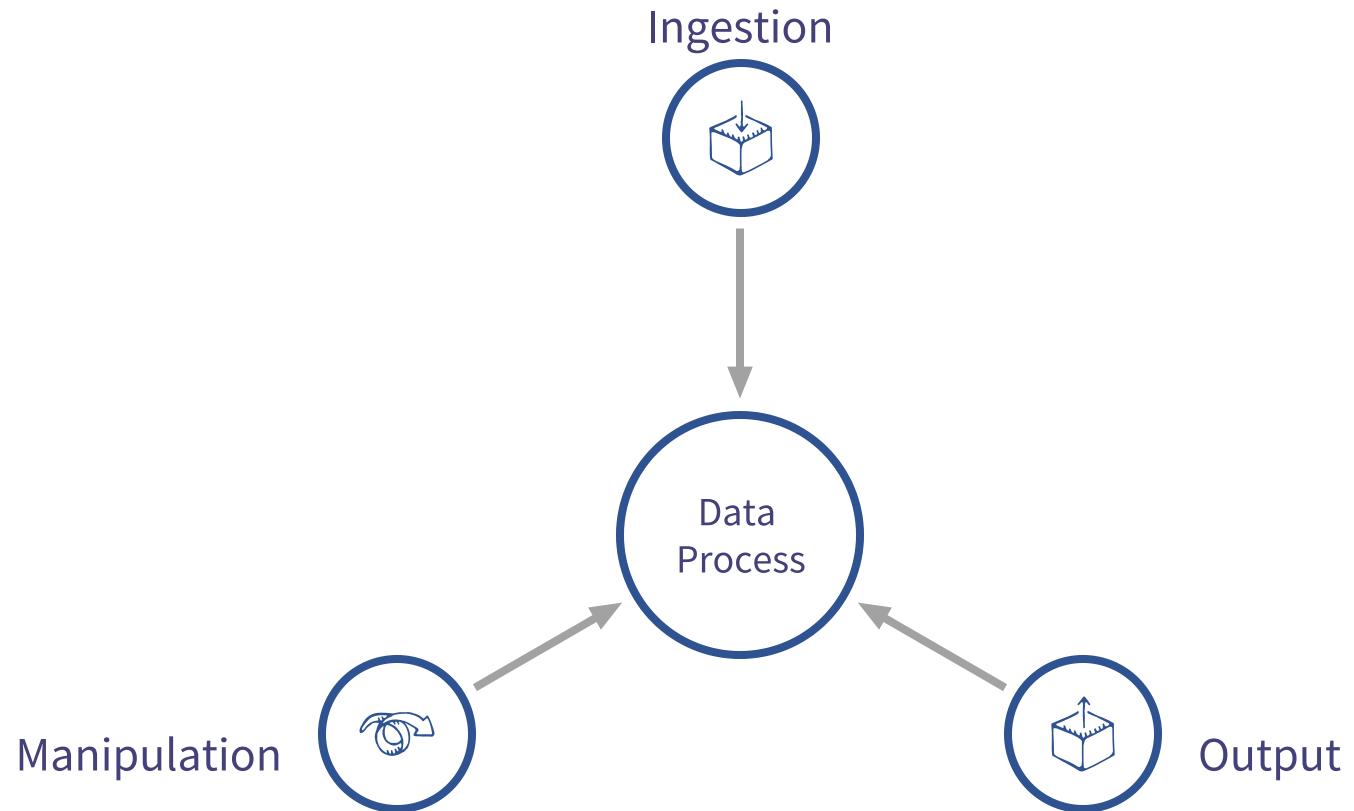


Simple commercial pipeline

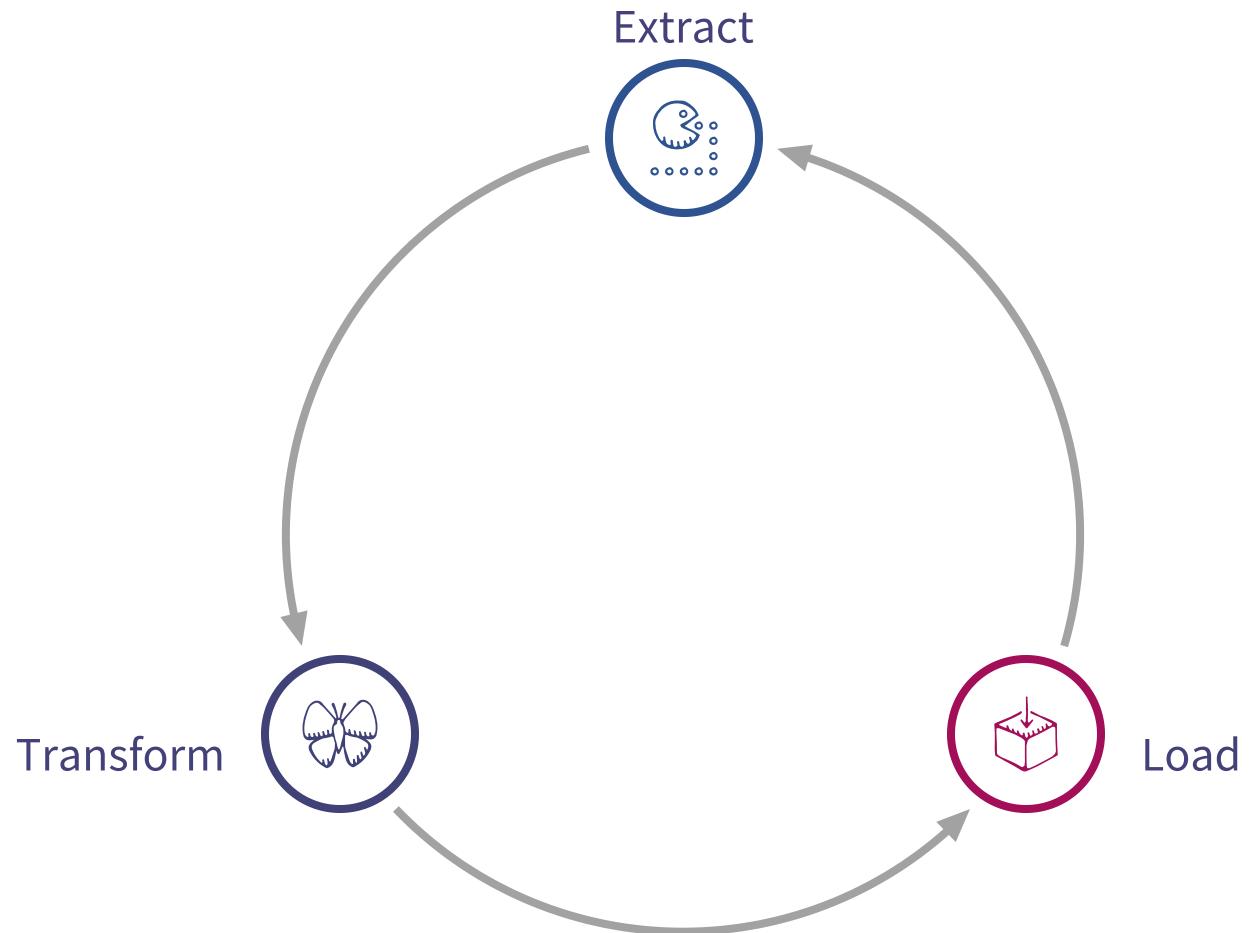


What do data processes do?

Purposes of data handling



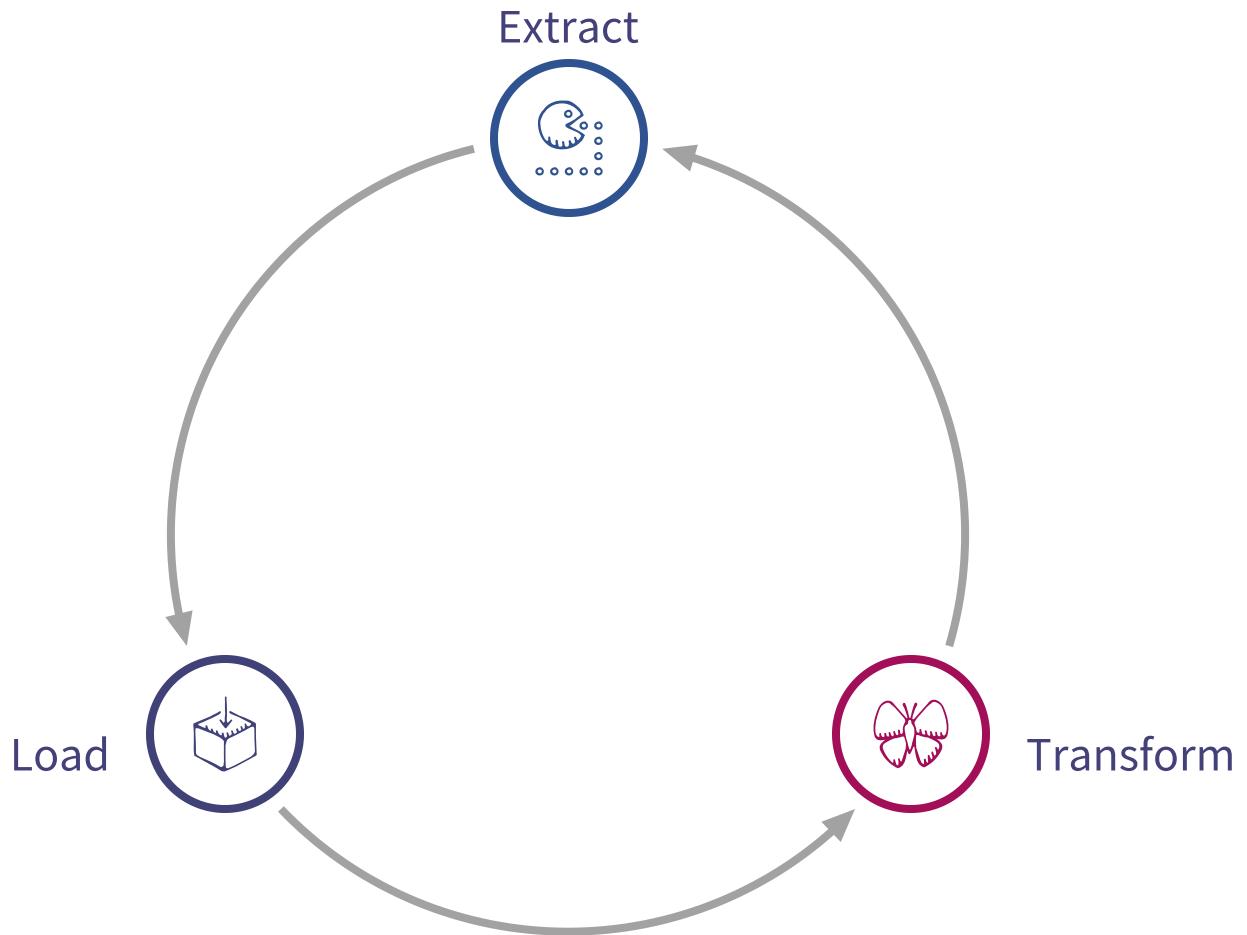
ETL



Extract - pull data from its source

Transform - change the data to a suitable state

Load - put the fixed data into an accessible location



ETL - Raw data exists in the source

ELT - Raw data exists in your system

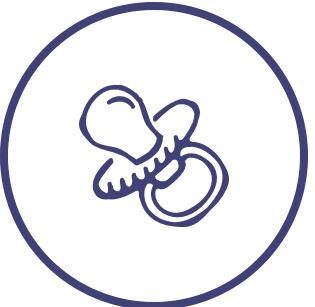
Tenets



Keep it raw

The rawest form of your data should always be kept stored, even if it is a duplicate of data elsewhere.

If anything goes wrong the rest of your data can be worked from first principles from the raw store.

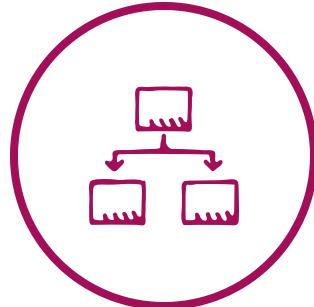


Keep it simple

Every step in a flow should aim to only do one major task.

Every step in a script should aim to do one minor task.

Every staging file should aim to record a change to the underlying data.



Keep it copied

Intermediate data from script or spreadsheet stages can be dumped into flat files.

Being able to evaluate every output makes diagnosing issues and testing steps simple, keeping you sane.

Break

What do you call a group of
stormtroopers playing monopoly?

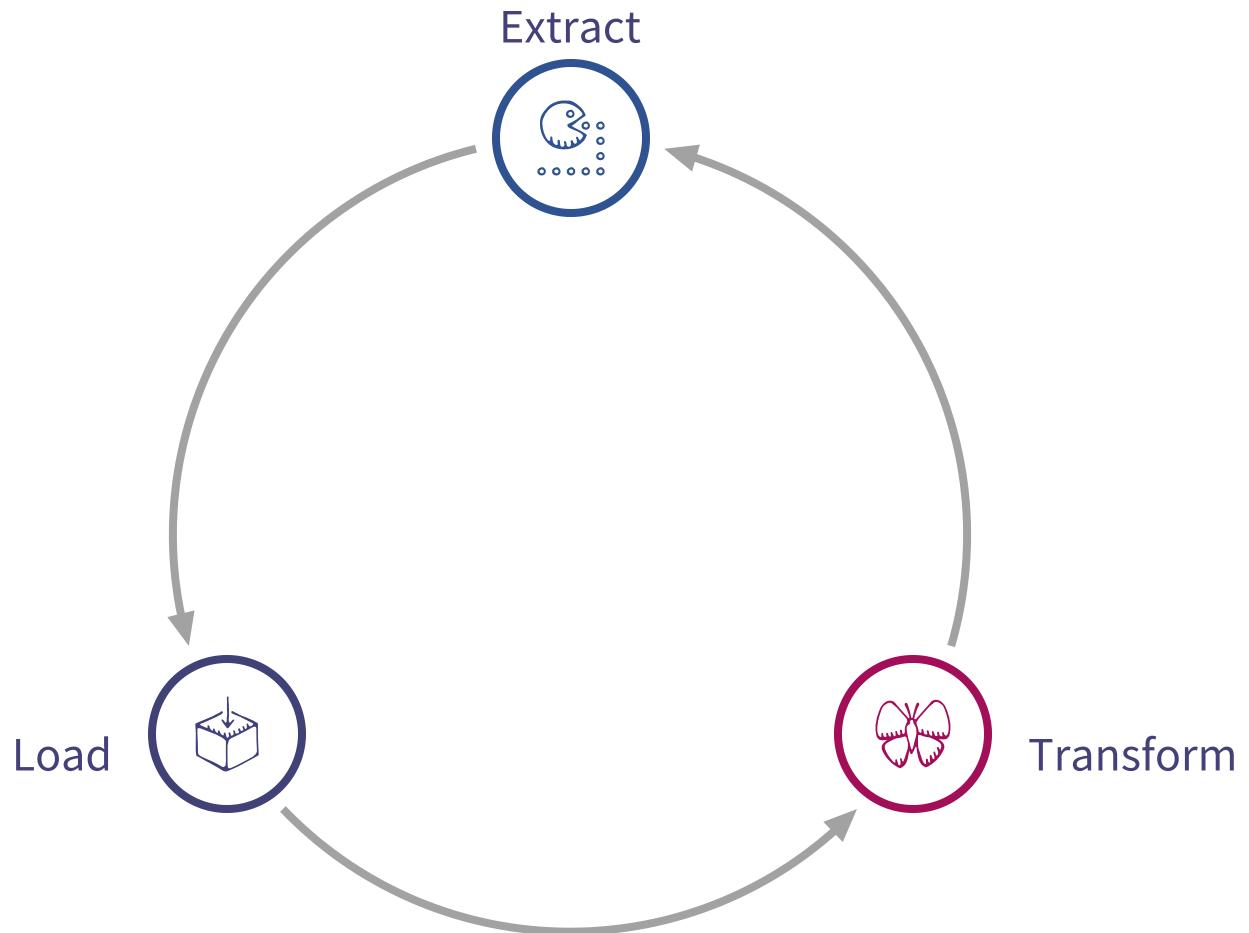


What do you call a group of
stormtroopers playing monopoly?

Game of Clones

Section 2

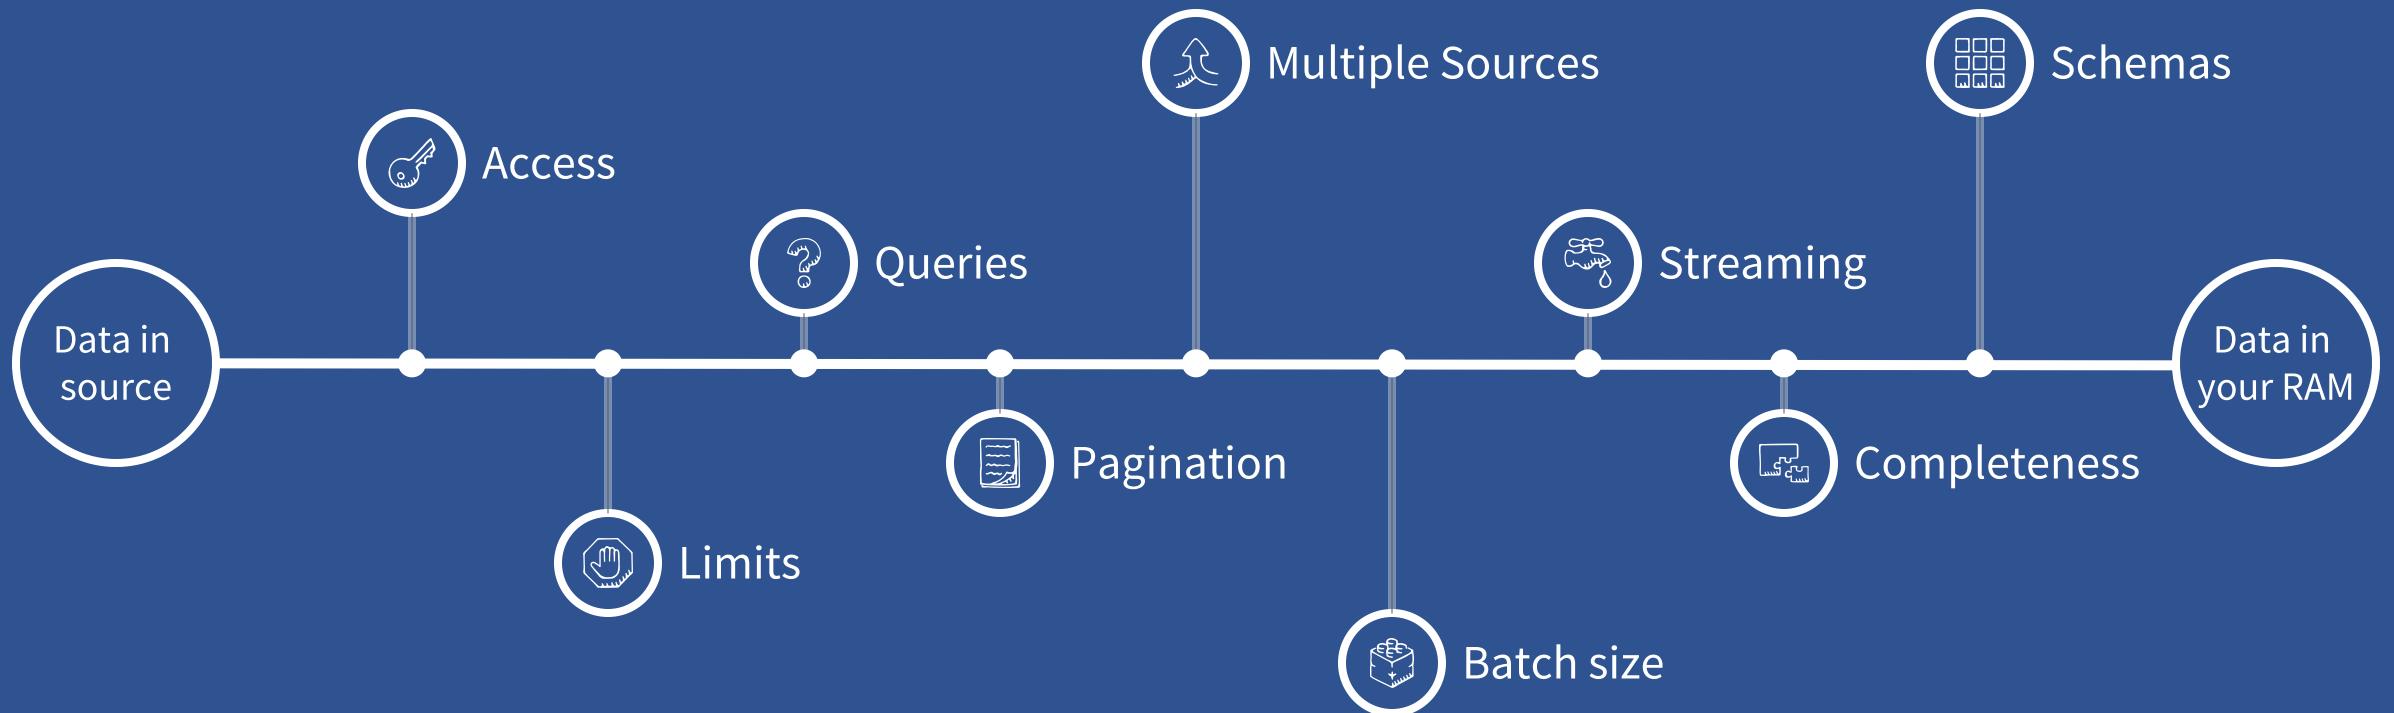
Pitfalls in data management



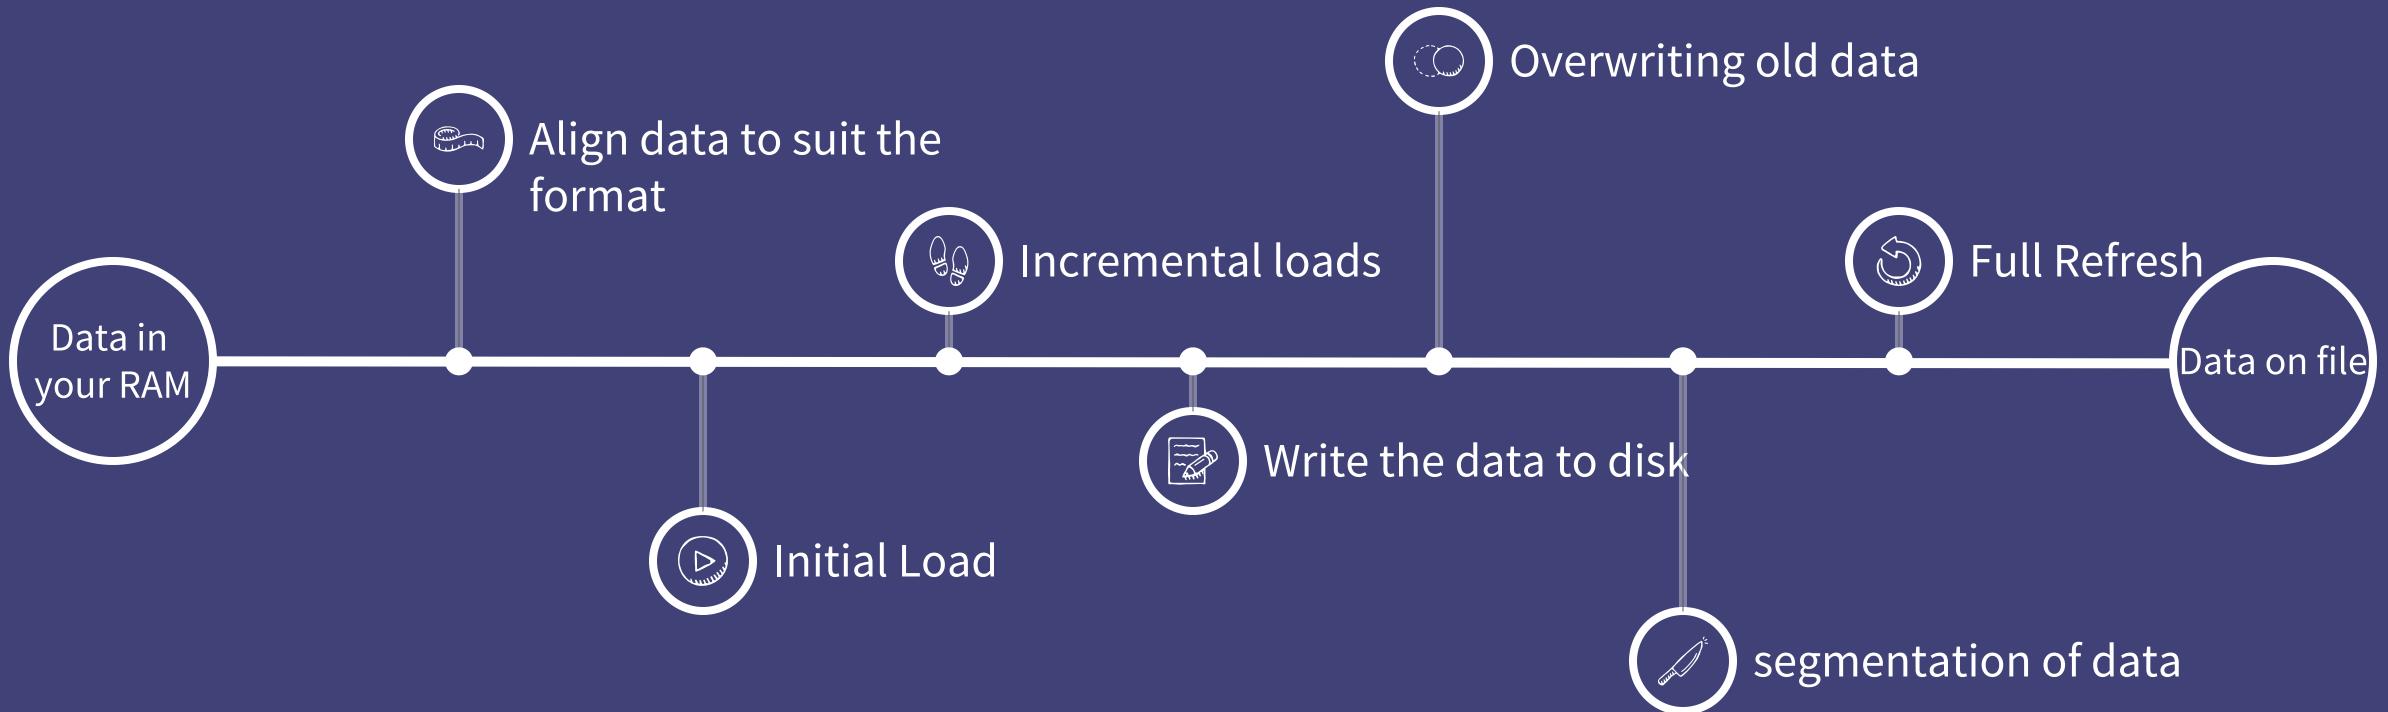
ETL - Raw data exists in the source

ELT - Raw data exists in your system

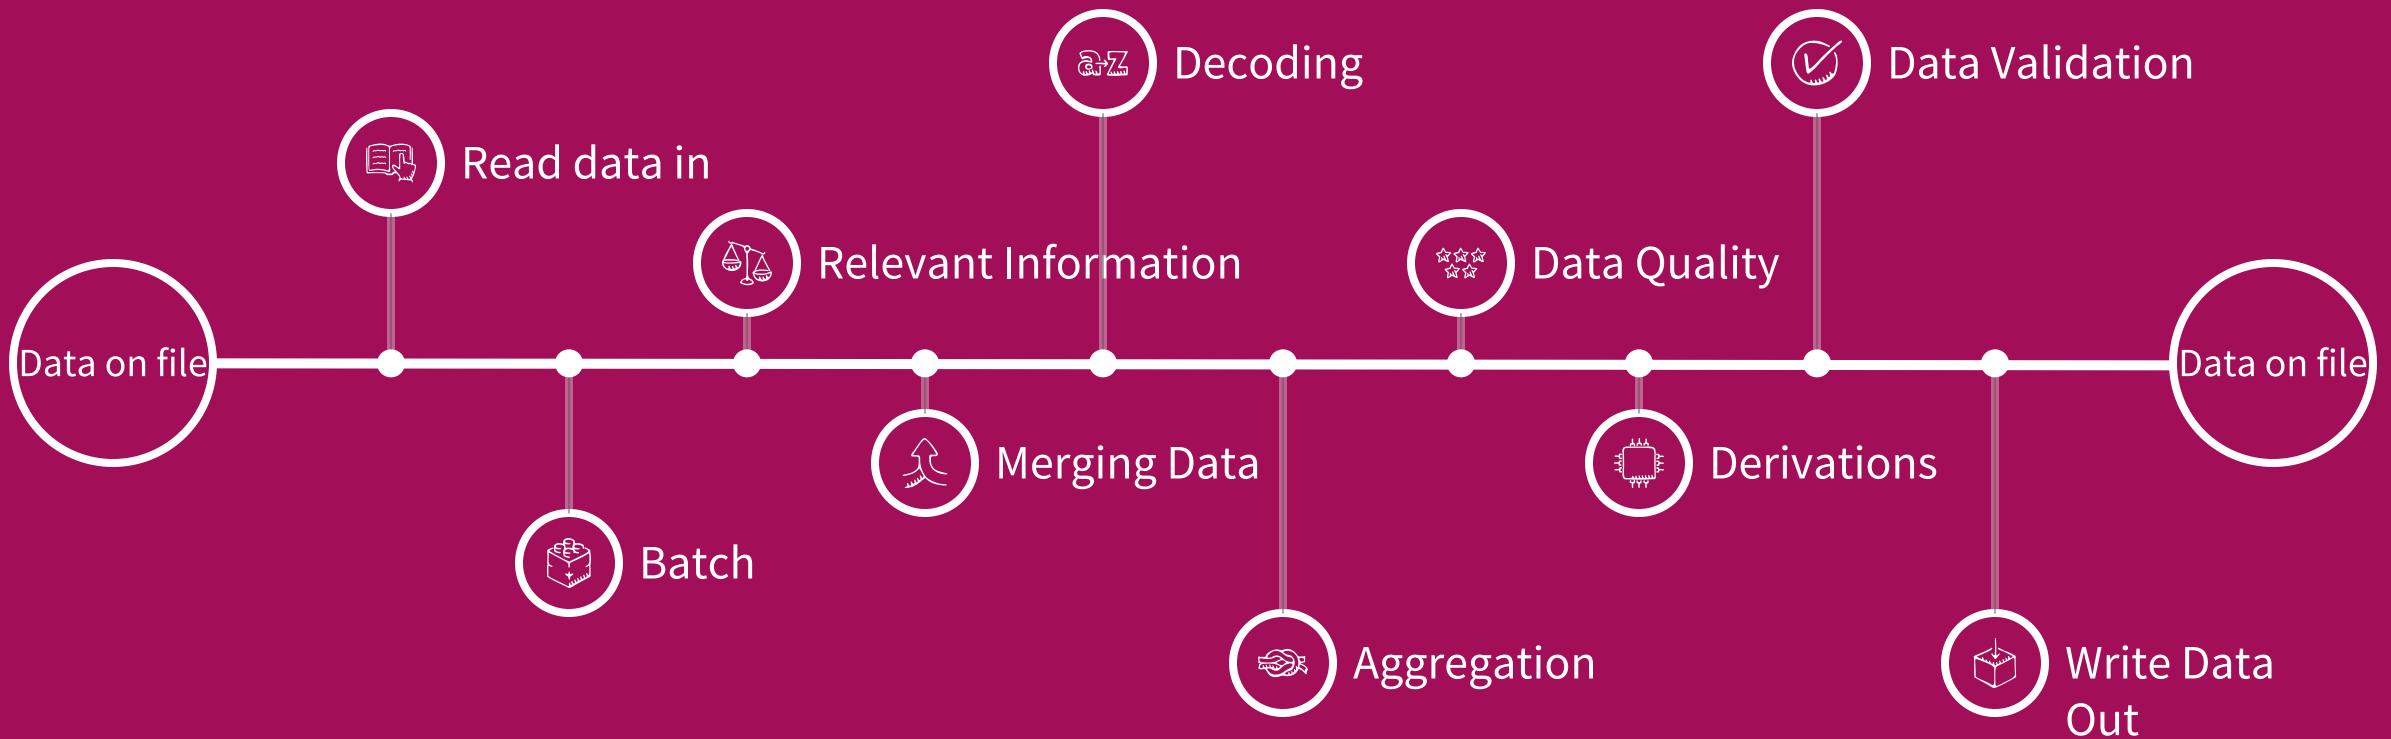
Extract



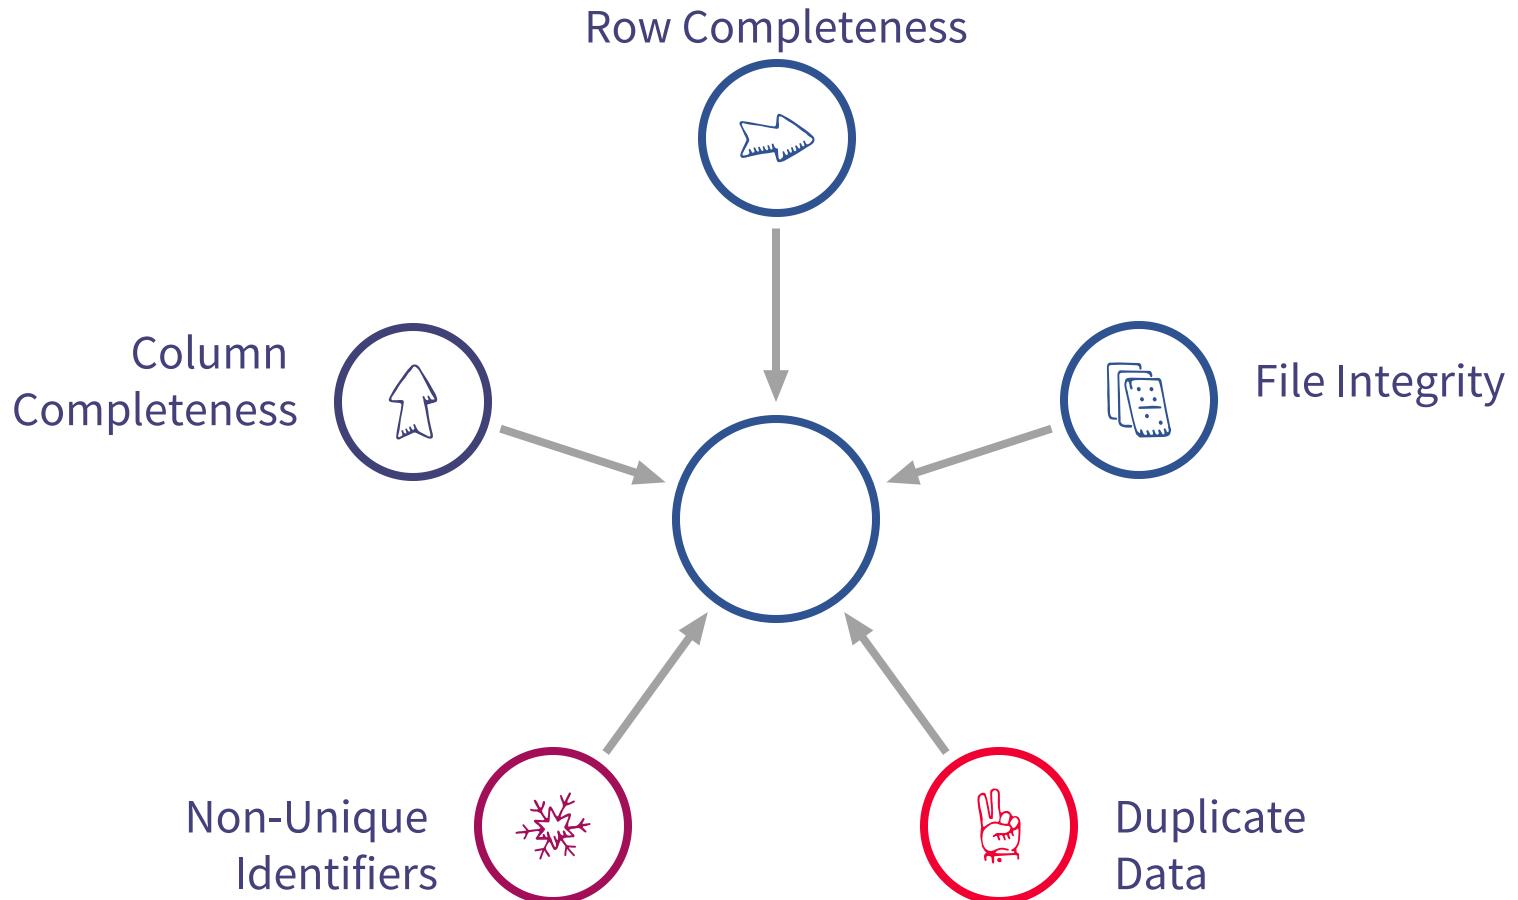
Load



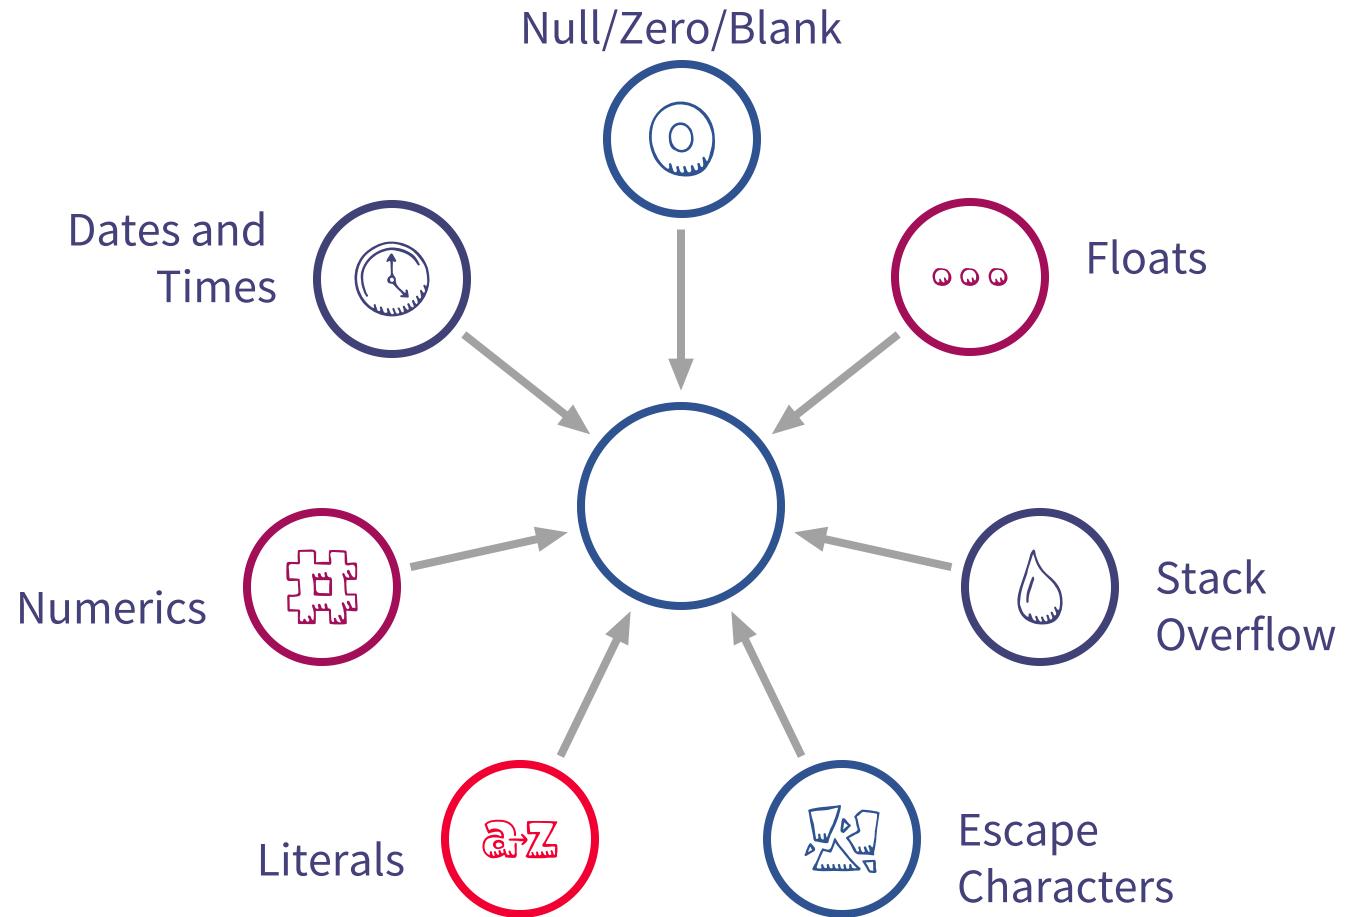
Transform



Out-of-cell issues

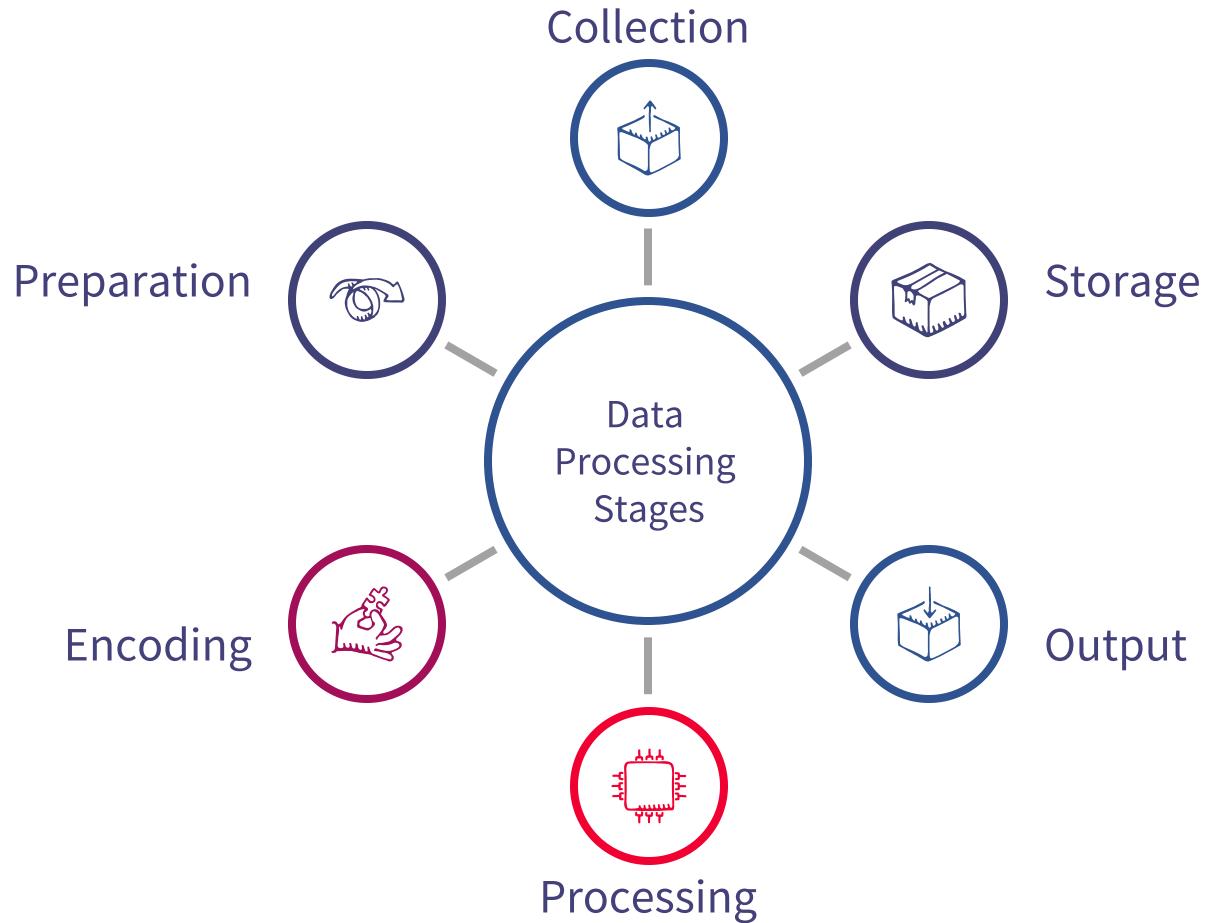


In-cell issues



Your turn!

Areas to consider...



Lunch

What did the Lego Alien say?



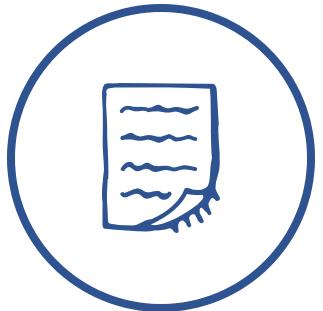
What did the Lego Alien say?

I come in pieces

Section 3

ELT Building Blocks

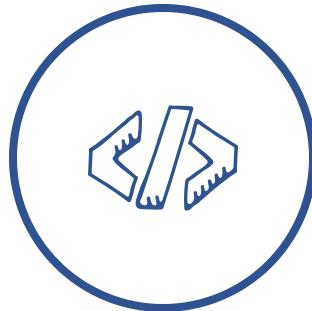
Technique Areas



Files



Shell

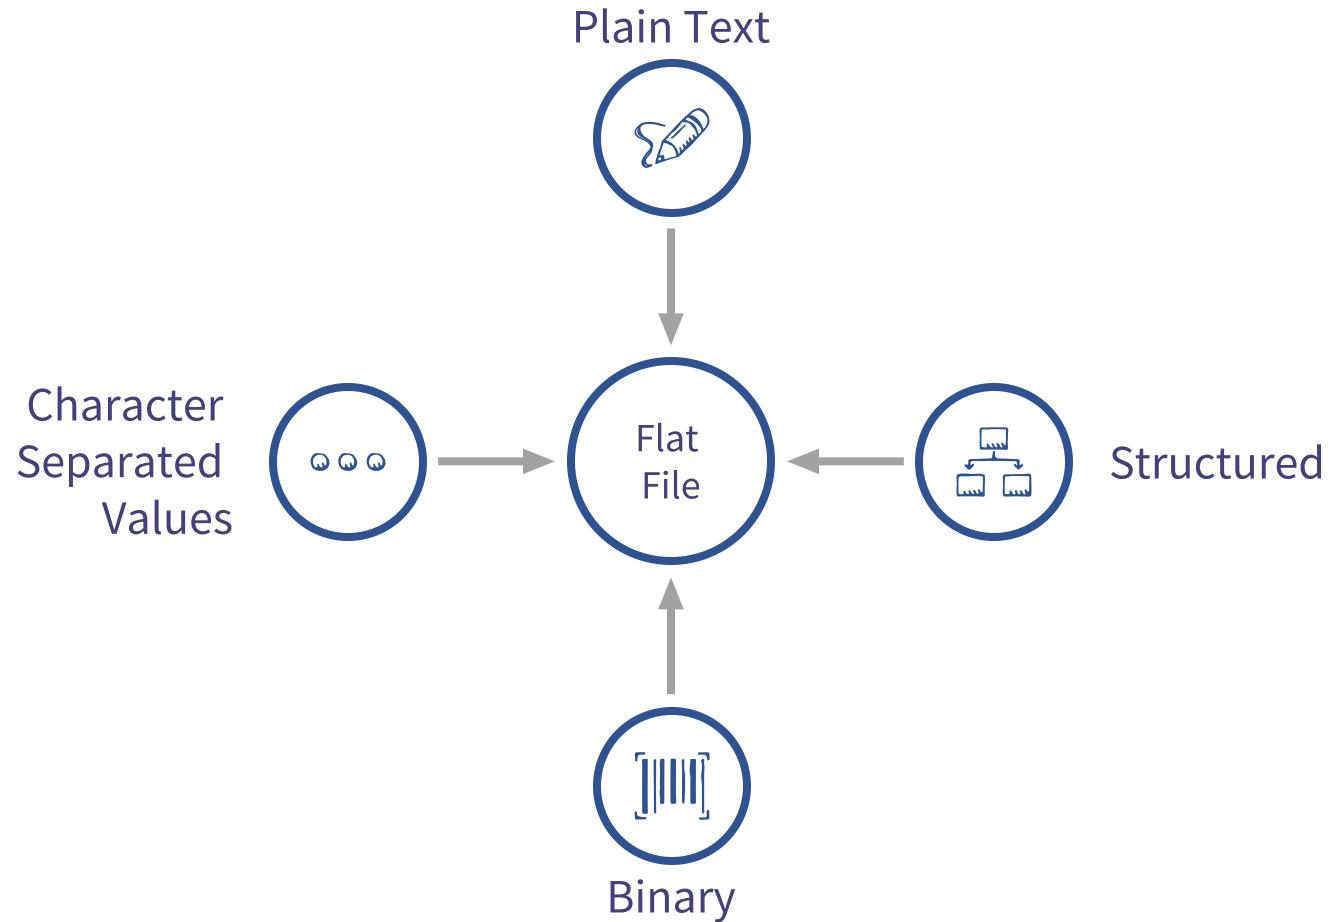


Language



Database

Flat Files



Plain Text

```
VARIABLE : Mean TS from clear sky composite (kelvin)
FILENAME : ISCCPMonthly_avg.nc
FILEPATH : /usr/local/fer_dsets/data/
SUBSET   : 93 points (TIME)
LONGITUDE: 123.8W(-123.8)
LATITUDE : 48.8S
                  123.8W
                  23
16-JAN-1994 00 / 1: 278.9
16-FEB-1994 00 / 2: 280.0
16-MAR-1994 00 / 3: 278.9
16-APR-1994 00 / 4: 278.9
16-MAY-1994 00 / 5: 277.8
16-JUN-1994 00 / 6: 276.1
```

Simplest type of file.
Often difficult to ingest,
and may rely on header
codes to define a
schema.

Can be fixed-width.

Character Separated Values , | ; \n ' " / \ .

```
Year,Make,Model,Description,Price
1997,Ford,E350,"ac, abs, moon",3000.00
1999,Chevy,"Venture ""Extended Edition"""","",4900.00
1999,Chevy,"Venture ""Extended Edition, Very Large""",,5000.00
1996,Jeep,Grand Cherokee,"MUST SELL!
air, moon roof, loaded",4799.00
```

marks the end of cells, rows,
decimals and text with a
character

Easier to ingest as a table of
data

Vulnerable to escape

Can also include Excel
Spreadsheets

Binary Files

Not human-readable
Program-specific dump
of information

Often intermediary data formats

Used to encode or protect information

Structured

```
{  
    "firstName": "John",  
    "lastName": "Smith",  
    "isAlive": true,  
    "age": 27,  
    "address": {  
        "streetAddress": "21 2nd Street",  
        "city": "New York",  
        "state": "NY",  
        "postalCode": "10021-3100"  
    },  
    "phoneNumbers": [  
        {  
            "type": "home",  
            "number": "212 555-1234"  
        },  
        {  
            "type": "office",  
            "number": "646 555-4567"  
        },  
        {  
            "type": "mobile",  
            "number": "123 456-7890"  
        }  
    "children": [],  
    "spouse": null  
}
```

Hold structured hierarchical data and metadata

Can hold class data well

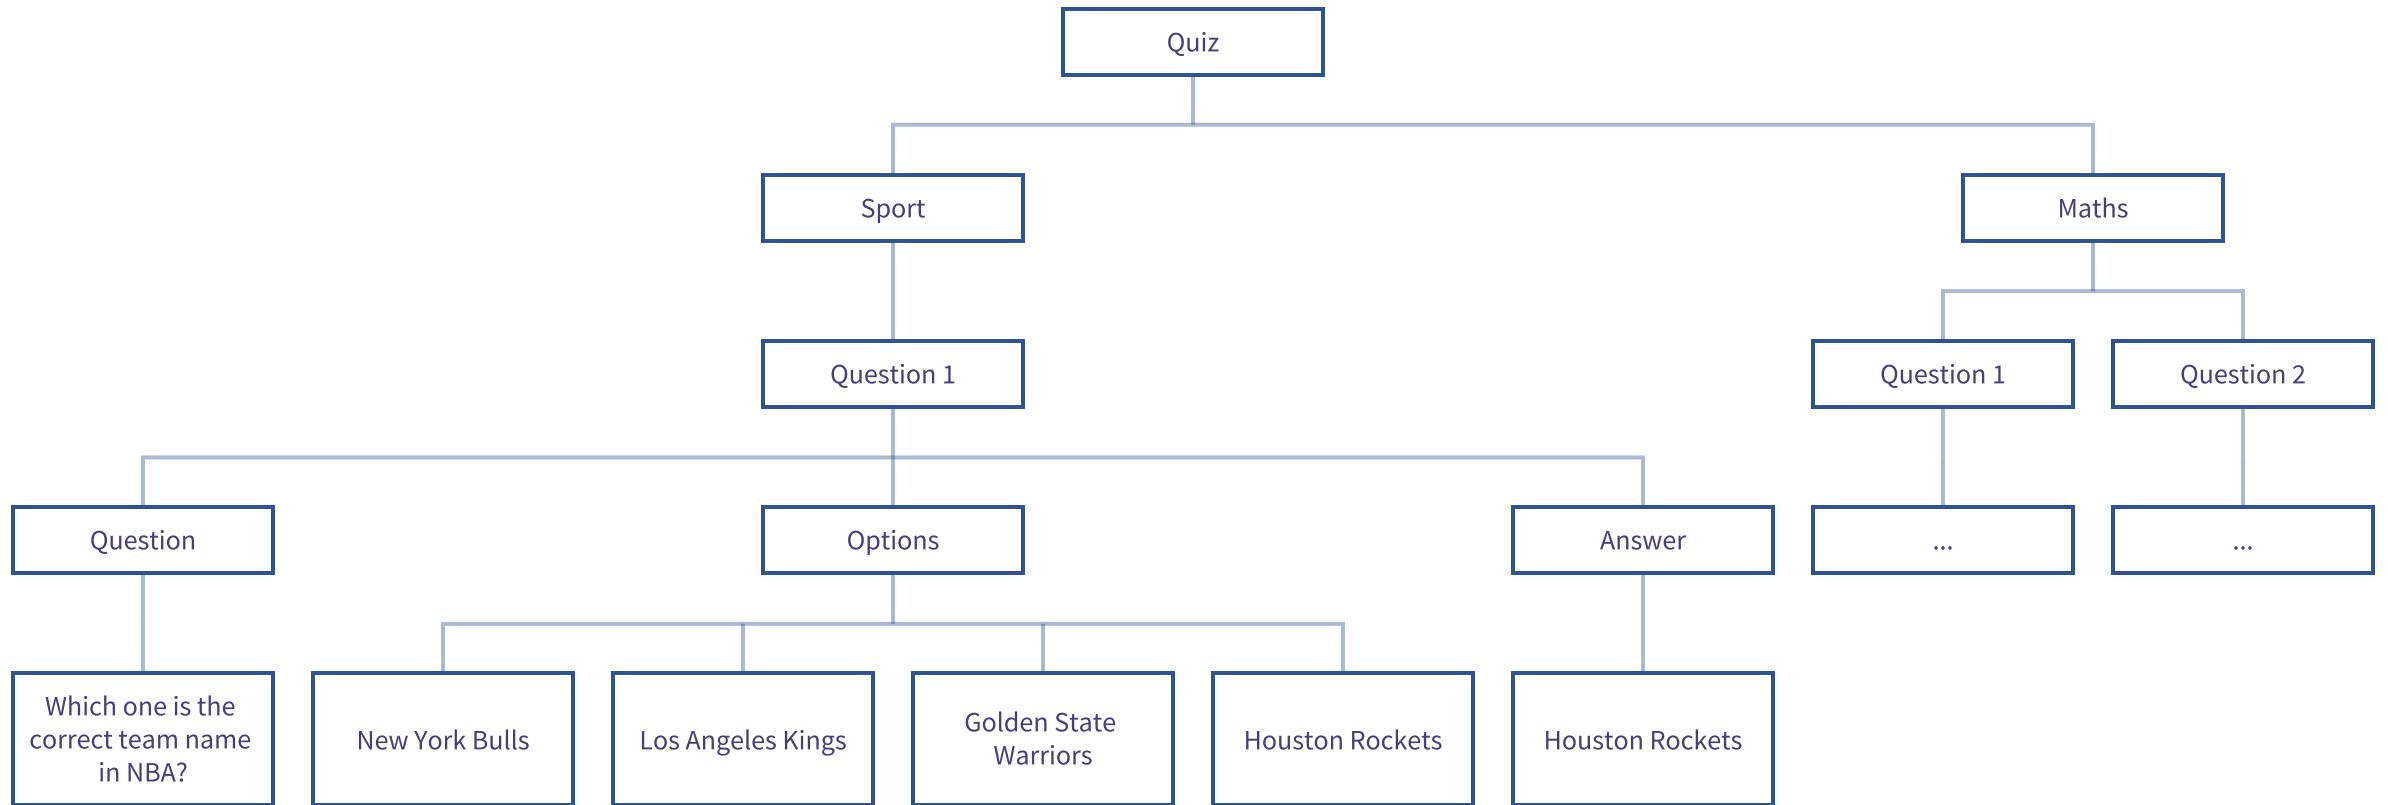
requires a schema to ingest

capable of holding data that does not fit within a two dimensional table structure well

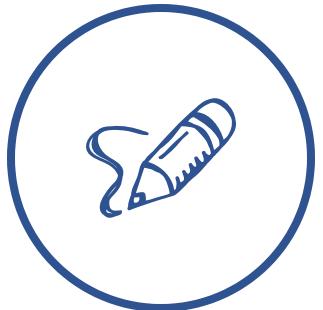
Example JSON file

```
{  
    "quiz": {  
        "sport": {  
            "q1": {  
                "question": "Which one is correct team name in NBA?",  
                "options": [  
                    "New York Bulls",  
                    "Los Angeles Kings",  
                    "Golden State Warriros",  
                    "Huston Rocket"  
                ],  
                "answer": "Huston Rocket"  
            }  
        },  
        "maths": {  
            "q1": {  
                "question": "5 + 7 = ?",  
                "options": [  
                    "10",  
                    "11",  
                    "12",  
                    "13"  
                ],  
                "answer": "12"  
            },  
            "q2": {  
                "question": "12 - 8 = ?",  
                "options": [  
                    "1",  
                    "2",  
                    "3",  
                    "4"  
                ],  
                "answer": "4"  
            }  
        }  
    }  
}
```

Structure

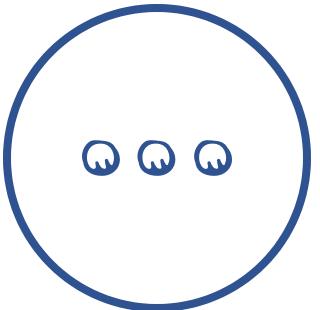


Filename Extensions



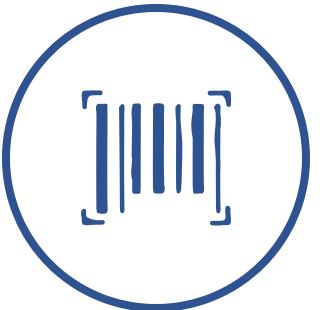
Plain Text

.txt
.log
no extension



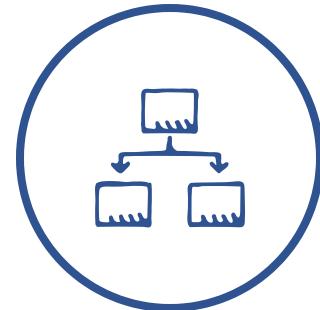
Character Separated
Value

.csv
.psv
.tsv
.log



Binary

.feather
.pickle
.exe
.dat
.jpg
.gif
.png



Structured

.xml
.json
.avro
.parquet
.orc

File Overview

	Structure	Ease of use	Size
Plain Text	Poor	Fair to Poor	Fair to Good
CSV	Fair	Fair to Good	Fair to Good
Binary	Poor to Good	Good	Poor to Fair
Structured	Good	Fair to Good	Poor

Terminal (Shell) scripting

Good for simple piping

Historically easier in the UNIX terminal, but Windows command-line is quite capable now.

Many shells have been developed over the last 50 years, but the predominant linux shell is the Bourne Again Shell (bash).

Alternatives include the C Shell (csh), and tenex c shell (tcsh), Hamilton c shell and KornShell (ksh)

Windows runs on its own terminal, but can now support bash, and offers powershell (but beware!)

MacOS typically runs with bash

if in doubt - bash it

common data manipulation bash commands

- **cat**

concatenate - print each row of all specified files

- **wget**

web-get. Download data from a URL

- **sort**

sorts data. use -n for a numeric sort

- **uniq**

assimilates duplicated values in a list. only sorts adjacent rows
so sort first!

- **head**

prints the top of the file (default top 10 rows)

- **tail**

prints the bottom of the file (default 10 rows)

- **wc**

word count. Can count by file and row.

- **grep**

Global Regular Expression Print:filter to select specific rows by a search pattern

- **awk**

file pattern scanning and processing

- **touch**

update the modification date on a file

- **sed**

Stream Editor: text transformations on a file or input stream

Data flow in Bash

| pipe

connects the output of one command to the next (STDOUT to STDIN)

> insert

Writes the output of one command to a text file, typically, but can also be used to insert parameters to a further command.

|& pipe error

connects the error and output of one command to the next (STDOUT and STDERR to STDIN)

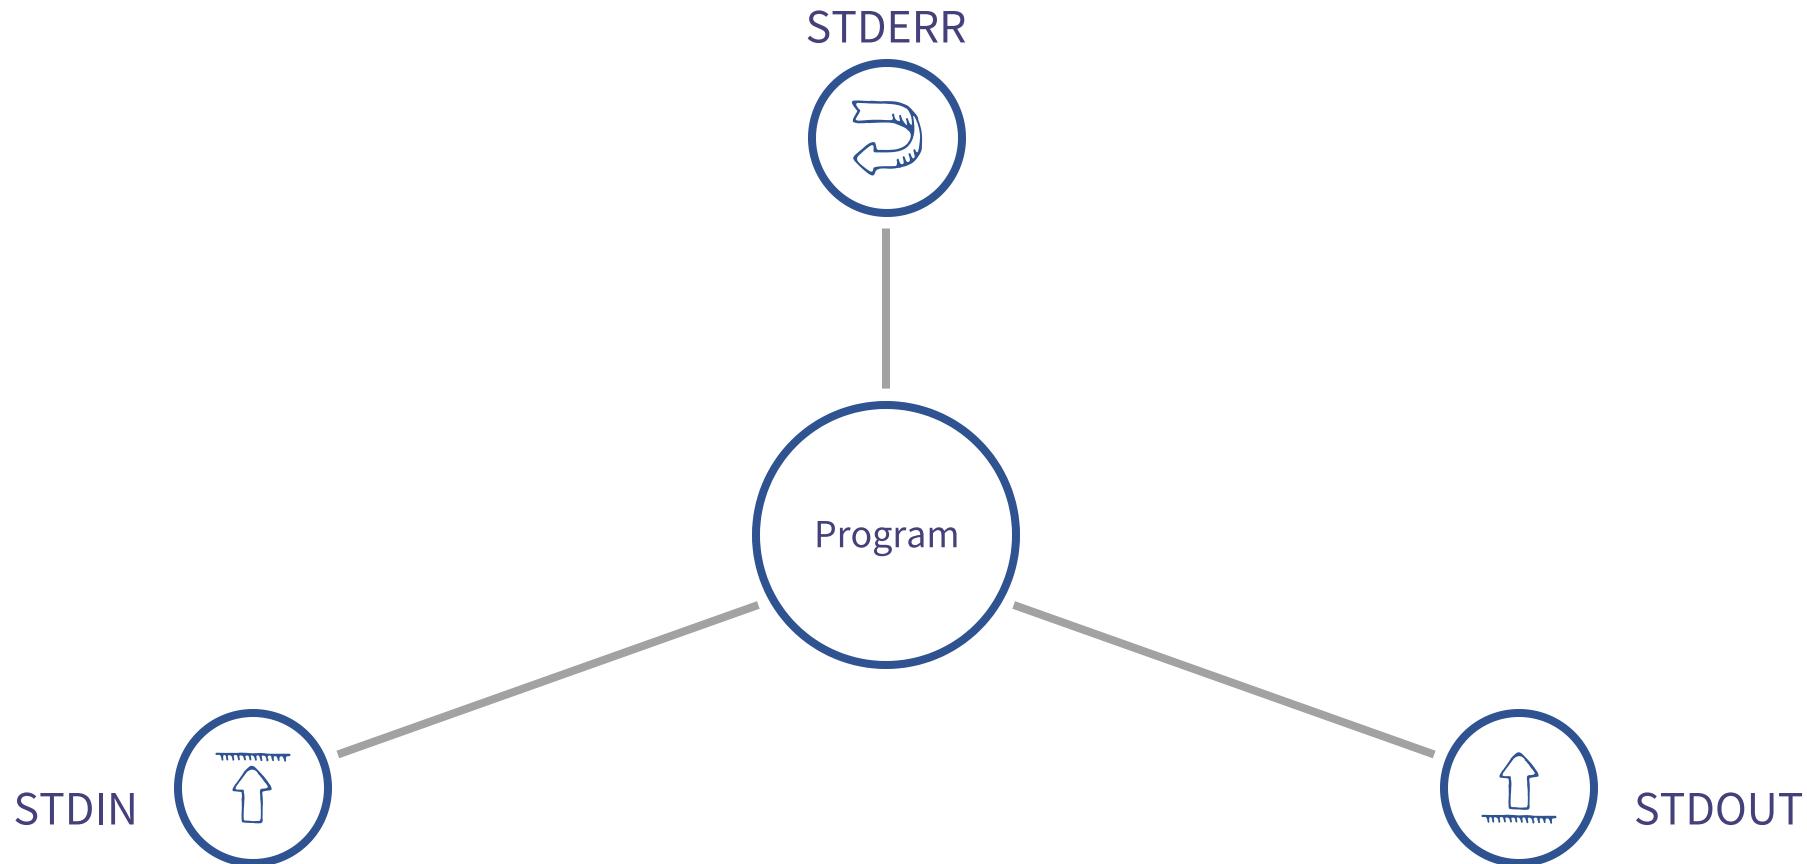
< insert back

Applied the output of one command back to the previous step. Can be used to insert data, parameters or to write to the previous file.

>> insert append

When data is streamed into a text file, any new data will append not overwrite.

POSIX compliance



Shell (BASH)

```
cat ./data_directory/datafile.csv | head -5 | grep 'search_string' > results.txt
```

Shell scripts allows simple files to be manipulated from the command line.

allows simple error handling
runs quickly, and compatible with a wide range of programs

Difficult for complex operations

Advanced shell wizardry

```
file_list = os.listdir(".")

f = FloatProgress(min=0, max=len(file_list))
display(f)
total_no_files = len(file_list)
for item in file_list:
    awk_command="""awk 'NR > 30 {}' {} | tr -d '\n' | awk 'gsub(" +", "\n") {}' > ./headless_only/{}""".format('{ print }', item, '{ print }', item)
    os.system(awk_command)
    f.value += 1
```

Common Data management Languages

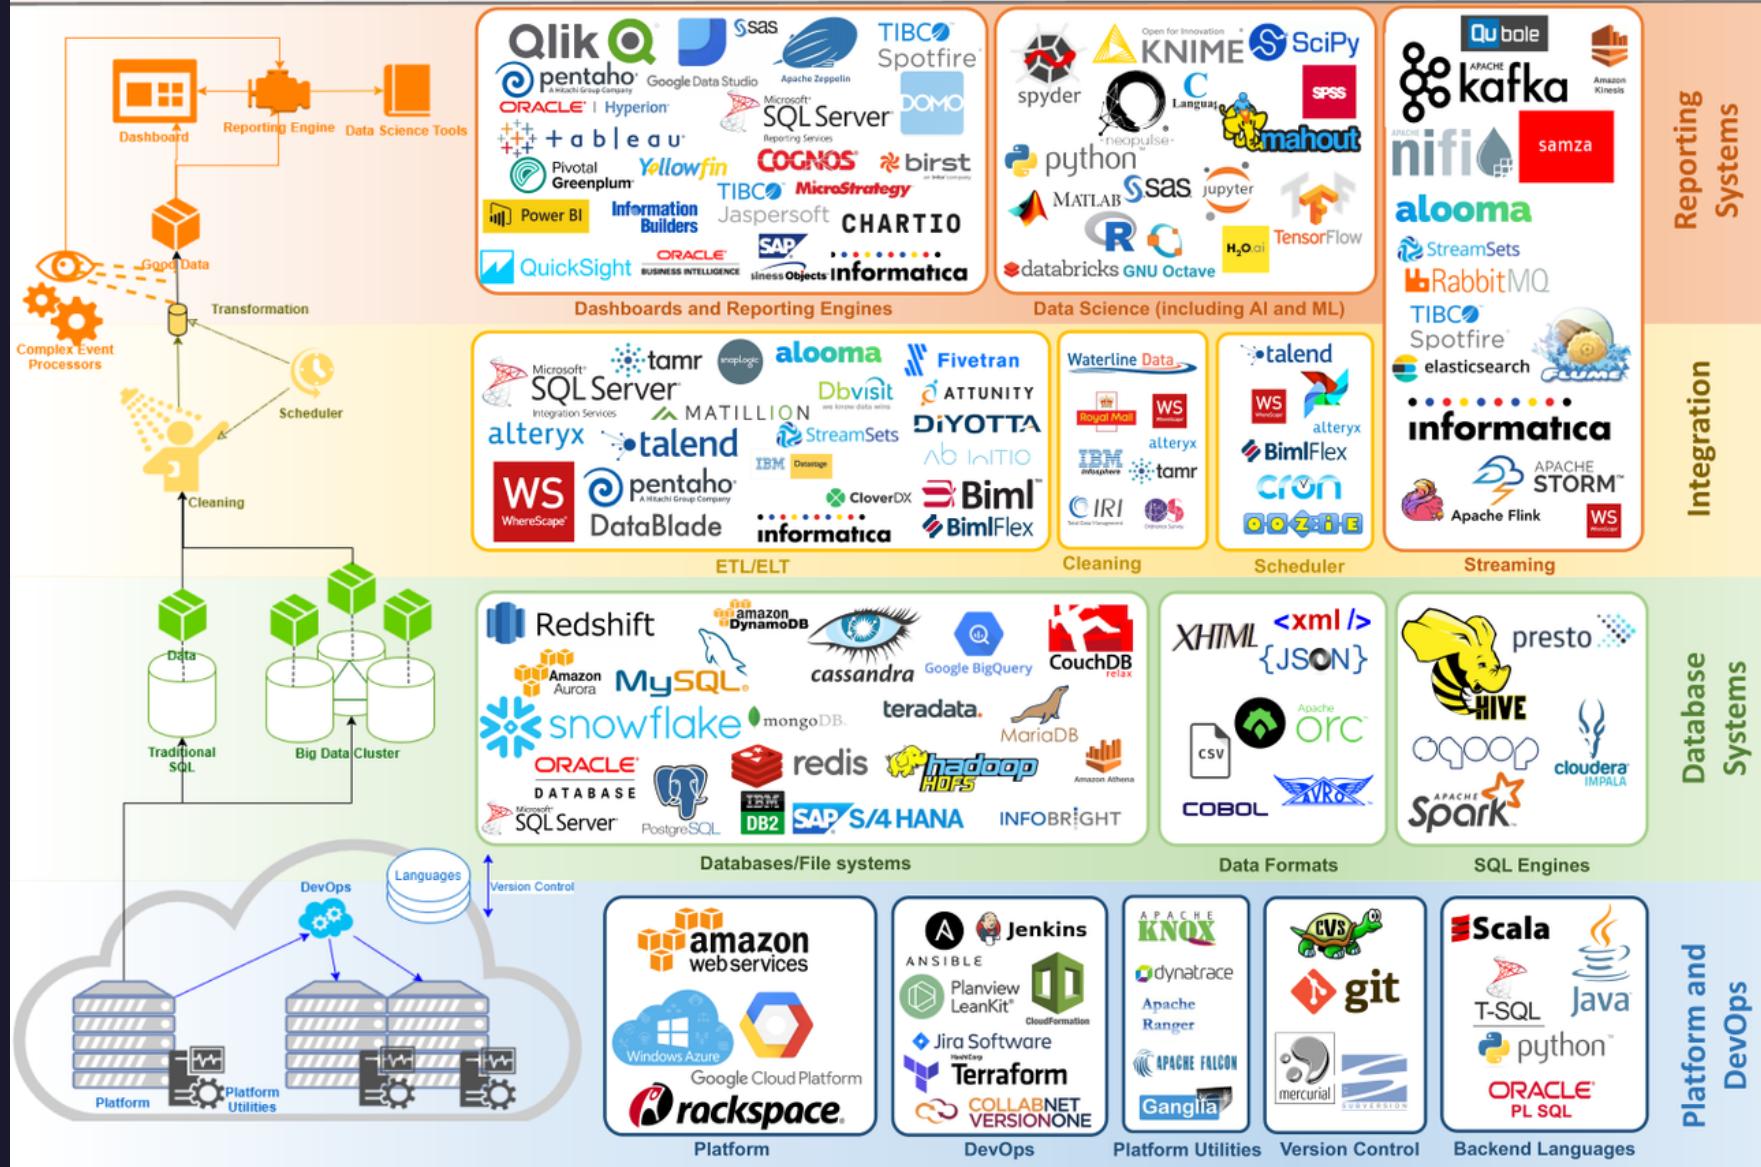


More Languages

- Python
- R
- TSQL, PL/SQL, pgSQL...
- SCALA
- Java
- Julia
- Perl
- MatLab
- Octave
- C/C++

Data Technology Landscape

Winter 2018/2019



Working with python

- **Use notebooks**

Store notes and narrative in addition to code during development.

- **Use Pandas**

There are wide-ranging data profiling and cleaning tools available

- **Use Numpy**

NumPy changes how python handles variables for scientific purposes. No reason not to use it

- **Use Python 3.x**

Major libraries now exist in Python 3. Python 3 runs faster and cleaner.

- **Watch out for old code**

Python 2 code \neq Python 3 code

- **Love your libraries**

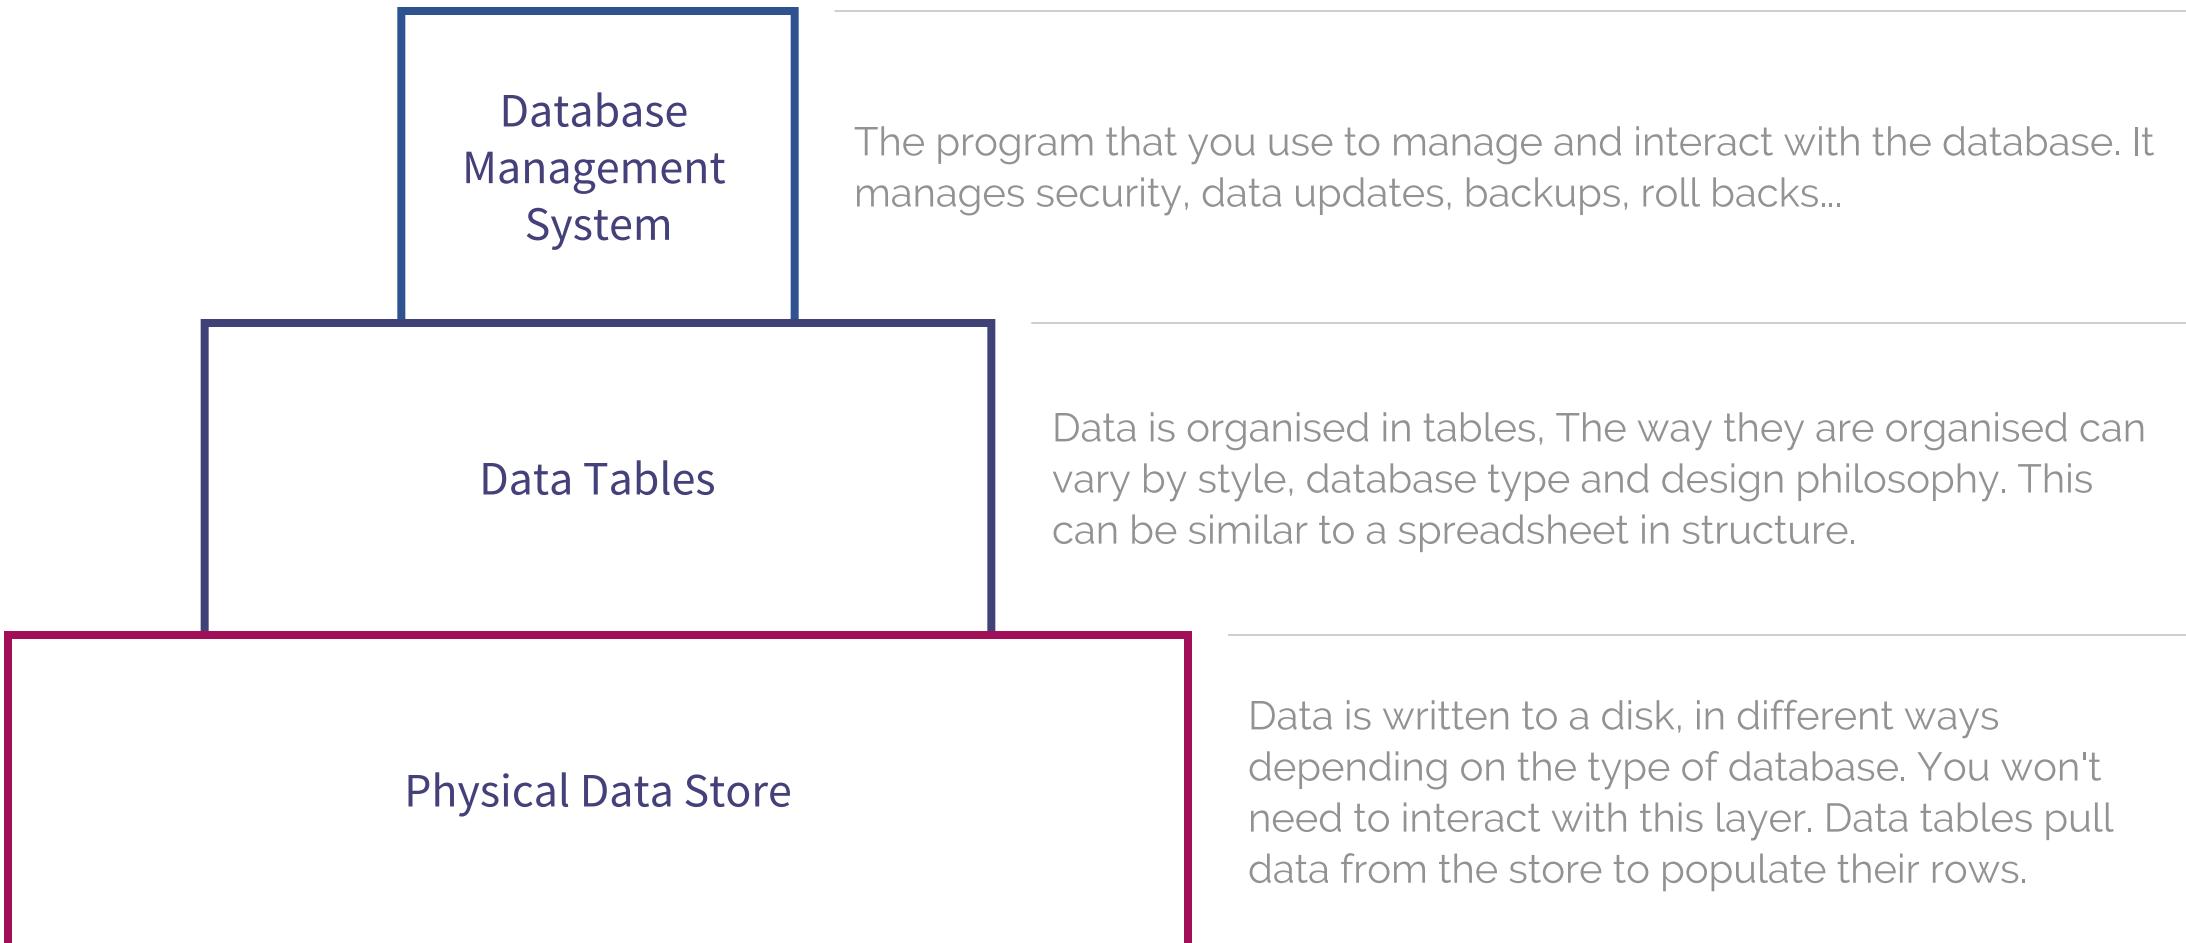
There are thousands of libraries available. Explore PyPi, PIP, Conda, Github.... to find them

File handling in python

```
In [ ]: file_list = os.listdir("../DATA/ASCII/AllData/headless_only/TimeseriesAndAmplitude/")[:-1] #to remove the FFT_db directory  
at the end of the list  
input_file_list = map(lambda x: "../DATA/ASCII/AllData/headless_only/TimeseriesAndAmplitude/" + x, file_list)  
output_file_list = map(lambda x: "../DATA/ASCII/AllData/headless_only/TimeseriesAndAmplitude/FFT_db/" + x[:-3] + "csv", file_list)  
  
i = 0  
for file_path in tqdm(file_list):  
    input_filepath = input_file_list[i]  
    output_filepath = output_file_list[i]  
  
    fft_kaiser_db(input_filepath=input_filepath, output_filepath=output_filepath, kaiser_betamax_value=14, time_interval=0.01)  
    i+=1
```

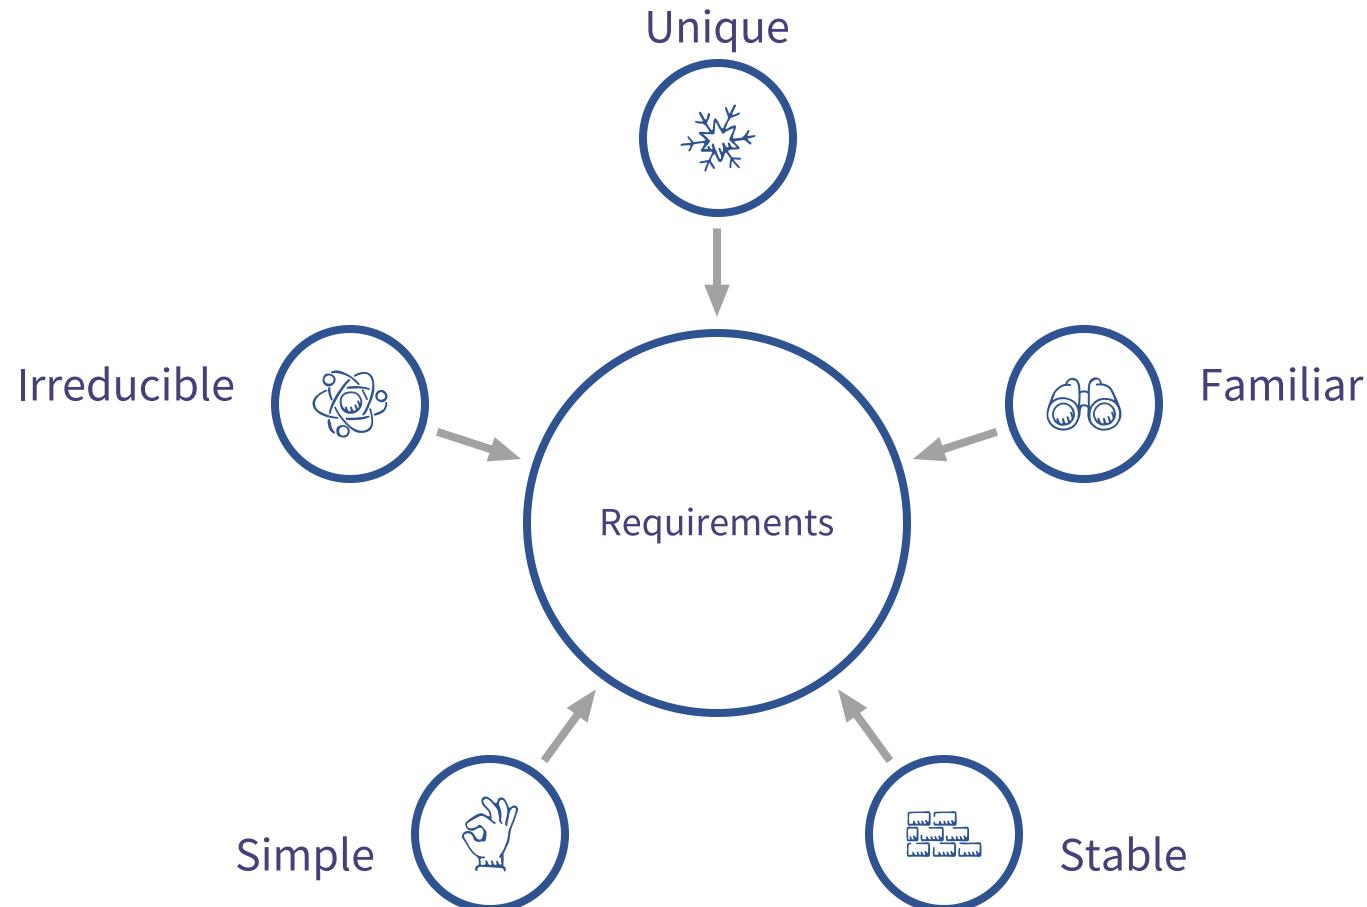
What is a database?

Database Layers



What do you need in a table?

Keys - Requirements



Keys - Types

Primary Key

the principal identifier for a table. Every piece of information should be associated with a PK. Values must be unique and NonNull

Foreign Key

The Primary Key of another table, used to identify related data across a database. Helps with joins.

Natural Key

A type of identifier that uses a meaningful attribute to identify each row. Examples include NI number, Student Number, ISBN. Can reduce the number of joins needed.

Surrogate Key

As opposed to a natural key, a surrogate key is a randomly generated unique value that acts as a key. Usually requires a new column to be added, but changing the data structure later won't break the key.

Normal Forms



1st Normal Form

Data is in a table

Values in the table are “atomic” - they can't be subdivided

Columns do not repeat

Data is being organised into meaningful structures!

Normal Forms



2nd Normal Form

All the columns in the table have to rely on the Primary Key

Each table holds information about only one thing!

Normal Forms



3rd Normal Form

No columns of data are transitively dependent on the Primary Key

Each piece of information about the subject has no hidden links!

UnNormalised Data

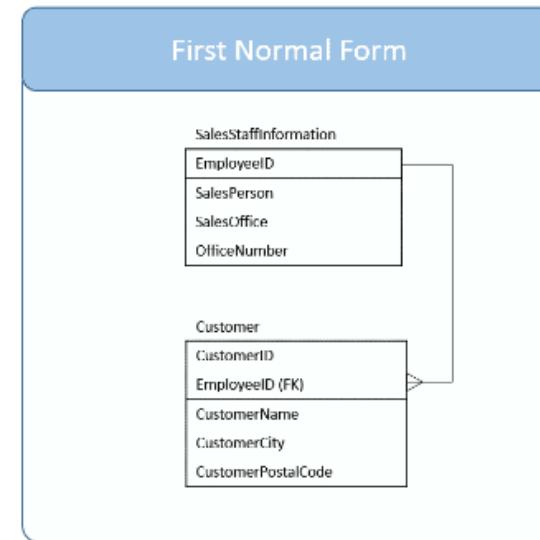
SalesStaff						
<u>EmployeeID</u>	SalesPerson	SalesOffice	OfficeNumber	Customer1	Customer2	Customer3
1003	Mary Smith	Chicago	312-555-1212	Ford	GM	
1004	John Hunt	New York	212-555-1212	Dell	HP	Apple
1005	Martin Hap	Chicago	312-555-1212	Boeing		

1st Normal Form

SalesStaffInformation			
<u>EmployeeID</u>	<u>SalesPerson</u>	<u>SalesOffice</u>	<u>OfficeNumber</u>
1003	Mary Smith	Chicago	312-555-1212
1004	John Hunt	New York	212-555-1212
1005	Martin Hap	Chicago	312-555-1212

Note: Primary Key: EmployeeID

Customer				
<u>CustomerID</u>	<u>EmployeeID</u>	<u>CustomerName</u>	<u>CustomerCity</u>	<u>PostalCode</u>
C1000	1003	Ford	Dearborn	48123
C1010	1003	GM	Detroit	48213
C1020	1004	Dell	Austin	78720
C1030	1004	HP	Palo Alto	94303
C1040	1004	Apple	Cupertino	95014
C1050	1005	Boeing	Chicago	60601



2nd Normal Form

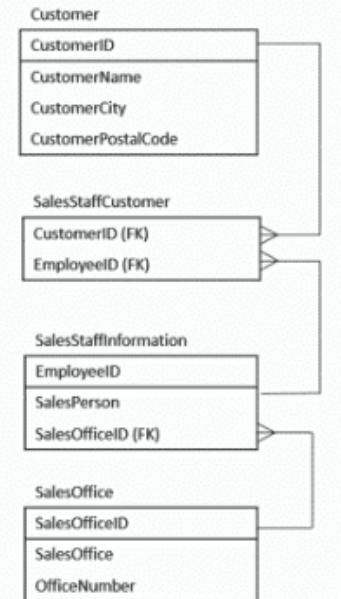
Customer			
CustomerID	CustomerName	CustomerCity	CustomerPostalCode
C1000	Ford	Dearborn	48123
C1010	GM	Detroit	48213
C1020	Dell	Austin	78720
C1030	HP	Palo Alto	94303
C1040	Apple	Cupertino	95014
C1050	Boeing	Chicago	60601

SalesStaffCustomer	
CustomerID	EmployeeID
C1000	1003
C1010	1003
C1020	1004
C1030	1004
C1040	1004
C1050	1005

SalesStaffInformation		
EmployeeID	SalesPerson	SalesOffice
1003	Mary Smith	S10
1004	John Hunt	S20
1005	Martin Hap	S10

SalesOffice		
SalesOfficeID	SalesOffice	OfficeNumber
S10	Chicago	312-555-1212
S20	New York	212-555-1212

Second Normal Form



3rd Normal Form

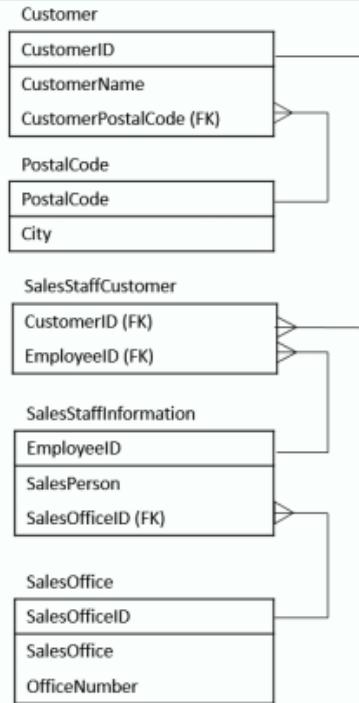
PostalCode	
PostalCode	City
48123	Dearborn
48213	Detroit
60601	Chicago
78720	Austin
94303	Palo Alto
95014	Cupertino

SalesStaffCustomer	
CustomerID	EmployeeID
C1000	1003
C1010	1003
C1020	1004
C1030	1004
C1040	1004
C1050	1005

SalesStaffInformation		
EmployeeID	SalesPerson	SalesOffice
1003	Mary Smith	S10
1004	John Hunt	S20
1005	Martin Hap	S10

SalesOffice		
SalesOfficeID	SalesOffice	OfficeNumber
S10	Chicago	312-555-1212
S20	New York	212-555-1212

Third Normal Form Issues



Structured Query Language

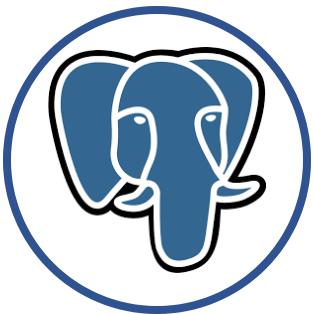
- Lets you modify data (Read, Write, Update, Delete)
- Lets you query data
- Human Readable
- Different Dialects! TSQL, PL/SQL...

Common Relational Database Programs



MySQL

Basic but functional database
uses SQL/PSM
Free



PostgreSQL

aka pgSQL and Postgres
uses PL/pgSQL
Stricter relational database
Free



Oracle DB

Enterprise grade database
uses T/SQL and PL/SQL
widely supported
Pricey!

Does have free variants (Oracle DB Express and more)

example Queries

- **SELECT * FROM Customer;**

Display all data from the Customer table

- **SELECT EmployeeID, SalesOffice FROM SalesStaffInformation**

Display the data from the Employee ID and The Sales Office columns ONLY from the Sales Staff Information table

- **SELECT * FROM SalesStaffInformation WHERE SalesPerson LIKE 'Mary*';**

Display data from all columns of the Sales Staff Information, but only the rows where the sales person is called Mary

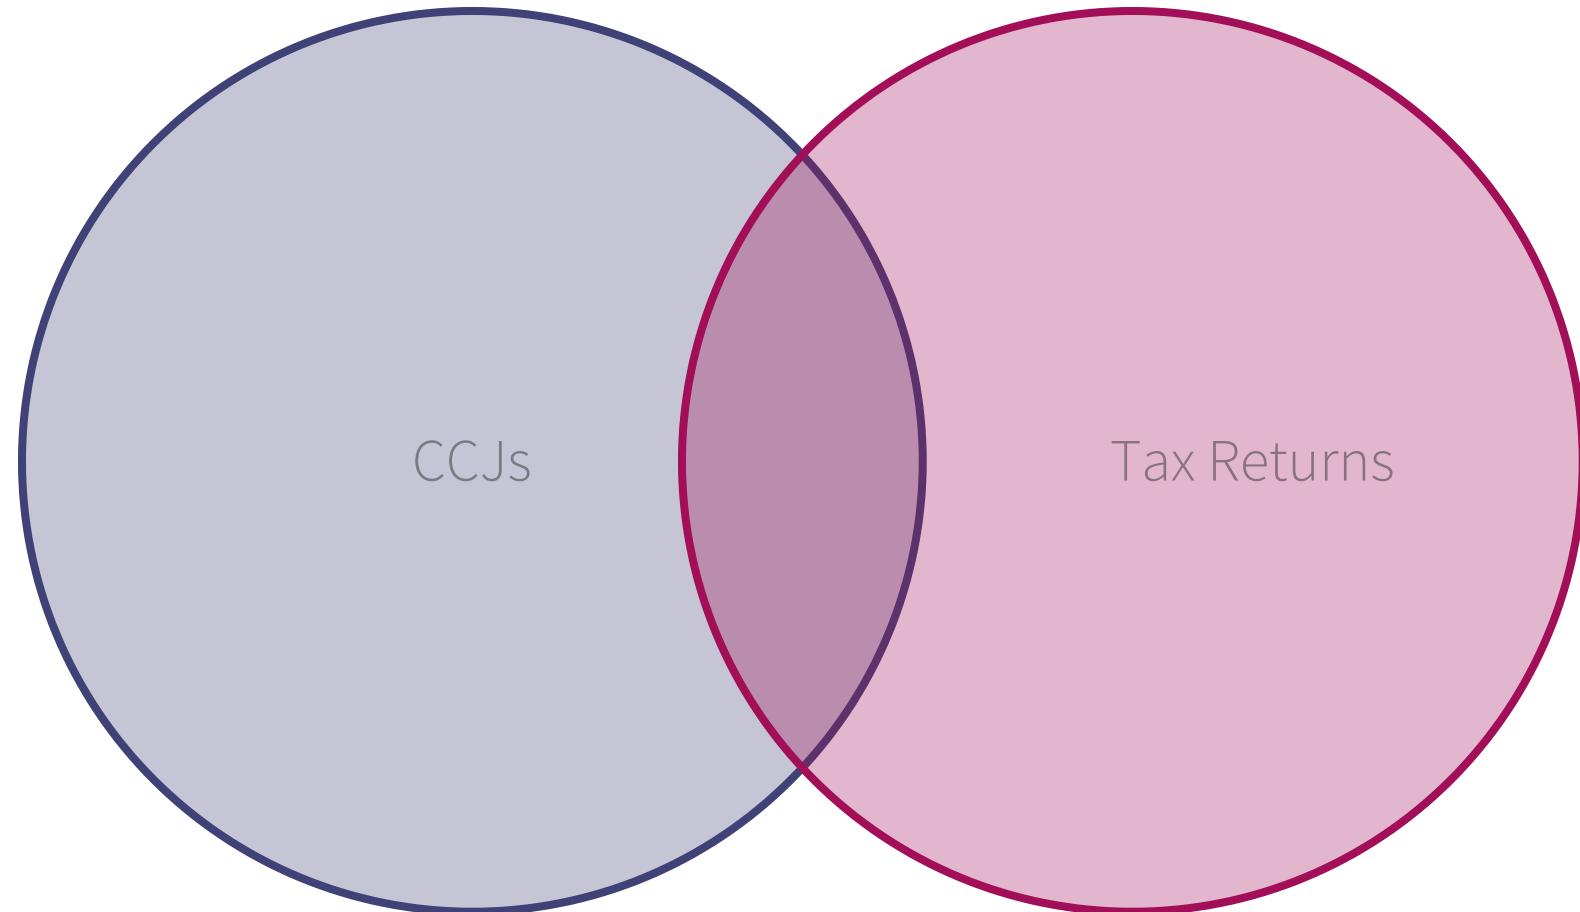
- **SELECT DISTINCT CustomerName FROM Customer;**

Display unique customer data from the customer table

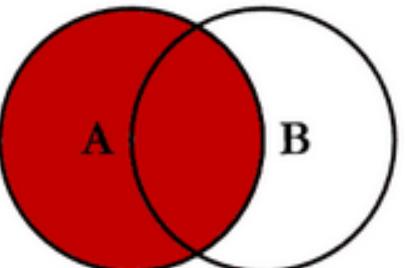
- **SELECT COUNT (DISTINCT CustomerName) FROM Customer;**

Returns the number of distinct customers in the Customer table

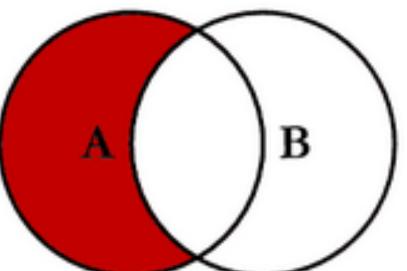
Joins



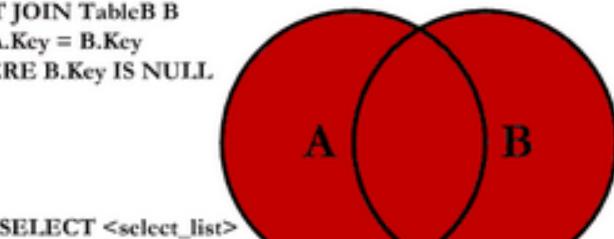
SQL JOINS



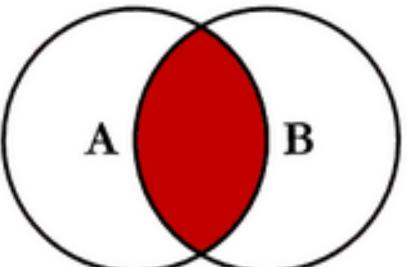
```
SELECT <select_list>
FROM TableA A
LEFT JOIN TableB B
ON A.Key = B.Key
```



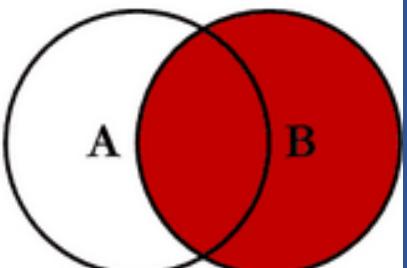
```
SELECT <select_list>
FROM TableA A
LEFT JOIN TableB B
ON A.Key = B.Key
WHERE B.Key IS NULL
```



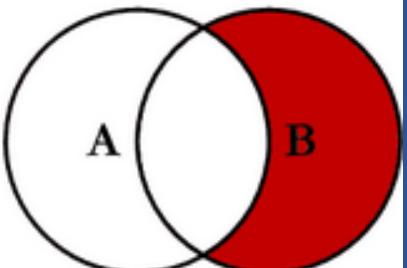
```
SELECT <select_list>
FROM TableA A
FULL OUTER JOIN TableB B
ON A.Key = B.Key
```



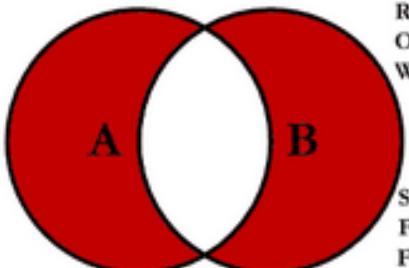
```
SELECT <select_list>
FROM TableA A
INNER JOIN TableB B
ON A.Key = B.Key
```



```
SELECT <select_list>
FROM TableA A
RIGHT JOIN TableB B
ON A.Key = B.Key
```



```
SELECT <select_list>
FROM TableA A
RIGHT JOIN TableB B
ON A.Key = B.Key
WHERE A.Key IS NULL
```



```
SELECT <select_list>
FROM TableA A
FULL OUTER JOIN TableB B
ON A.Key = B.Key
WHERE A.Key IS NULL
OR B.Key IS NULL
```

Common scenarios

Same Data, Different Name

Data gets duplicated, and inflates the row count.
This might encourage artificial trends in the data

Different Data, Same Name

Data might be over-written and lost. The identifier may be failing, and would cause issues if not considered when summing columns

Data only found in one table

Data may be missing in other tables, or simply does not exist. Can cause issues when joining tables

Different Keys, Same Data

Similar to the first scenario, but worse: the keys are the baseline identifier which makes differentiating between them difficult

General Tips

- **Use Snake Case**

when tables or column names cannot use spaces, complicated_column_name is easier to read than
complicatedcolumnname

- **make names (and variables) easy to understand**

sgl_arr_tme is difficult to understand.
signal_arrival_time is easier.

- **Append dates to table names if you are exporting or copying them**

table_name_vfinal_FINAL means less than
table_name_190329

- **Be careful where you leave your credentials**

Hard coding your access credentials into your scripts is a quick way to test your code and an even faster way to broadcast your keys

Break

Did you hear about the elf that lost his toys?



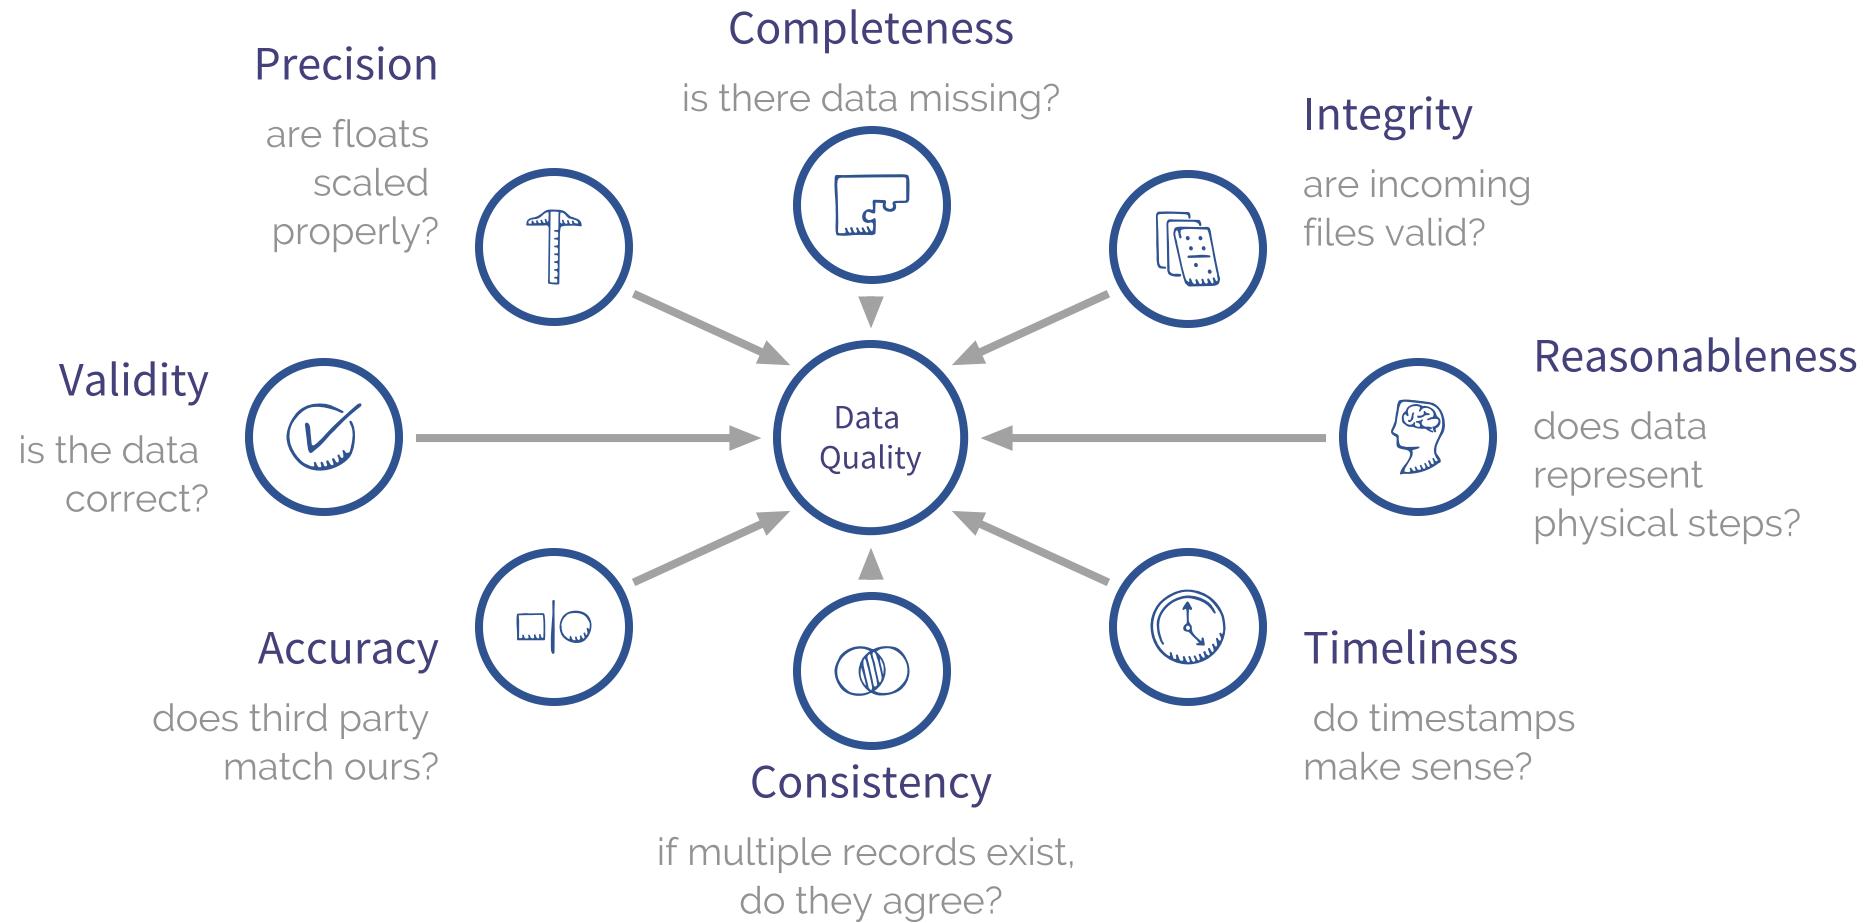
Did you hear about the
elf that lost his toys?

He was lego-less

Section 4

Data Profiling

Data Quality



DAMA Definition

Adapted from DAMA-DMBOK

① Timeliness

Does the data represent reality from the required point in time?

② Completeness

Does all the data for the event exist?

③ Uniqueness

Is anything recorded more than once?

④ Validity

Does the data conform to its syntax?

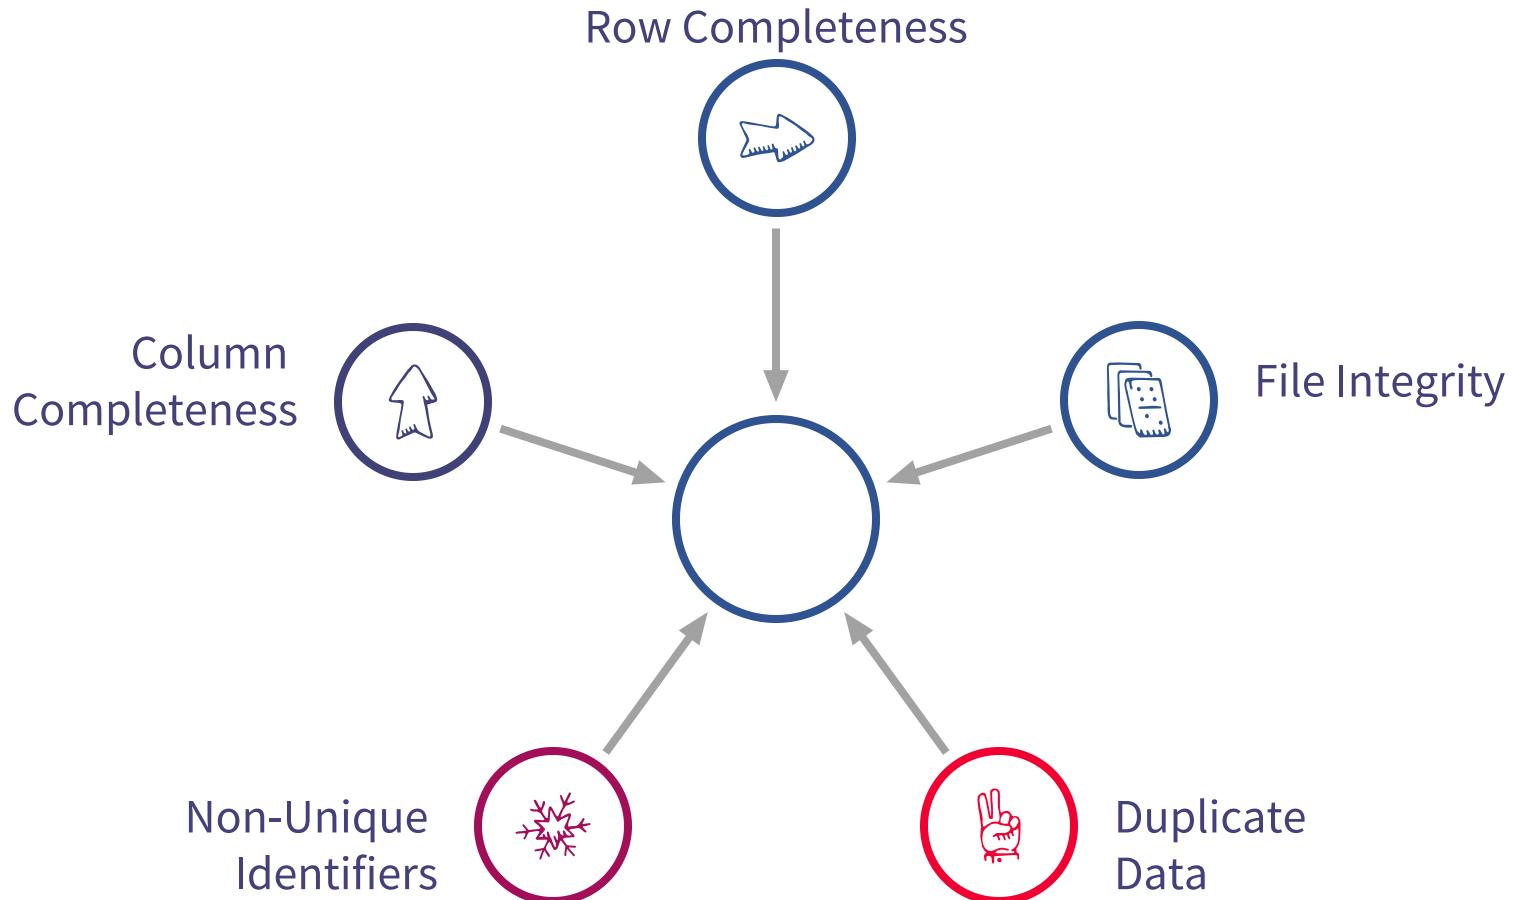
⑤ Consistency

Do multiple representations of a record differ?

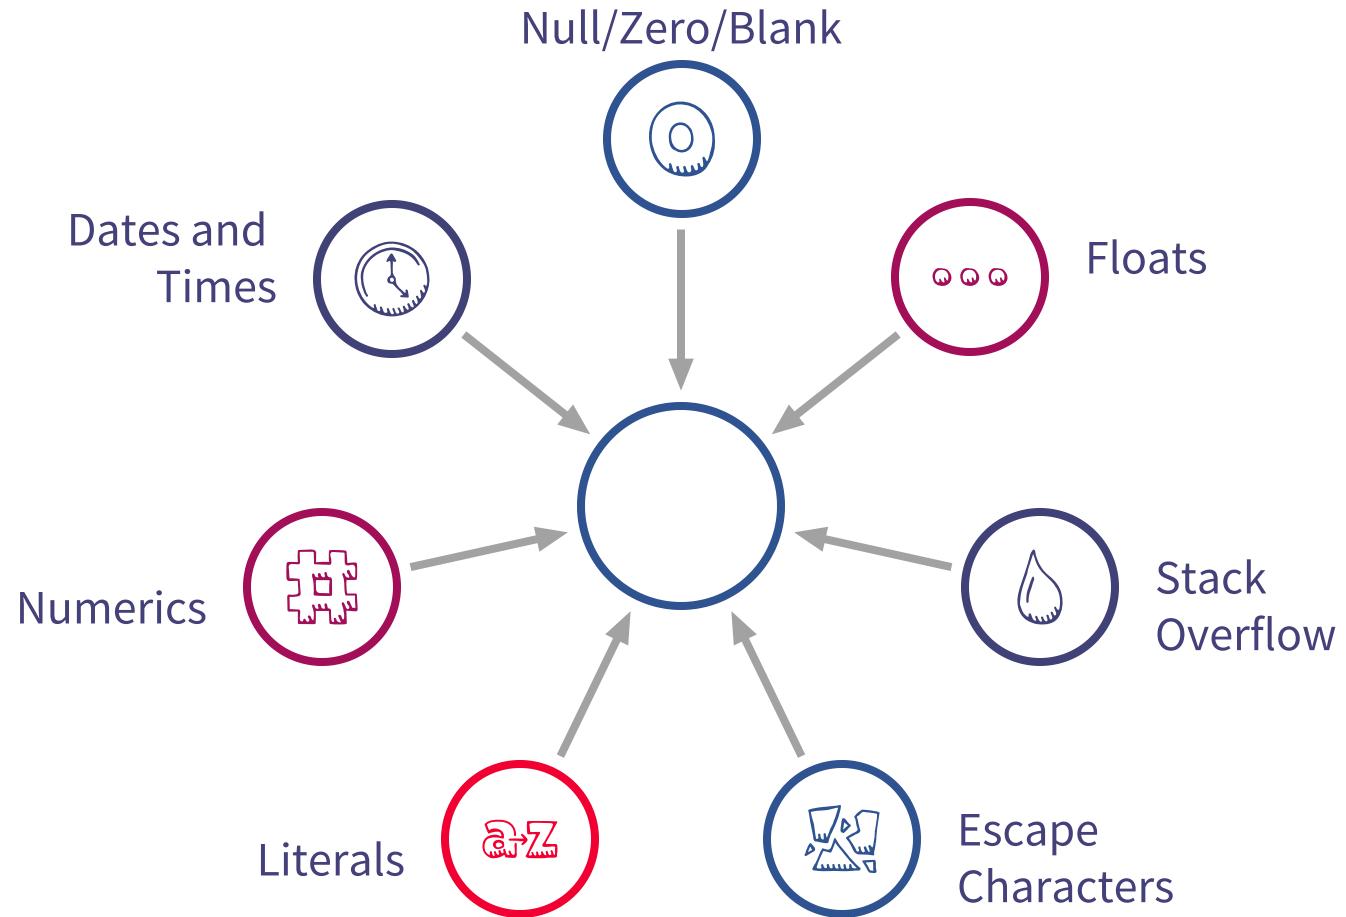
⑥ Accuracy

Does the data accurately describe the real world event or object?

Out-of-cell issues

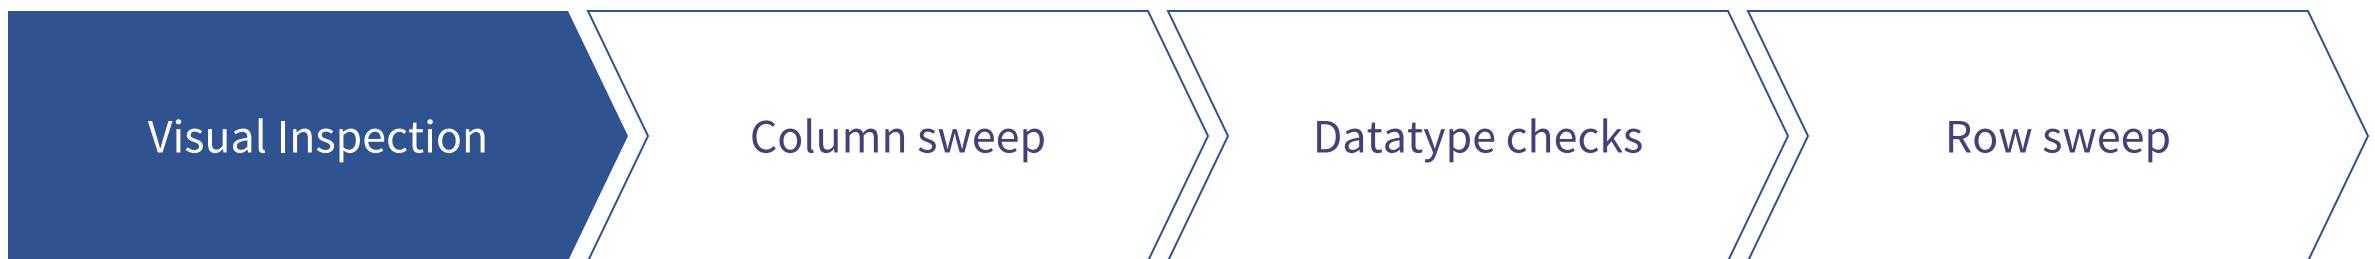


In-cell issues



Data Profiling

sample_df.head()



Data Profiling

sample_df.head()

Visual Inspection

Column sweep

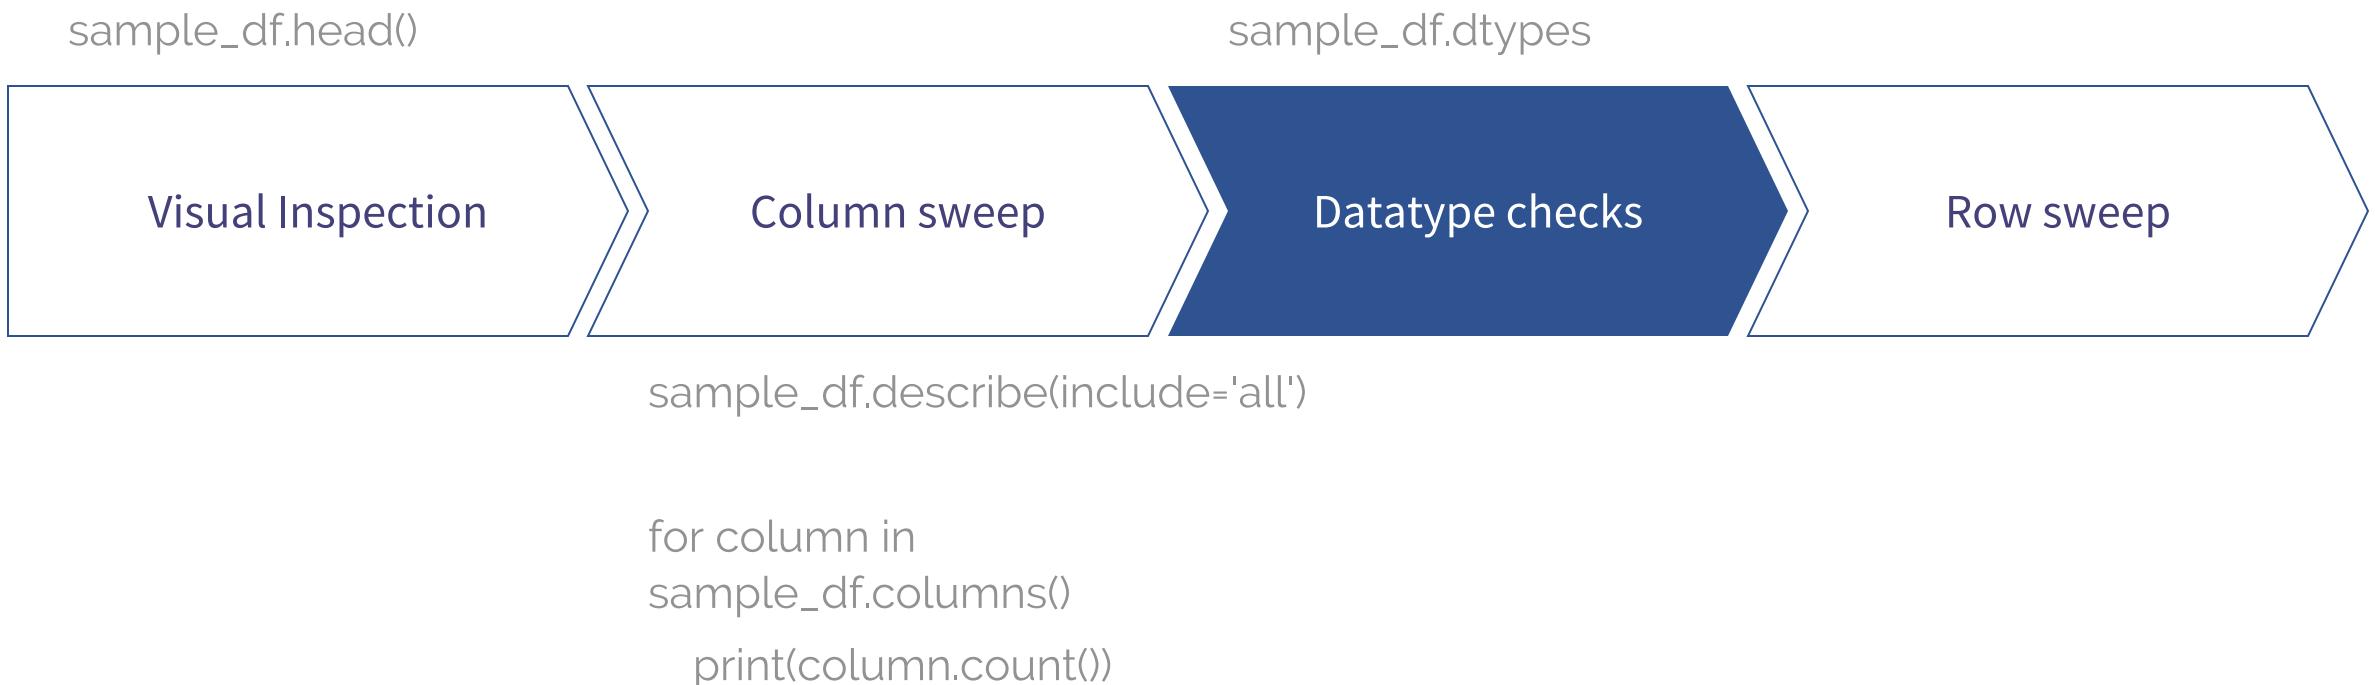
Datatype checks

Row sweep

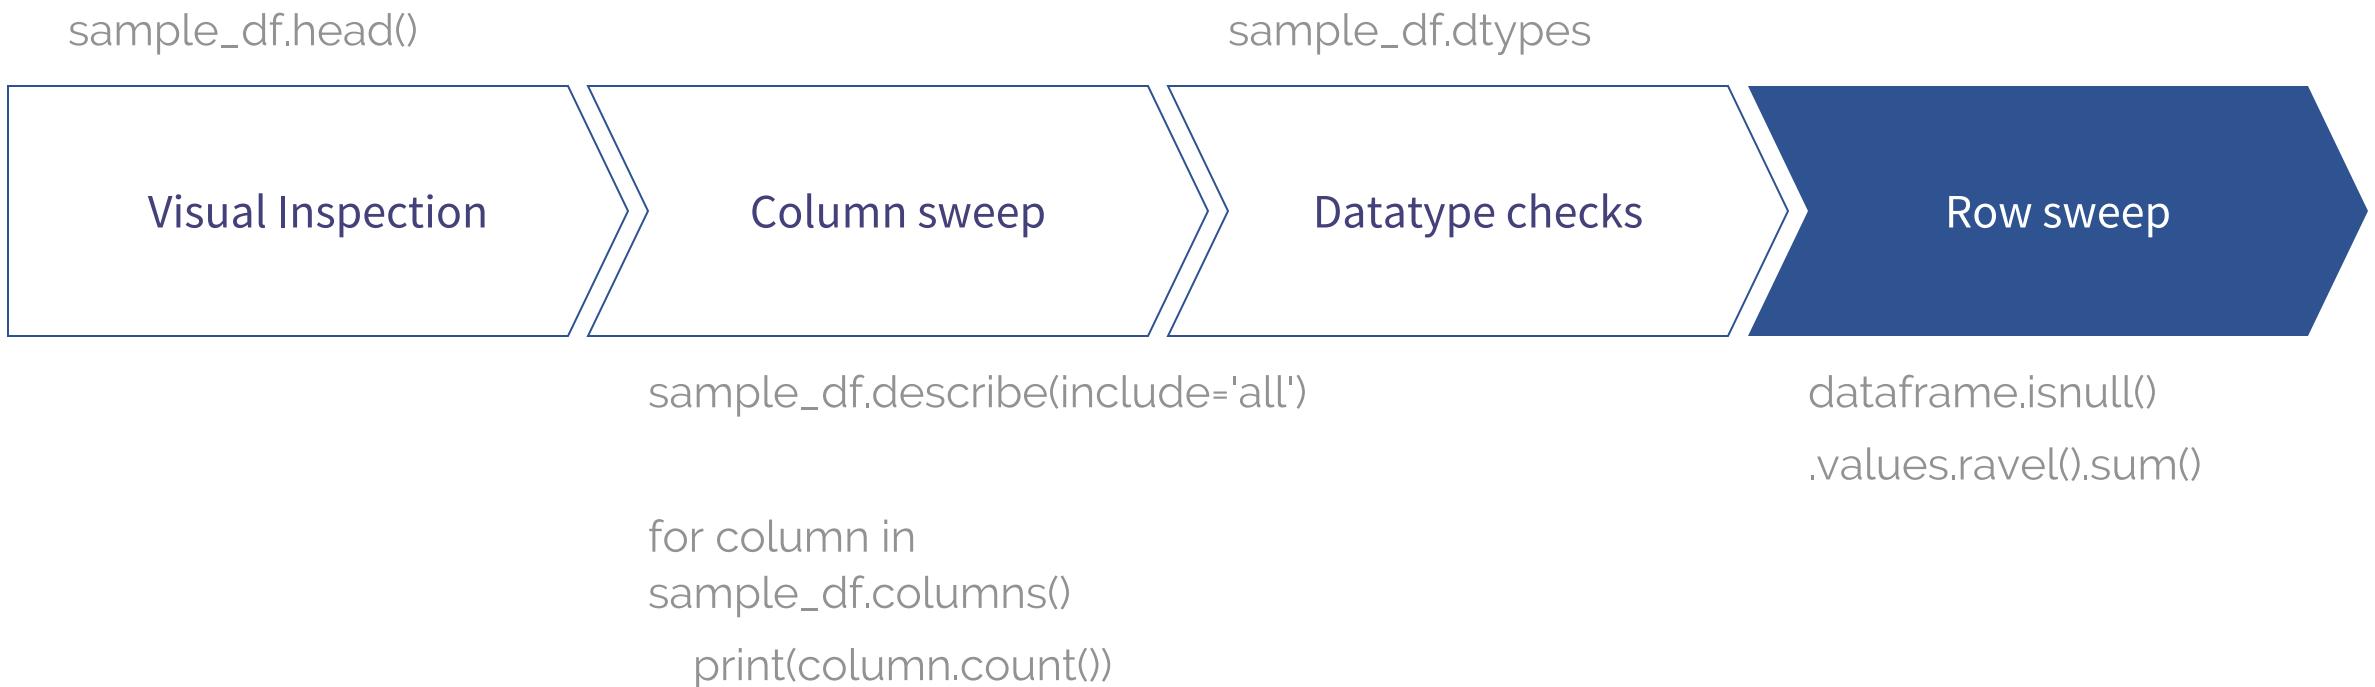
sample_df.describe(include='all')

```
for column in  
    sample_df.columns()  
        print(column.count())
```

Data Profiling



Data Profiling



Identifiers - unique?

- Pandas: `series.is_unique`
- Python: `len(x) == len(set(x))`
- SQL: `SELECT CASE WHEN count(distinct col)= count(col)
THEN 'pass' ELSE 'fail' END
FROM table;`

Null|NaN|Zero|None|“”[]{}

- **Zero**

is a value. Can be computed.

- **Null**

is not a value - nothing there

- **NaN**

Not a Number: can be a null or a dtype issue

- **None**

Python object for Null.

- “” “ ”

Empty - string. May be a space.

- []

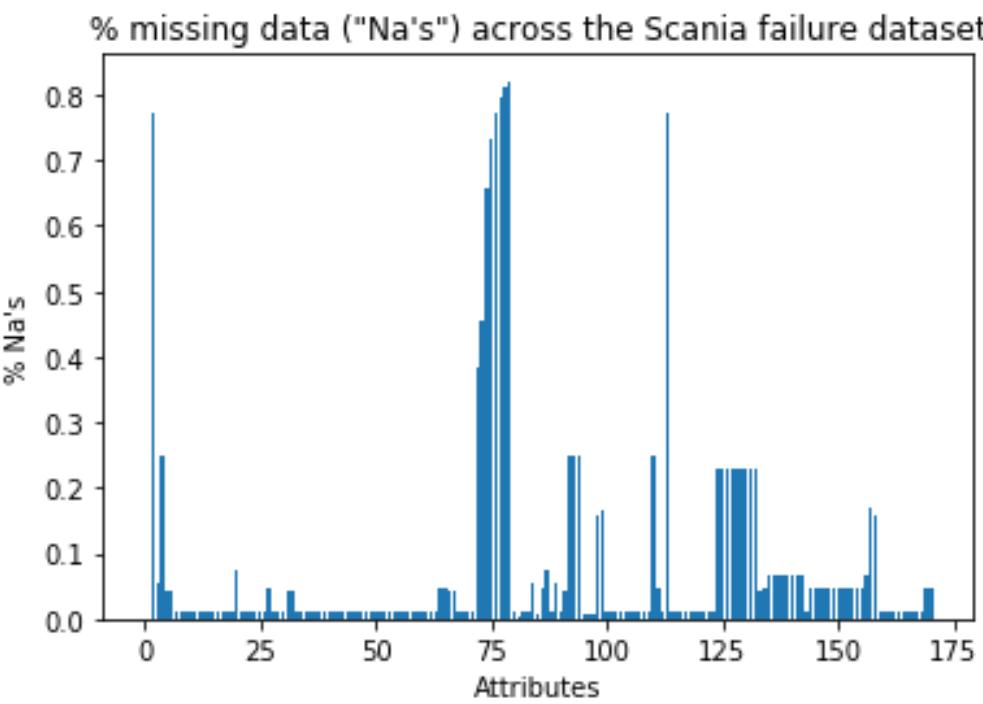
empty - list.

- {}

empty - dict.

Boolean Test

Element	True	False
0		✓
Null		✓
NaN		✓
None		✓
""		✓
""	✓	
[]		✓
{}		✓



Profiling Crib Sheet

Issue	When you know	how to check
Non Unique Identifiers/Duplicate Data	failure on load failure on join	count distinct
Completeness	failure on join failure on computation	count distinct count null values
Null/Zero	failure on computation failure to derive datatype	count distinct count null values targeted query
DateTime	failure on load failure on computation	check datatype regex for wrong format
Numerics	failure on load failure on computation	check datatype sample and plot values inspect
Escapes	failure on load failure on write	check for column length and row completeness regex/targeted query
Stack Overflow	failure on computation	plot range of values by column describe statistics/boxplot check datatype

Section 5

Data Cleaning

Casting - numerics

pushing one datatype into another

- pandas: pd.to_numeric(obj)
- python: int(), decimal(), float()
- SQL: CAST (x AS numeric(y))

Casting - datetimes

pushing one datatype into another

- pandas: pd.to_datetime(obj)
- python: strptime()
- SQL: TO_DATE (x, “format”)

YYYY-MM-DD HH:MM:SS:UU....

2018-04-08 10:30:00:00

2018-04-08 10:30:00:00 GMT

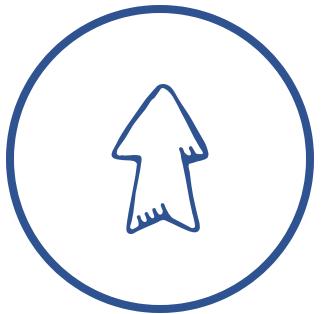
2018-04-08 05:30:00:00 -5

2018-04-08 05:30:00:00 EST

Missing Values

- Pandas: `df.skipna(axis="0 or 1",
how = "any")`
- Python: Boolean test
- SQL: `SELECT *
FROM table
WHERE column IS NOT NULL;`

Dropping missing/bad data

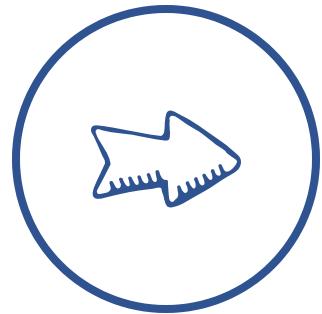


Drop by column

you do not lose rows of data

good for continuous time data

limits the information available for analysis



Drop by row

you do not lose information on each point

good for more complex/wide analysis

you have fewer datapoints

Use both

Drop very poor columns first

Follow with a row-wide sweep

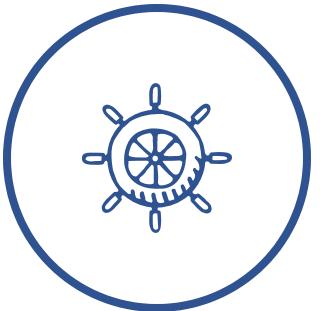
Identifiers - unique

- Pandas: `df.drop_duplicates`
- Python: `set(x)`
- SQL: `SELECT DISTINCT`

note

Database connectors

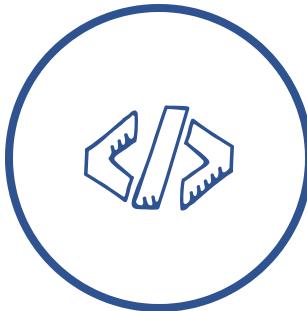
Connectors



Admin console

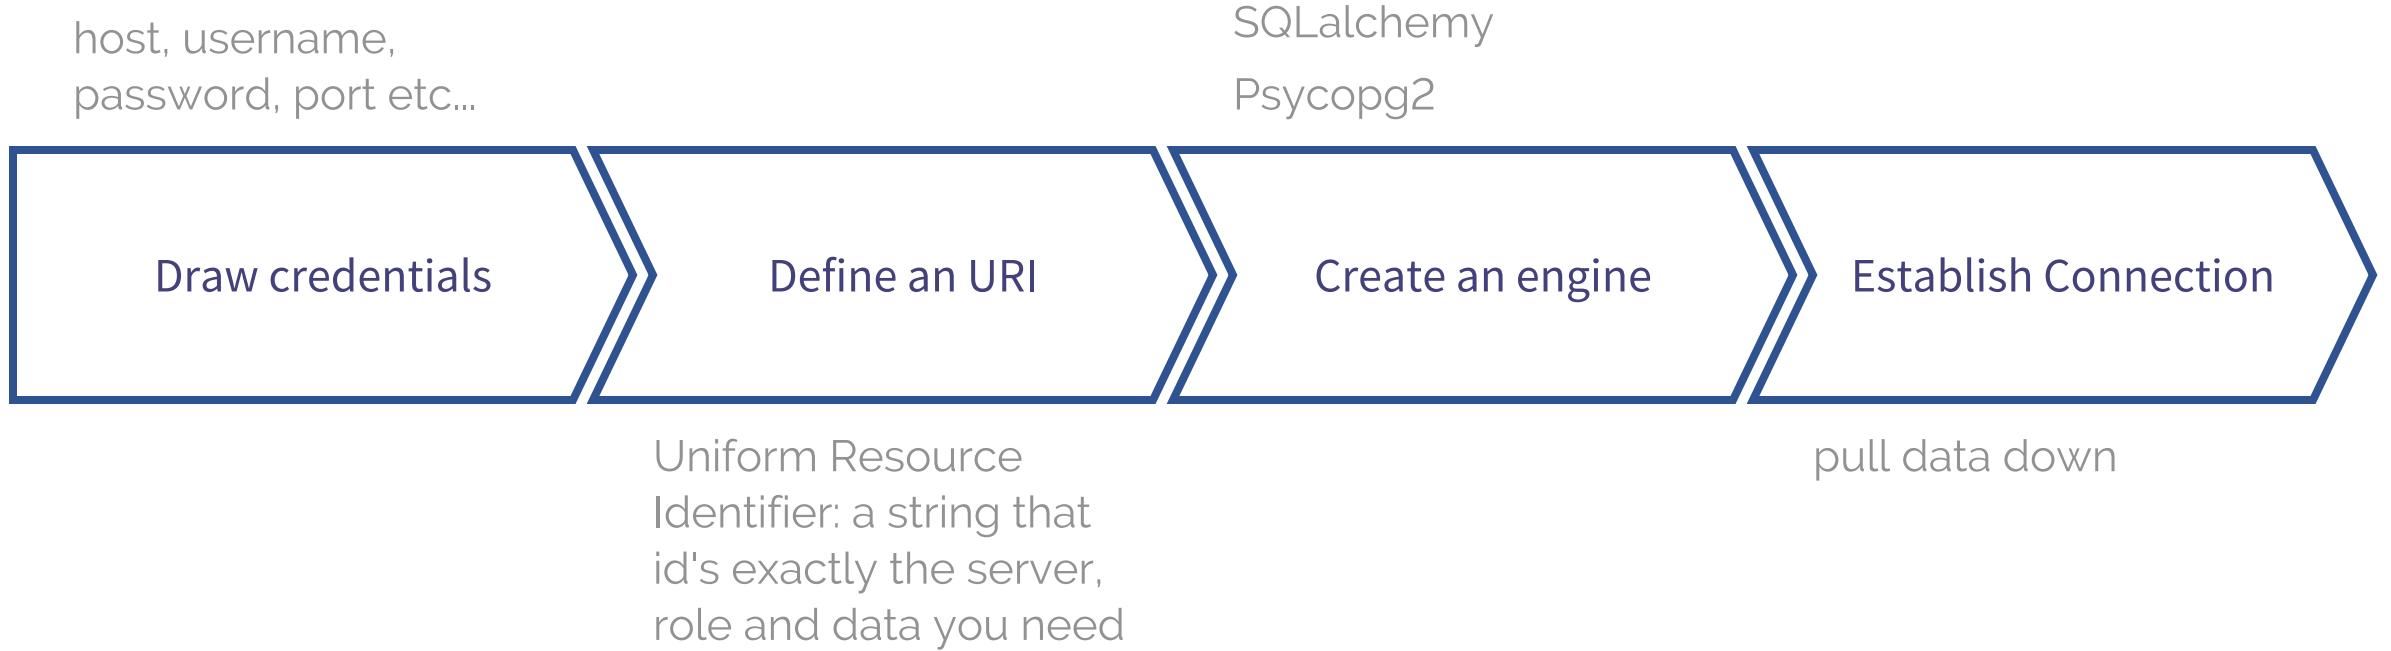


Terminal



Script

Connecting to an online database

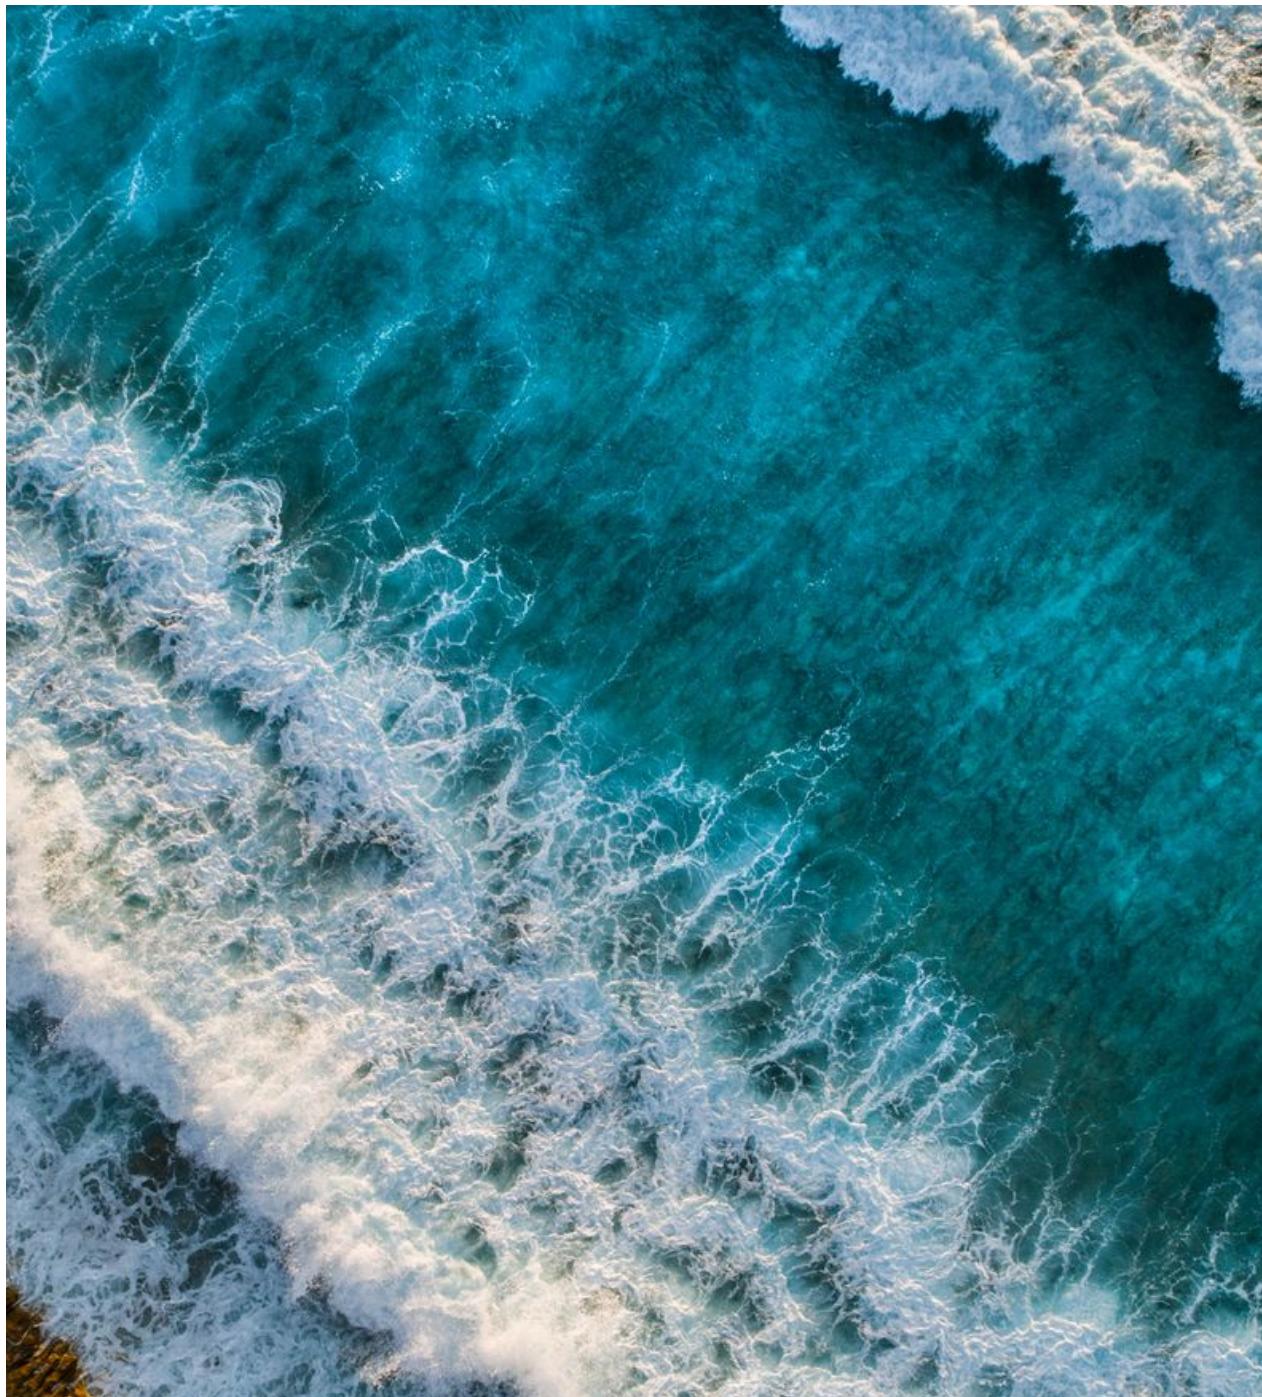


Your turn!

Objectives

- access a dataset from an online database
- profile the data
- if you find any issues, clean the data

[https://github.com/R-Strange/
Data_Wrangling_Course_Oxford](https://github.com/R-Strange/Data_Wrangling_Course_Oxford)



Recap

- What do pipelines look like?
- What do pipelines do?
- Extract-Load-Transform
- Flat Files
- Shell Scripting
- SQL Queries
- Database options
- Normal Forms
- Joins
- Profiling
- Cleaning
- DB Connectors

Fin

Go home - while you still can



Wrangling for Pipelines and Data Handling

8th April

Sorry!



Day outline

1 Data Flows Revisited

2 Testing

— Lunch —

3 Cloud

4 Big Data

5 IoT

6 APIs

7 API Lab

Section 1

Data Pipelines Revisited

Yesterday's Pipelines...

Break

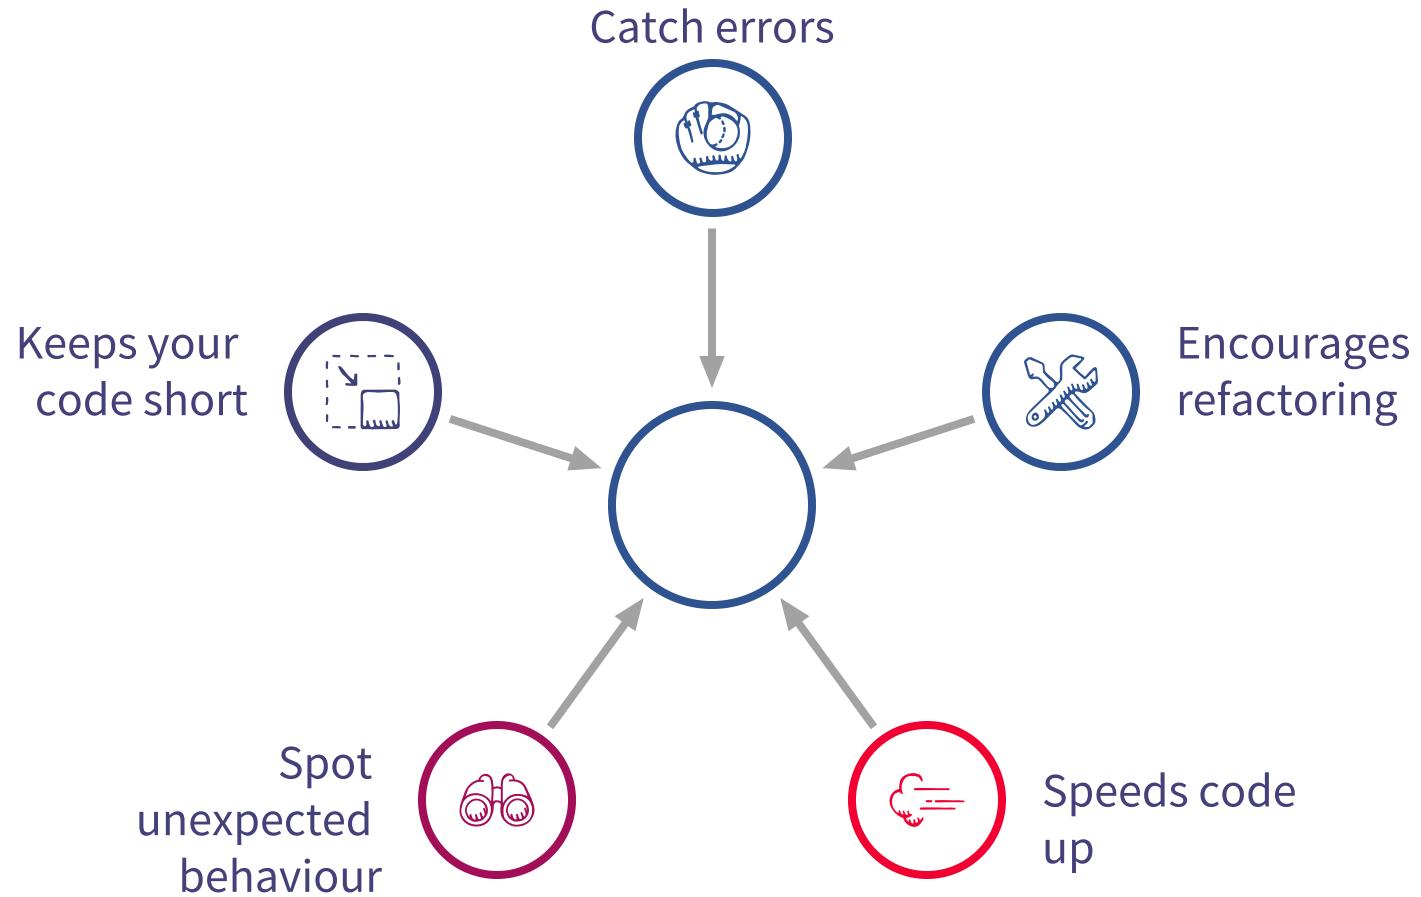


Section 2

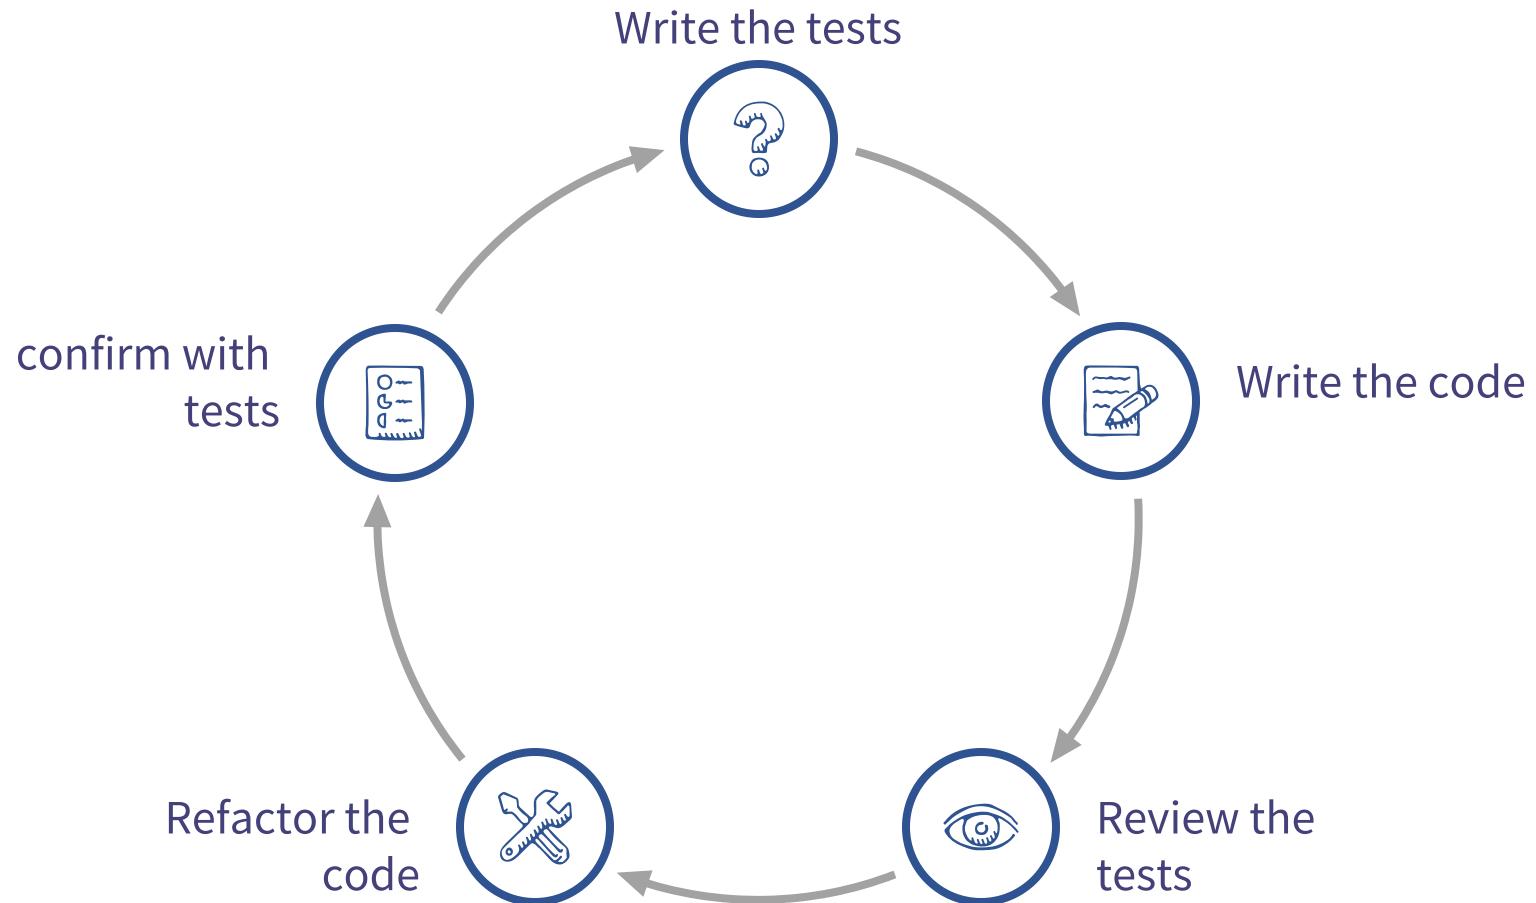
Testing



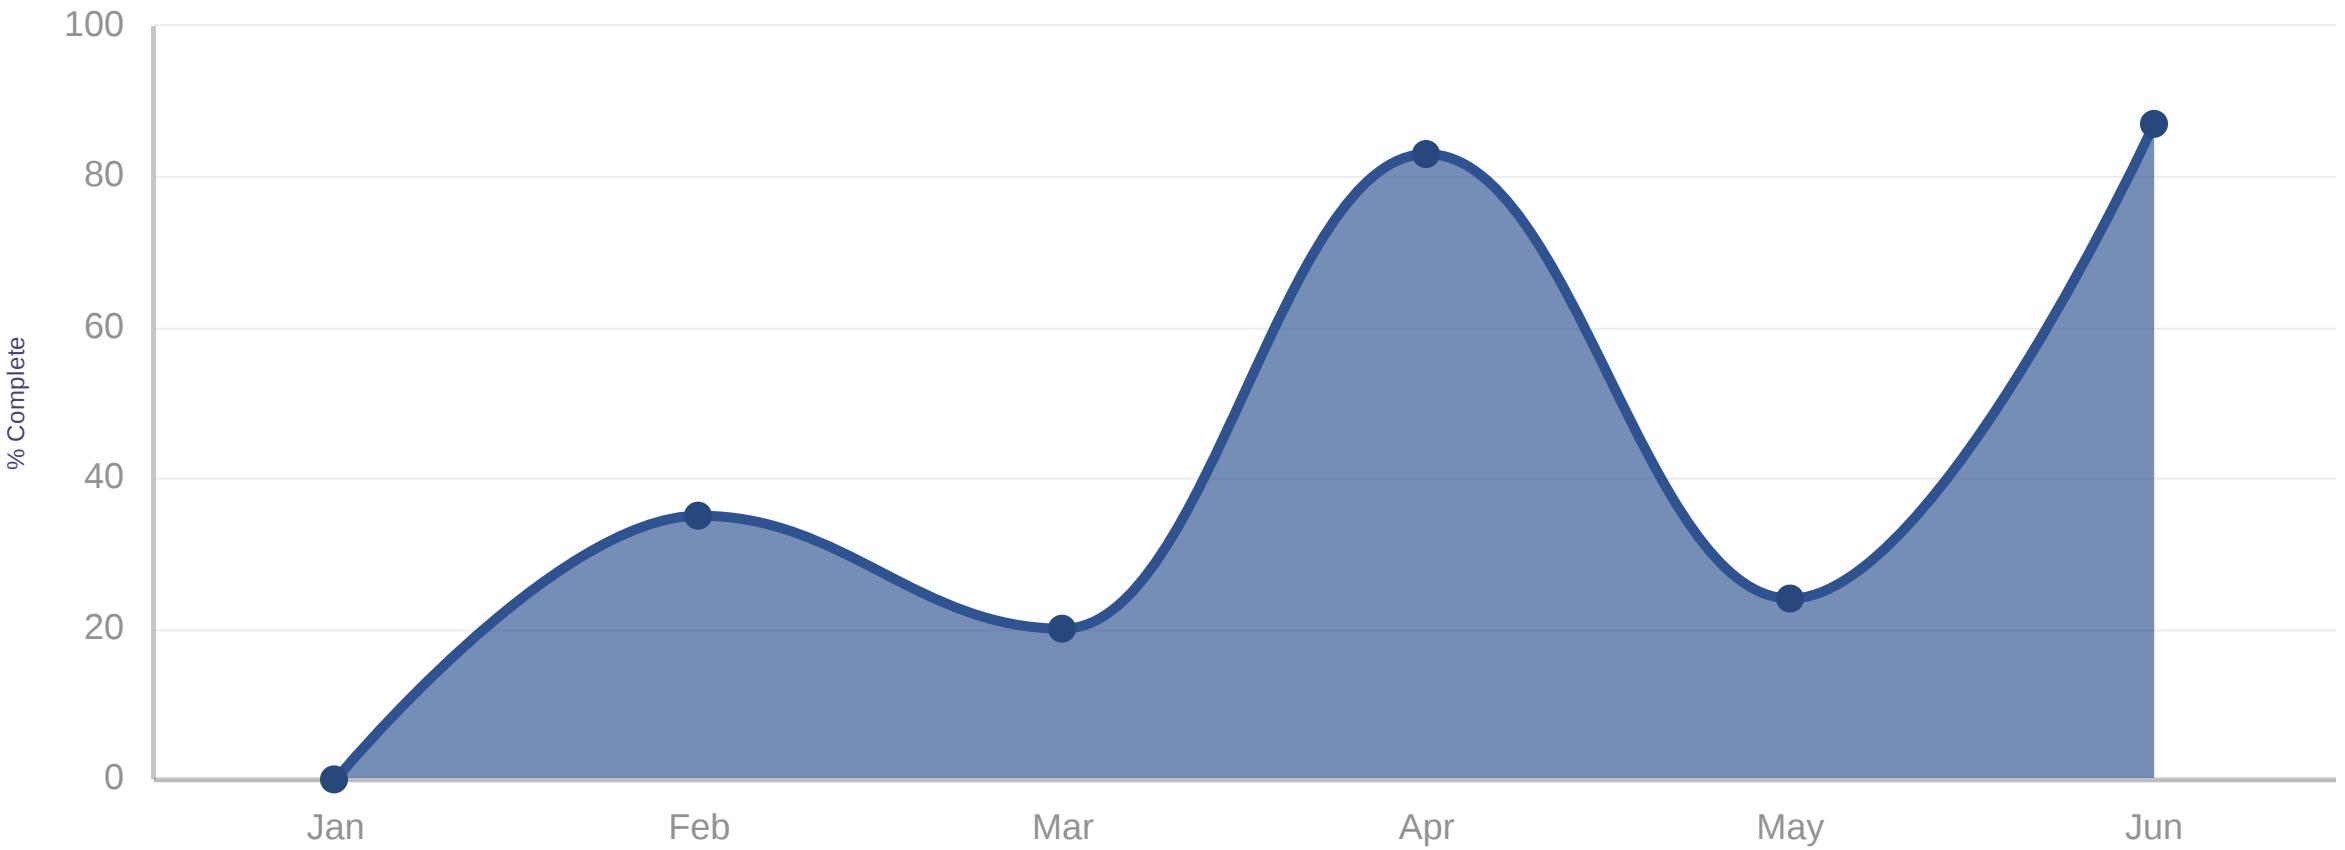
Why bother testing code?



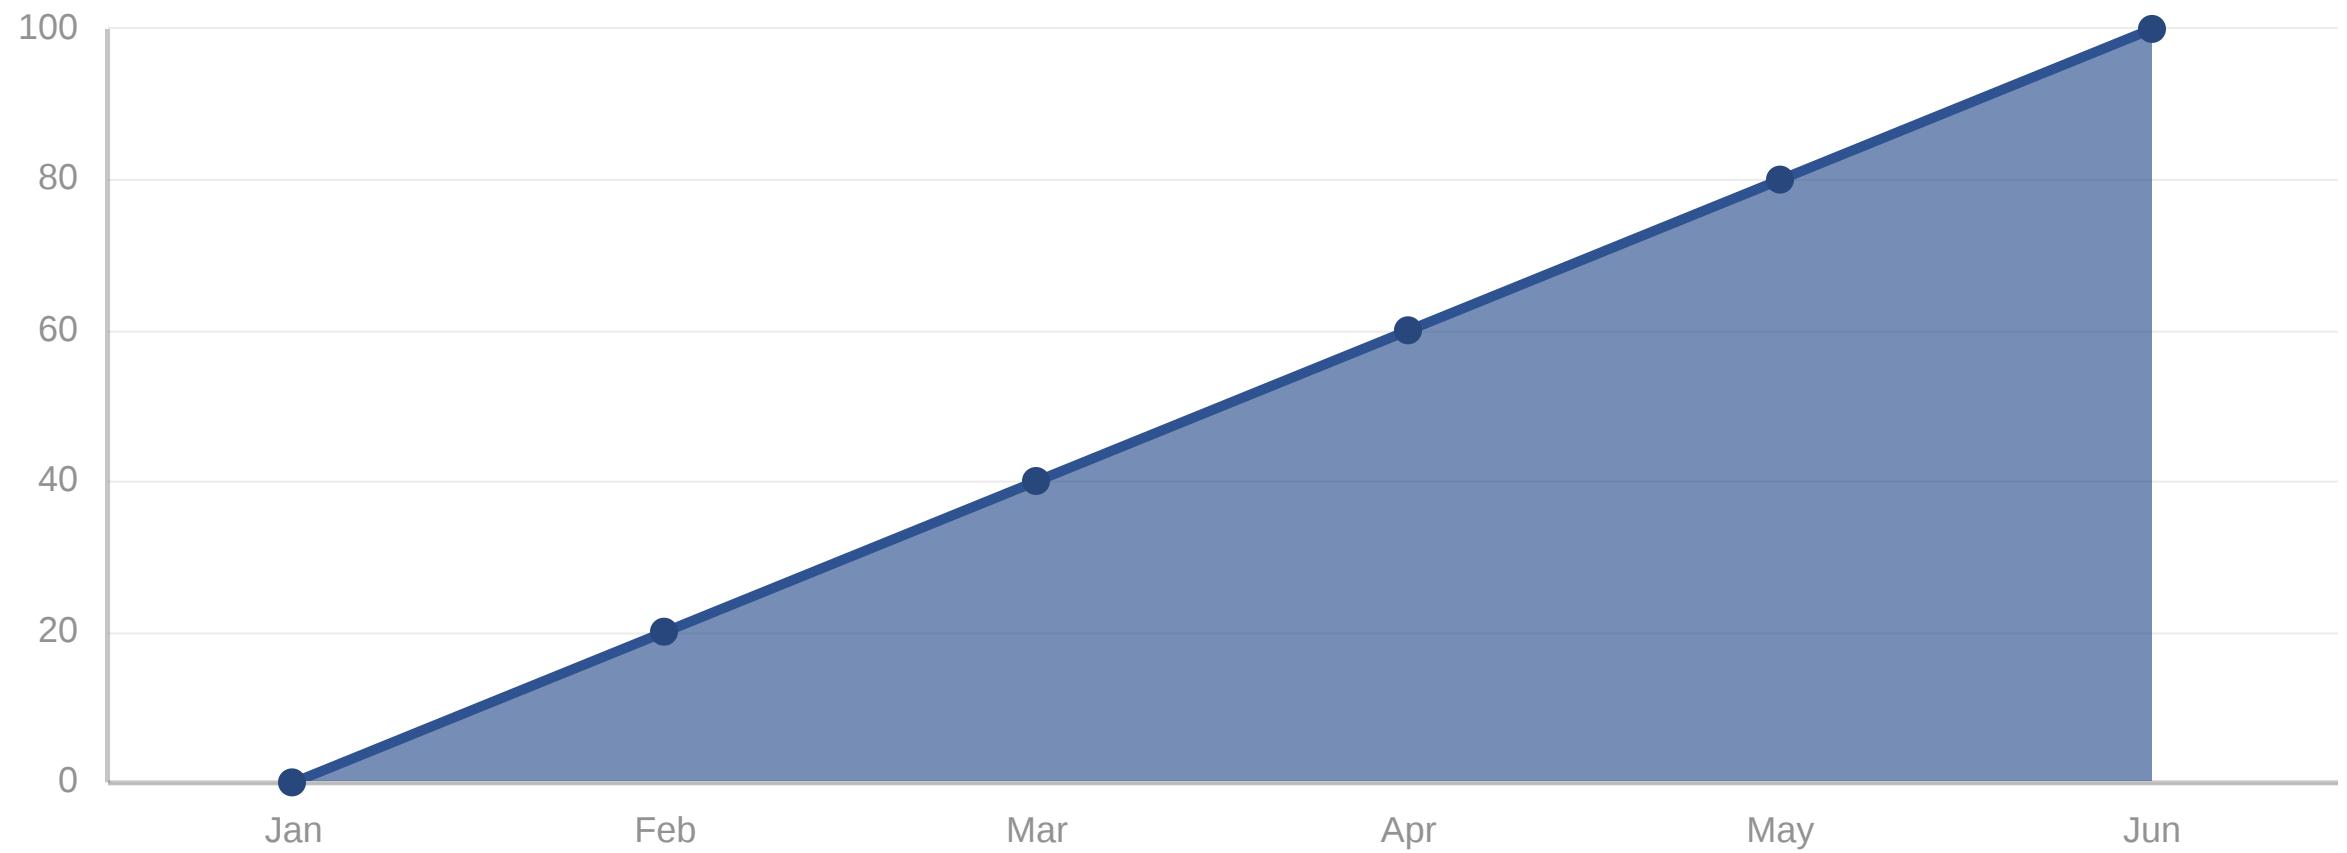
Test driven development?



Typical development



Test Driven



Benefits

- **Define the purpose of your code**

Defines inputs, outputs and functions before you write

Helps you design how the code fits together

Helps you keep code clean and simple

- **Noisy failure**

Fewer surprise failures down the line

More reliable code

Confidence in results

- **Granular progress**

You can track code progress more effectively

Easier to predict completion time

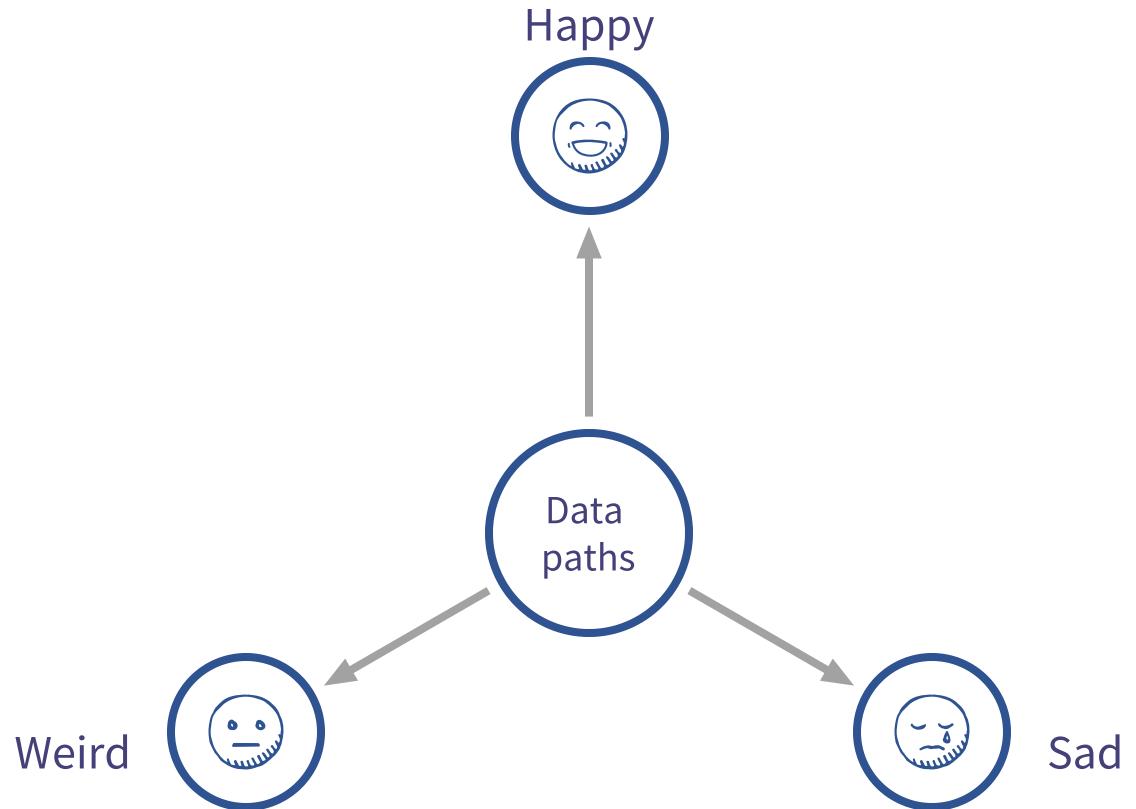
- **Better code**

Timing tests point out processing bottlenecks

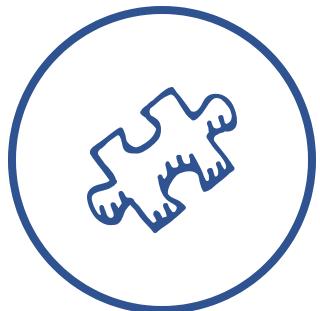
Stops feature bloat/creep

Safer refactoring

Paths



Further reading



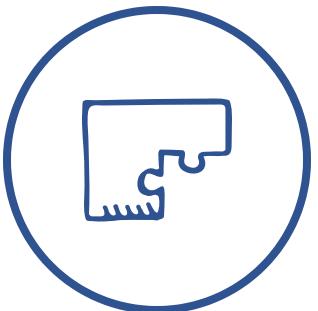
Unit Testing

Function-level testing of the smallest chunks of code

Automated code health checks

Code is tested outside of its environment

Looks for unneeded dependencies



Behavioural Testing

Tests pieces of code on their "behaviour"

Tests by what the program should produce, not just on individual "units" of code

Based on a framework with natural language tests



Test Driven Development

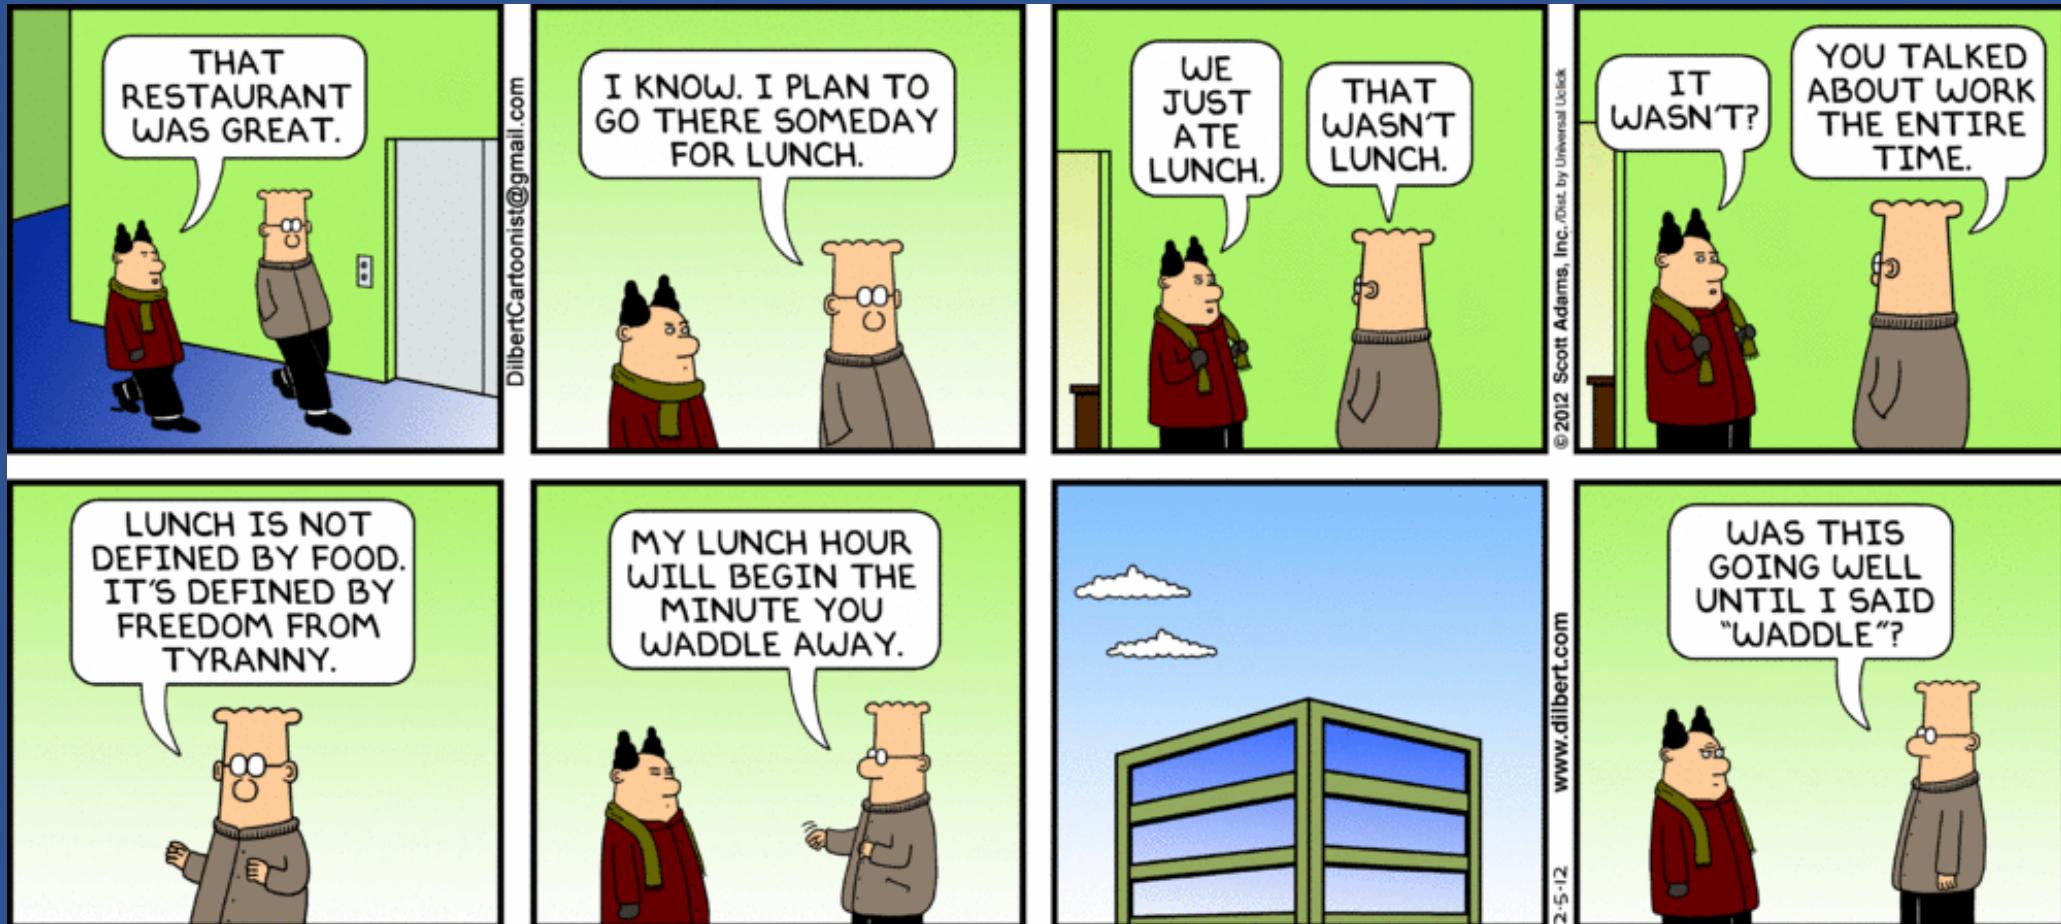
Framework for building code with tests

Everything the code must do is represented by a series of tests

All code should fit the tests

The program is complete when all tests pass

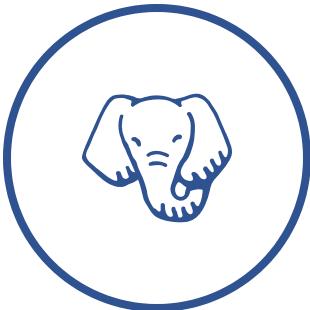
Lunch



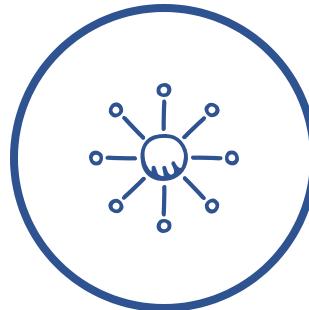
This Afternoon - “Buzzword” Techs



Cloud



Big Data



Internet of Things

Section 3

Cloud

What is the cloud?

Cloud resources:

- Online servers
- You rent machines
- You subscribe to services
- You store information remotely
- You compute remotely
- Prices and resources change dynamically

Platform Providers



Amazon Web Services



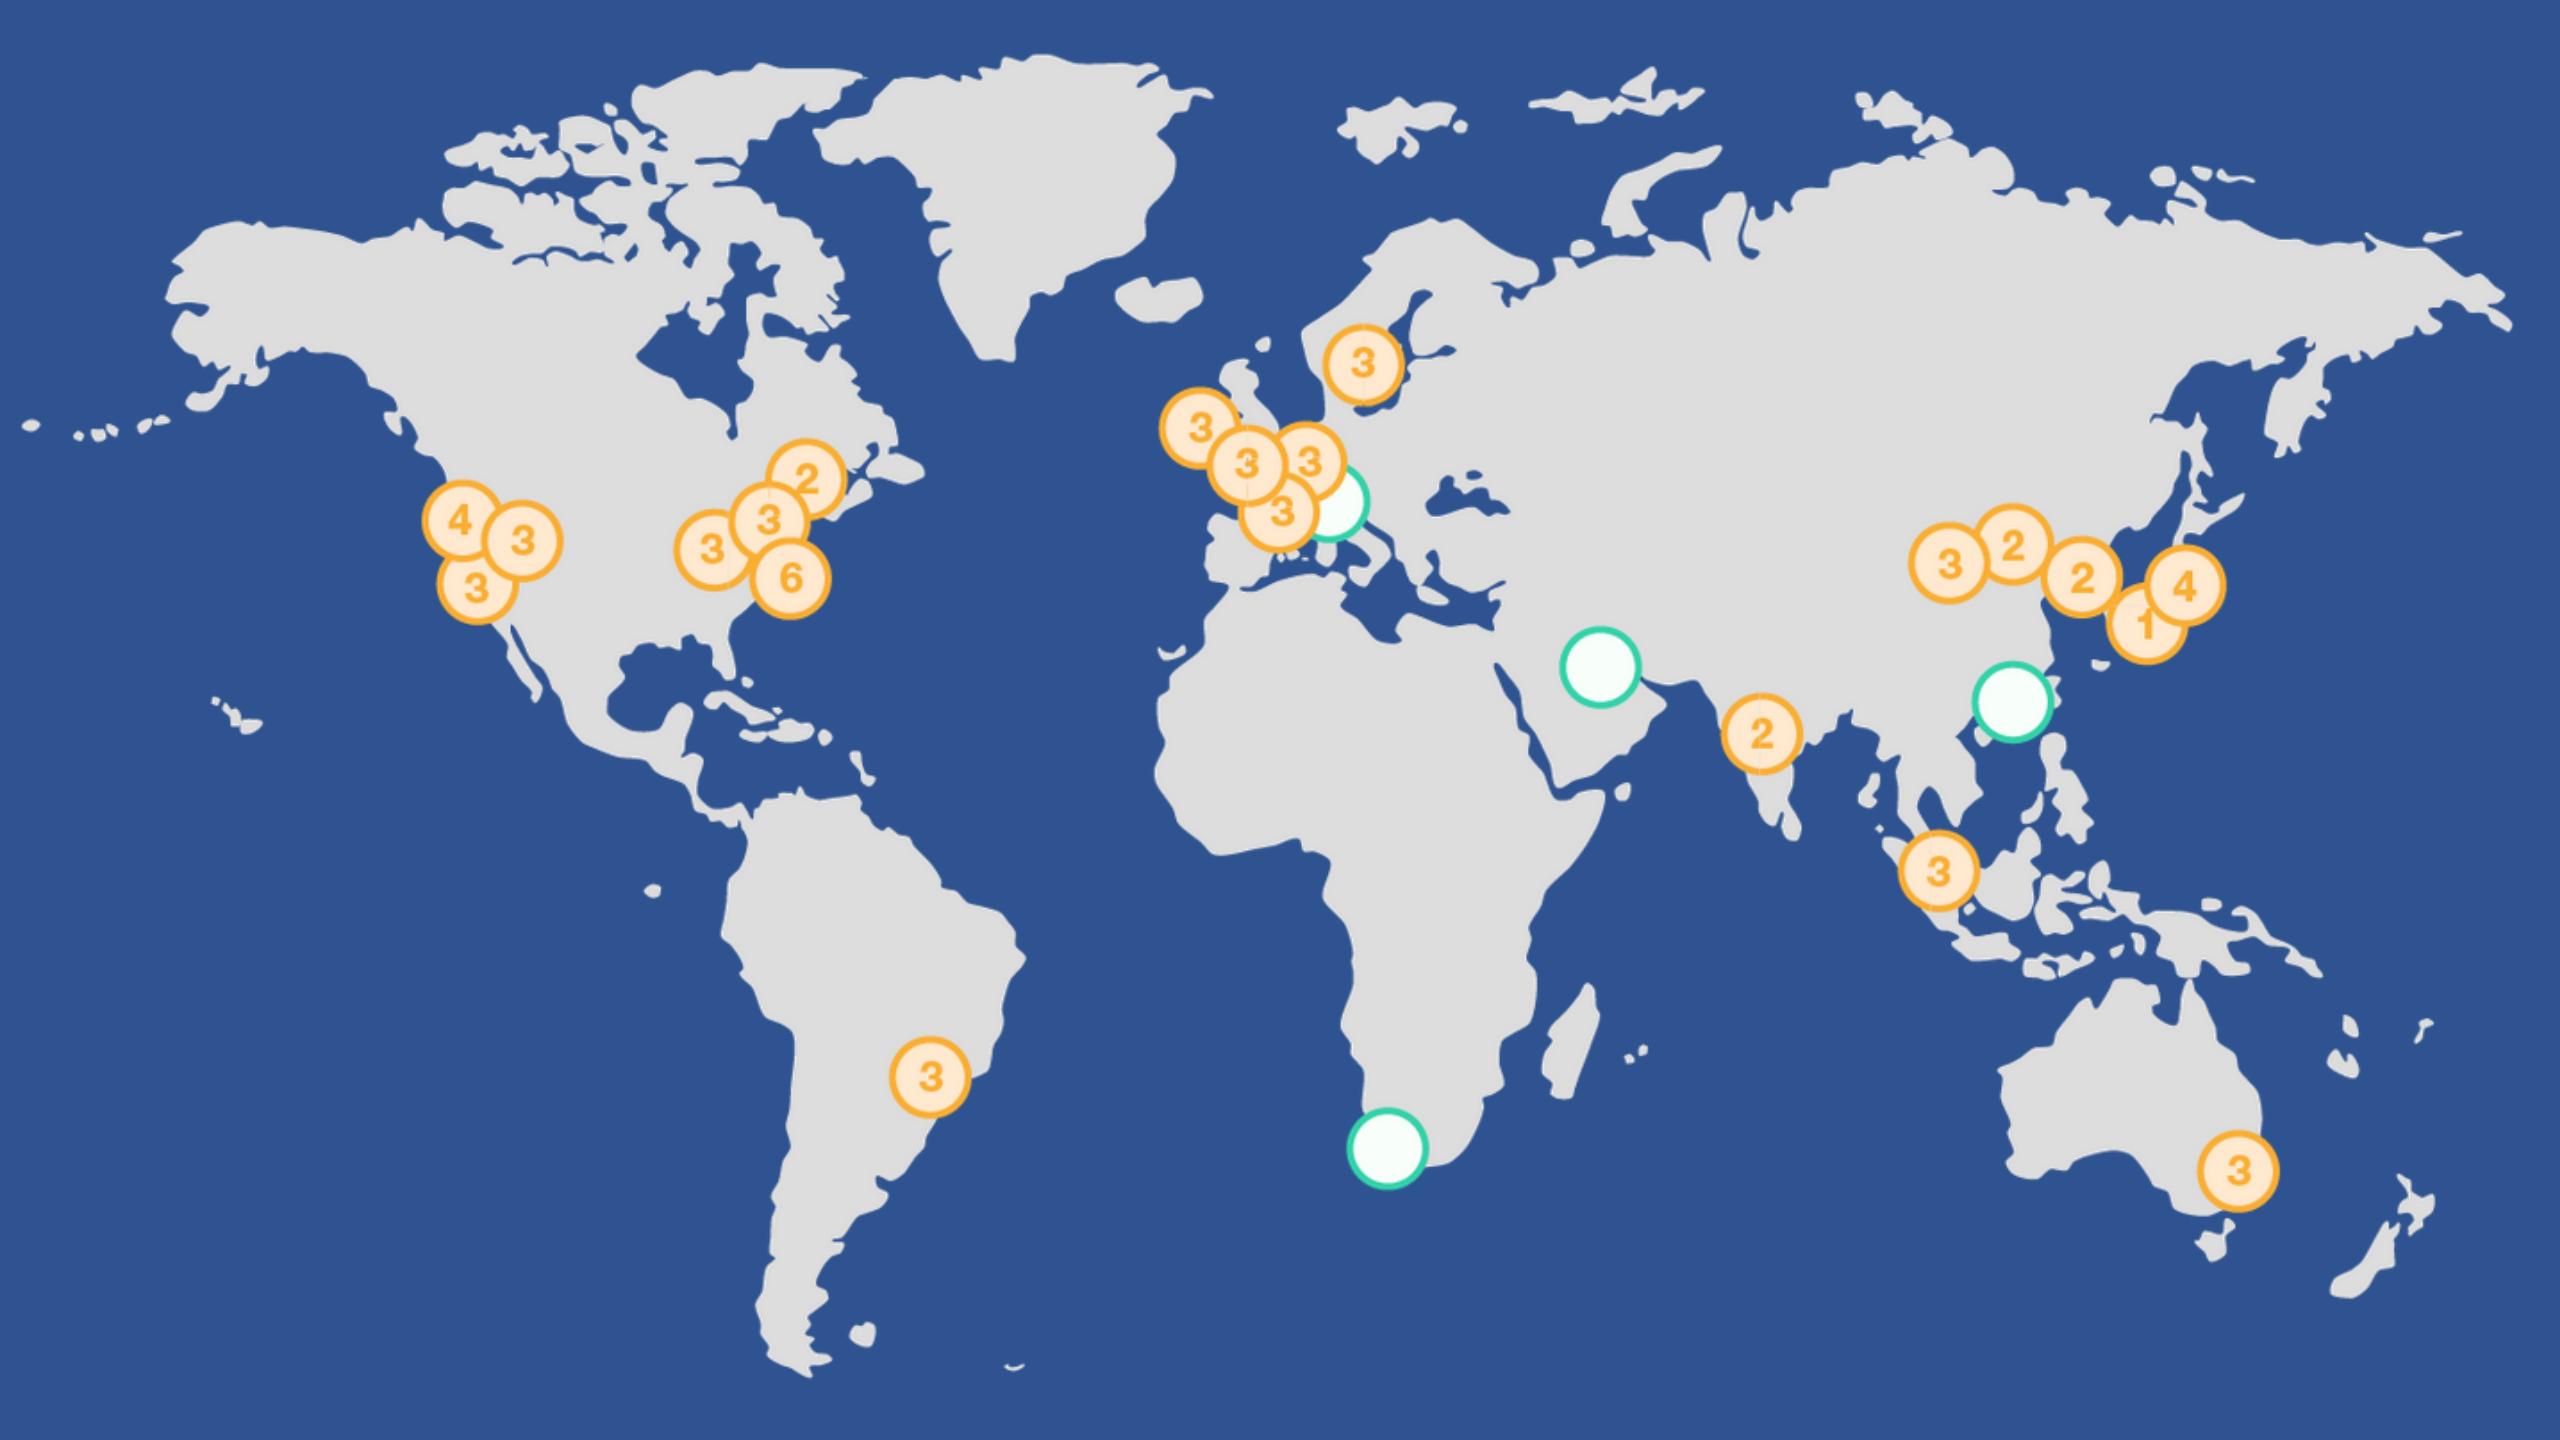
Google Cloud Platform

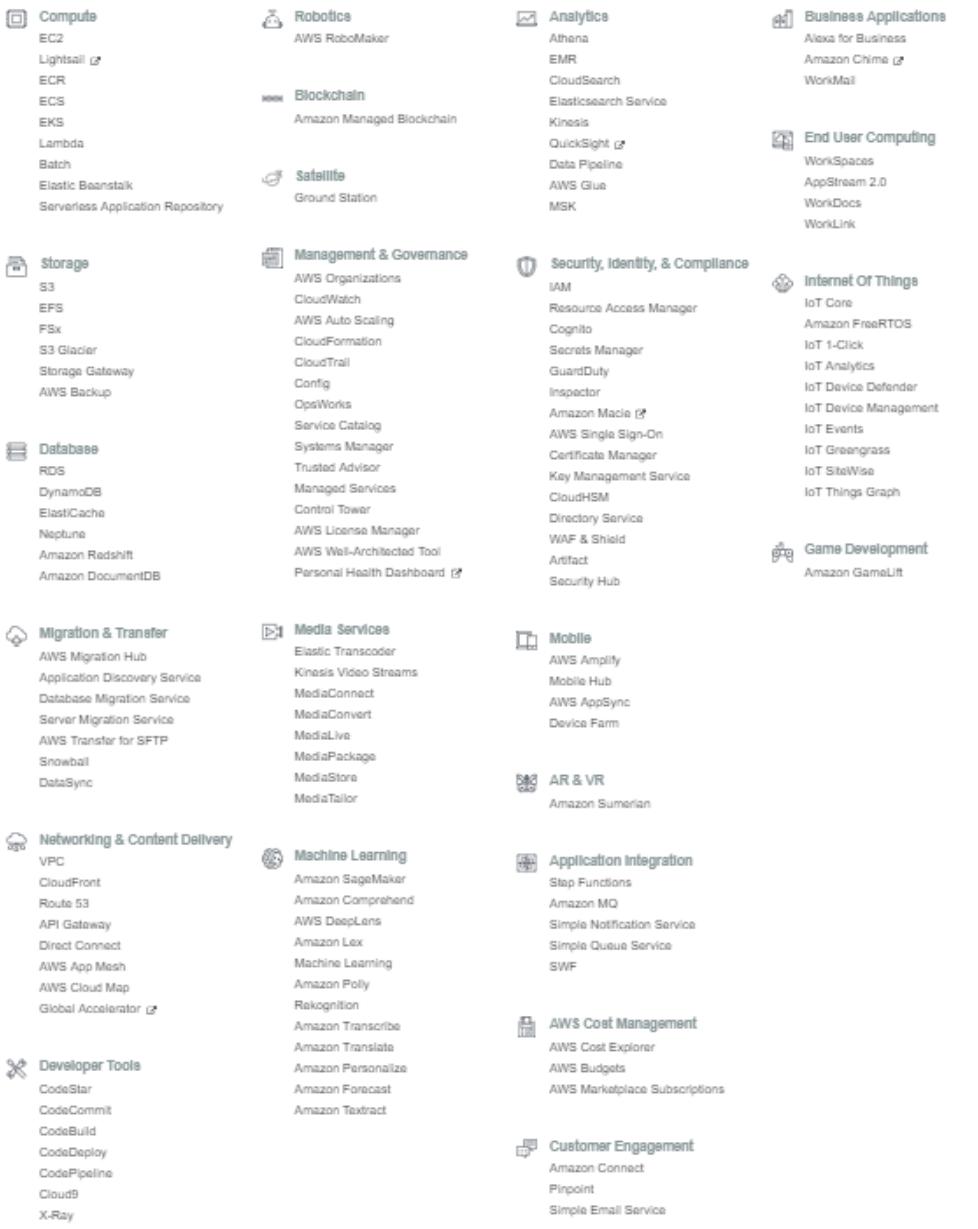


Azure

Market Share







Core services



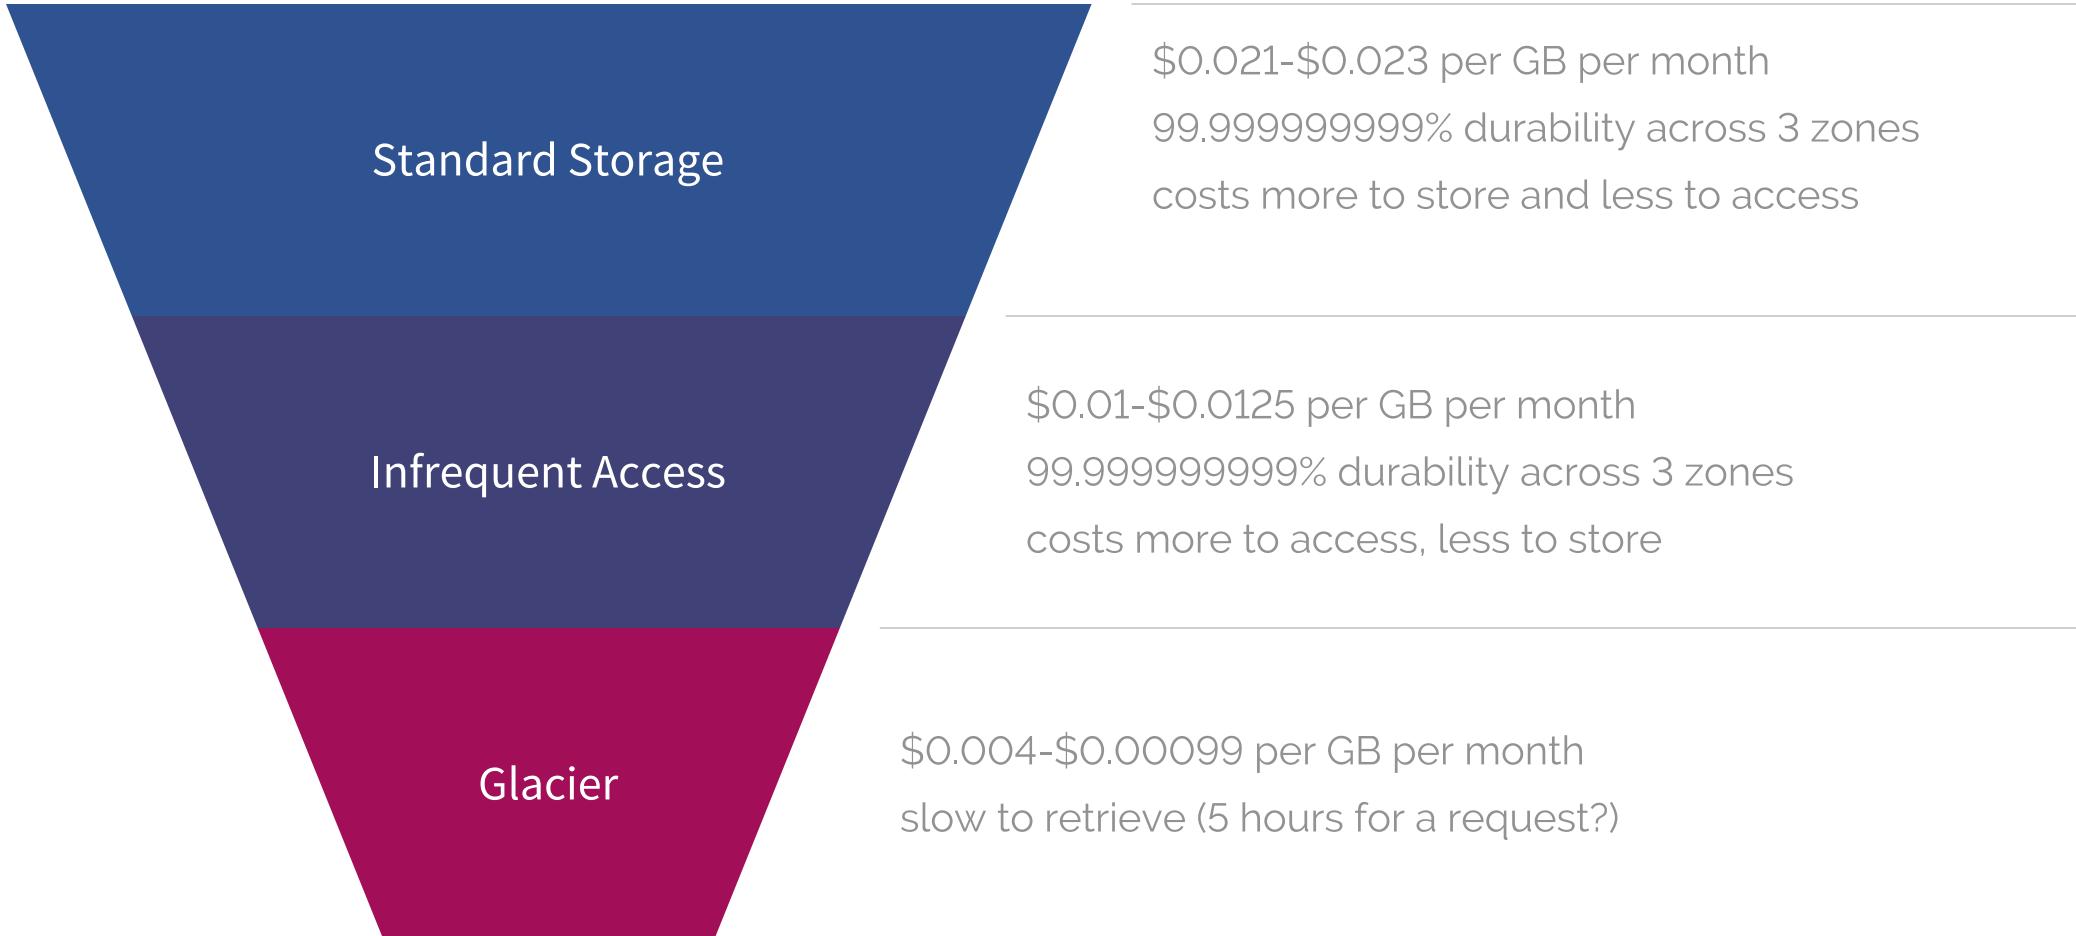
EC2

Elastic Compute Cloud



S3

Simple Storage Solution



EC2

- virtual server
- predefined builds
- various hardware options
- reserved/hourly/spot pricing
- replicable

EC2 instance types

Instance	Use	Notes	Cost (on demand hour)
T3	General Purpose	General Balanced Server	\$0.0052 - \$0.3341
M5	General Purpose	T, with higher network bandwidth	\$0.096-\$4.608
A1	General Purpose	ARM edge processing	\$0.0255 - \$0.408
C5	Compute Optimised	HPC - good for number crunching	\$0.085-\$3.06
X1	Memory Optimised	Extreme in-memory needs	\$6.669-\$26.688
R5	Memory Optimised	For in-memory calculation (e.g. in-memory DB or stream)	\$0.126-\$6.048
P3	Accelerated Computing	General Purpose GPU instances	\$3.06-\$31.212
G3	Accelerated Computing	Graphics intensive instances	\$1.14-\$4.56
F1	Accelerated Computing	allows low-level GPU logic alteration	\$1.65-\$13.20
I3	Storage Optimised	Resource intensive databases	\$0.156-\$4.992
D2	Dense-store Optimised	Massively-Parallel Processing	\$0.69-\$5.52

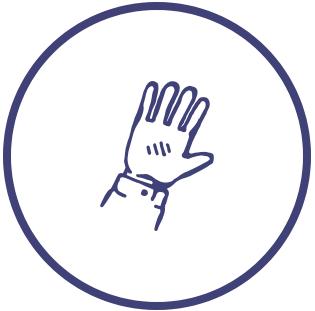
prices



On demand

Most expensive

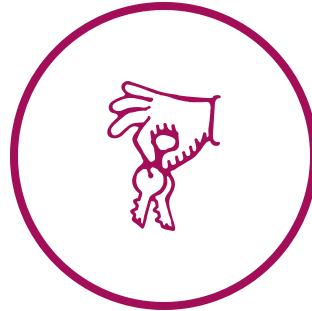
Run your server whenever you want
Only pay for the hours that your
server is online



Spot-Price

Name a price per hour

When server demand drops, "spot"
price falls. If you meet the spot price
you get compute time.



Reserved Price

You rent a server whole-hog for a
year or three years

Cheaper than on-demand
reliable

Some other services...

- **RDS**

Relational Data Store - automanaged Databases

- **Lambda**

Serverless computation

- **QuickSight**

Dashboard Analytics

- **Kinesis**

Data streaming

- **Data Pipeline**

Pipeline management and scheduling

- **Sumerian**

VR and AR development

- **Cloud9**

online IDE

- **GreenGrass**

IoT device management

- **Polly**

Text to Speech

- **SageMaker**

AI development

- **Rekognition**

Pre-built computer vision

- **GroundStation**

Satellite Management

Section 4

Big Data

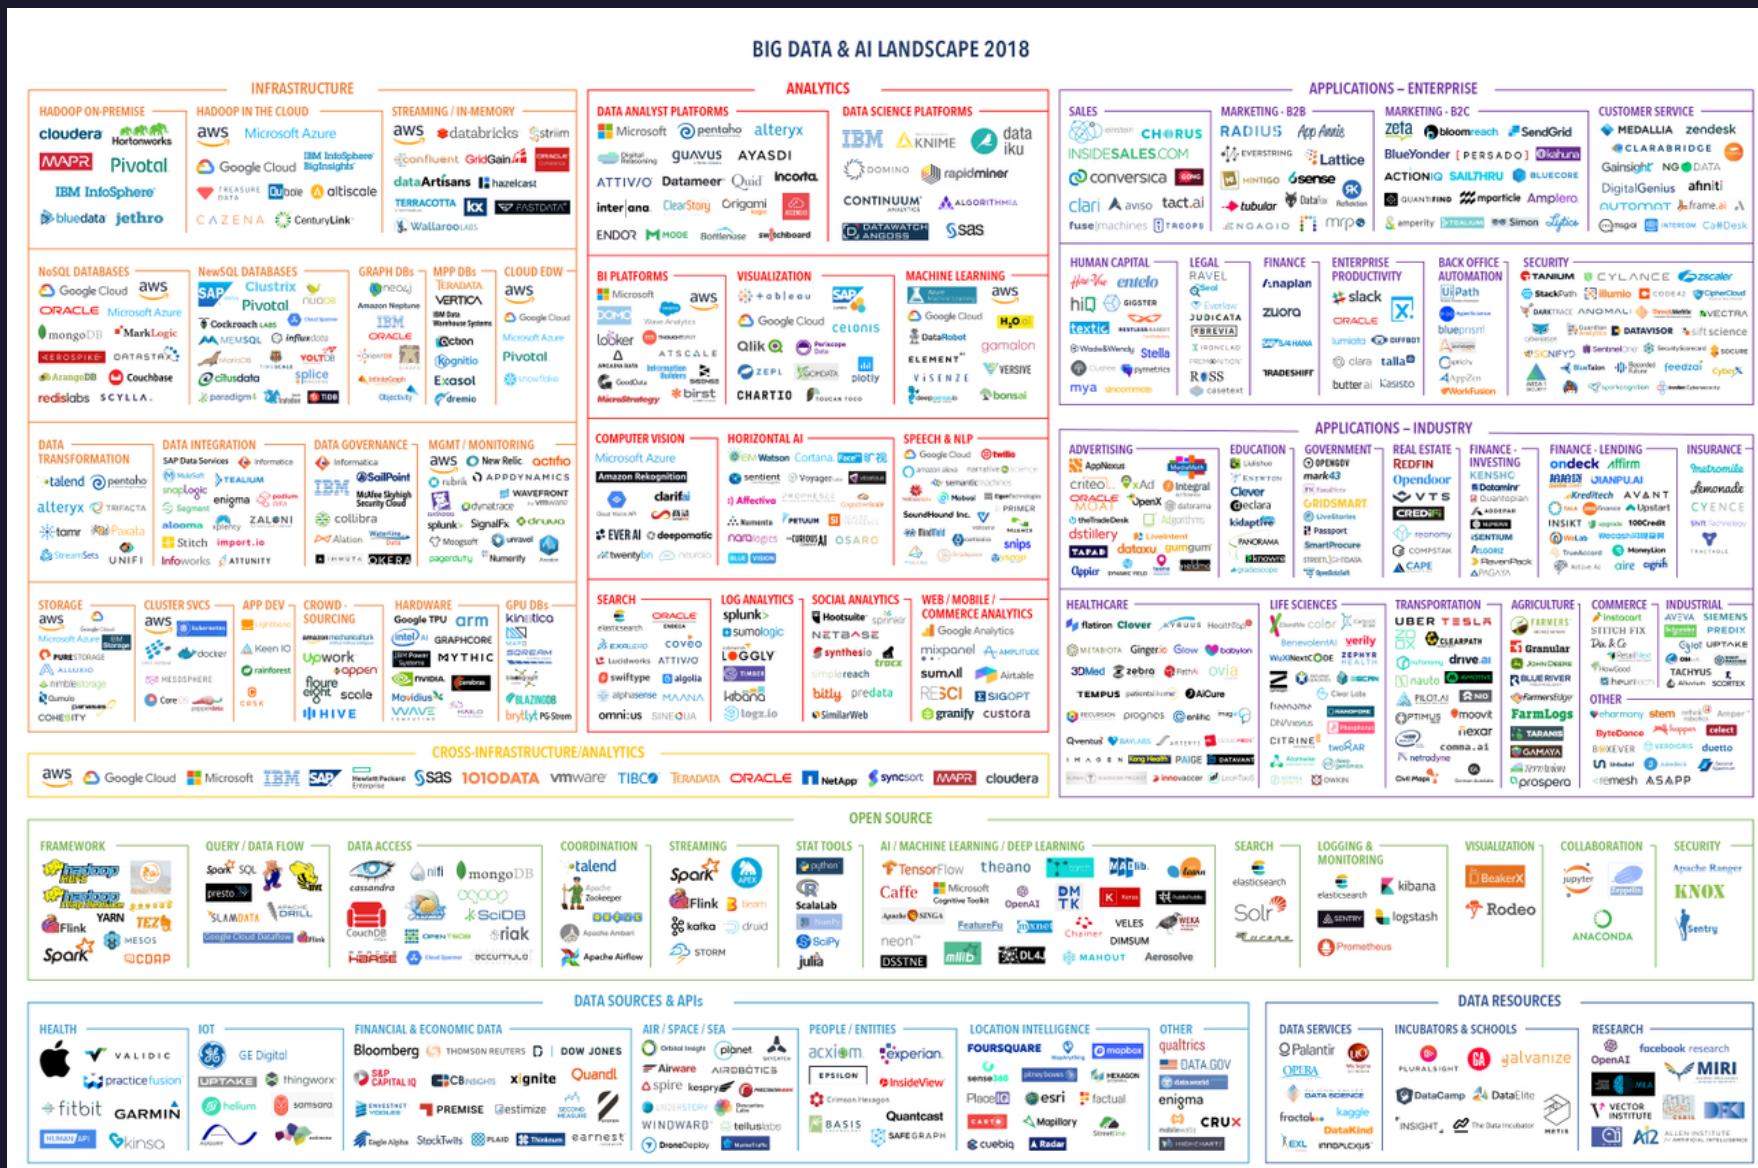


Major Technologies



TERADATA





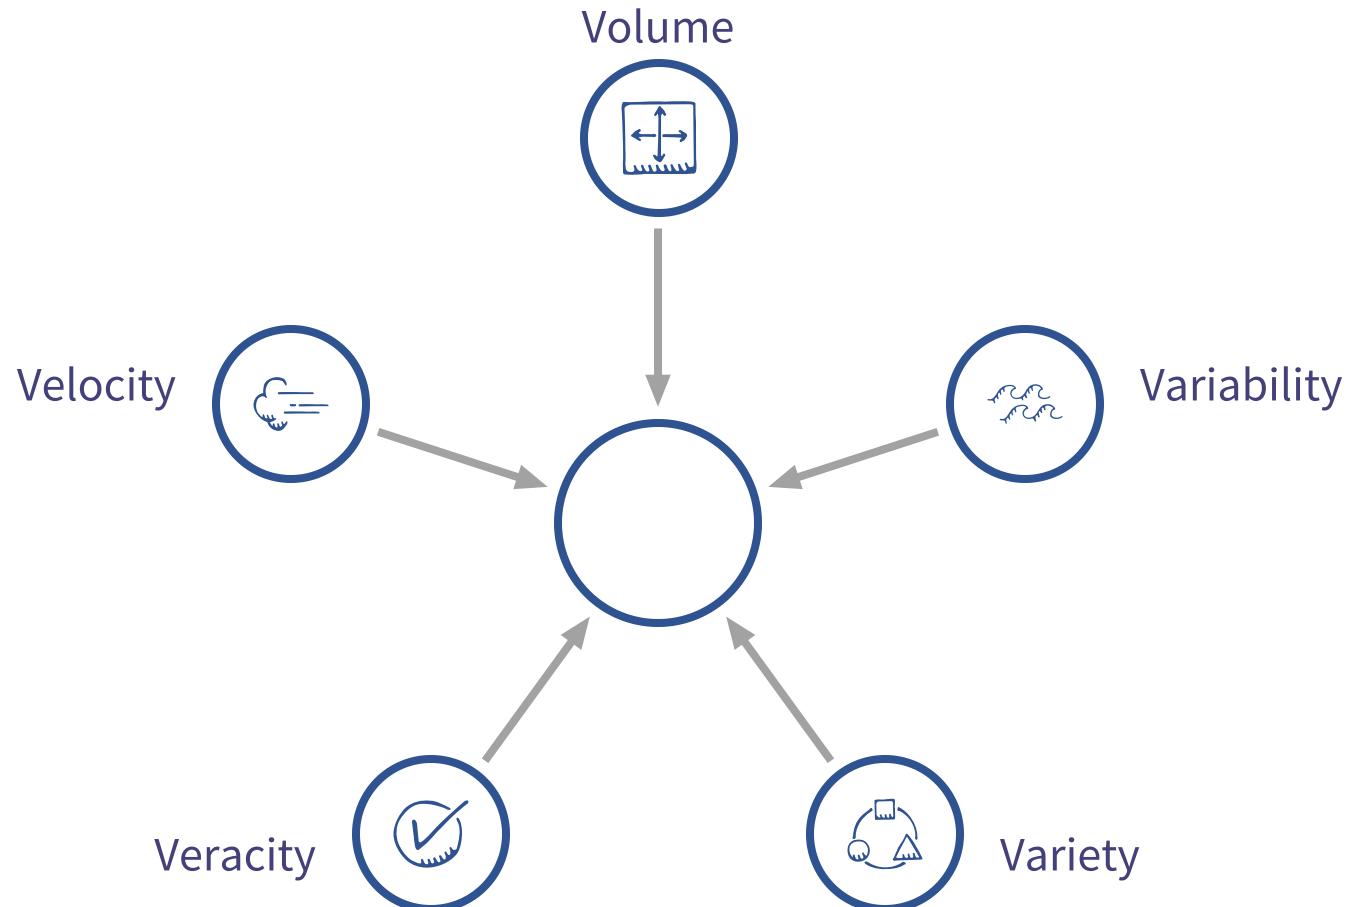
Final 2018 version, updated 07/15/2018

© Matt Turck (@mattturck), Demi Obavomi (@demi_ obavomi), & FirstMark (@firstmarkcap)

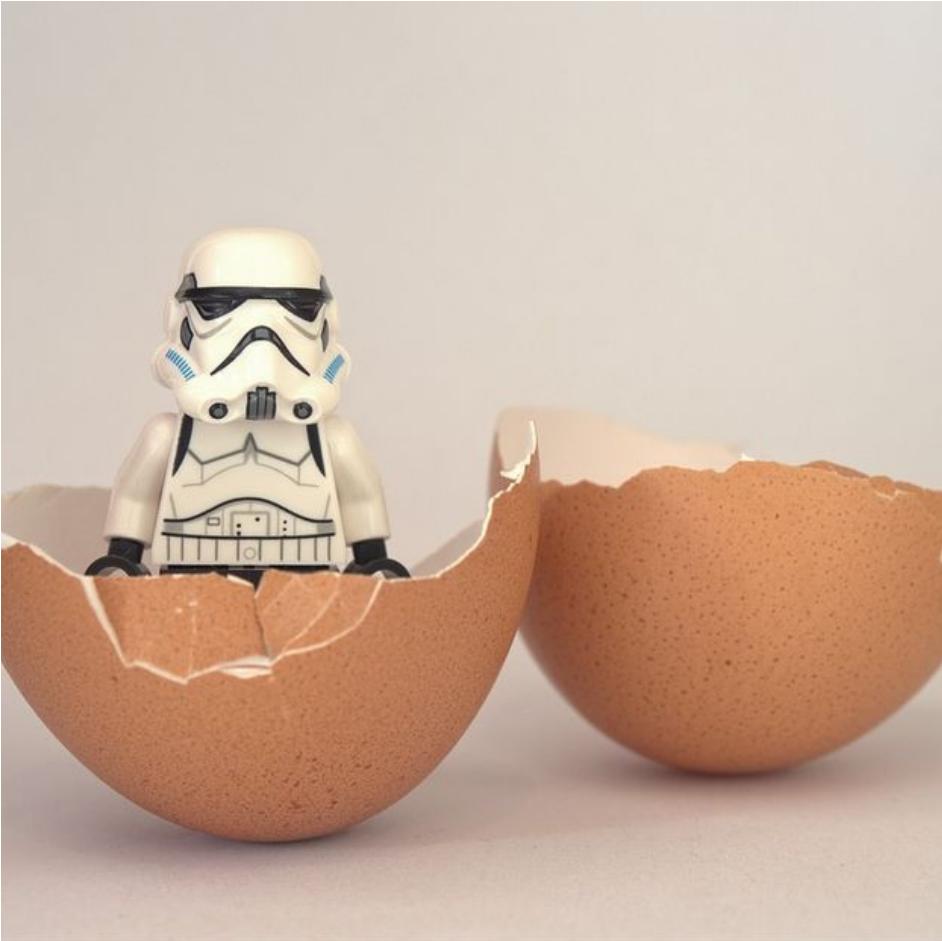
mattturck.com/bigdata2018

FIRSTMARK
EARLY STAGE VENTURE CAPITAL

The V's



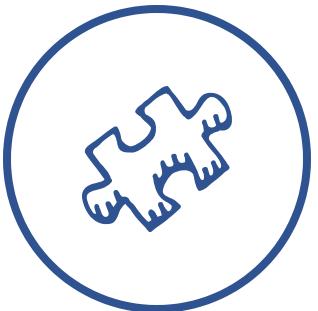
What's different?



Traditional approaches break down. No matter how big and powerful your servers are they can't keep up.

Instead of processing your data through big servers as fast as possible, split the data across hundreds of servers and process different bits in different places

Big Data Steps

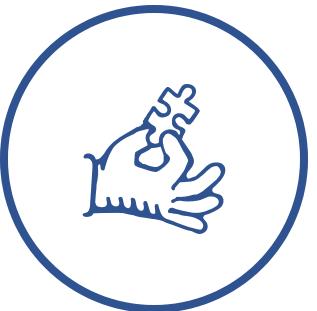


Distributed file system

Data is kept fragmented across a cluster of systems

Each node in a cluster has an incomplete copy of the data

The same piece of information will be stored in multiple servers

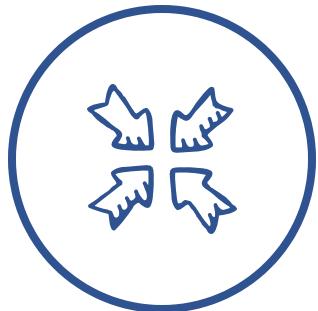


Computational mapping

Any queries need to be sent out to every node in a cluster

Data should be distributed to ensure even compute load

Parallel computation speeds up processing



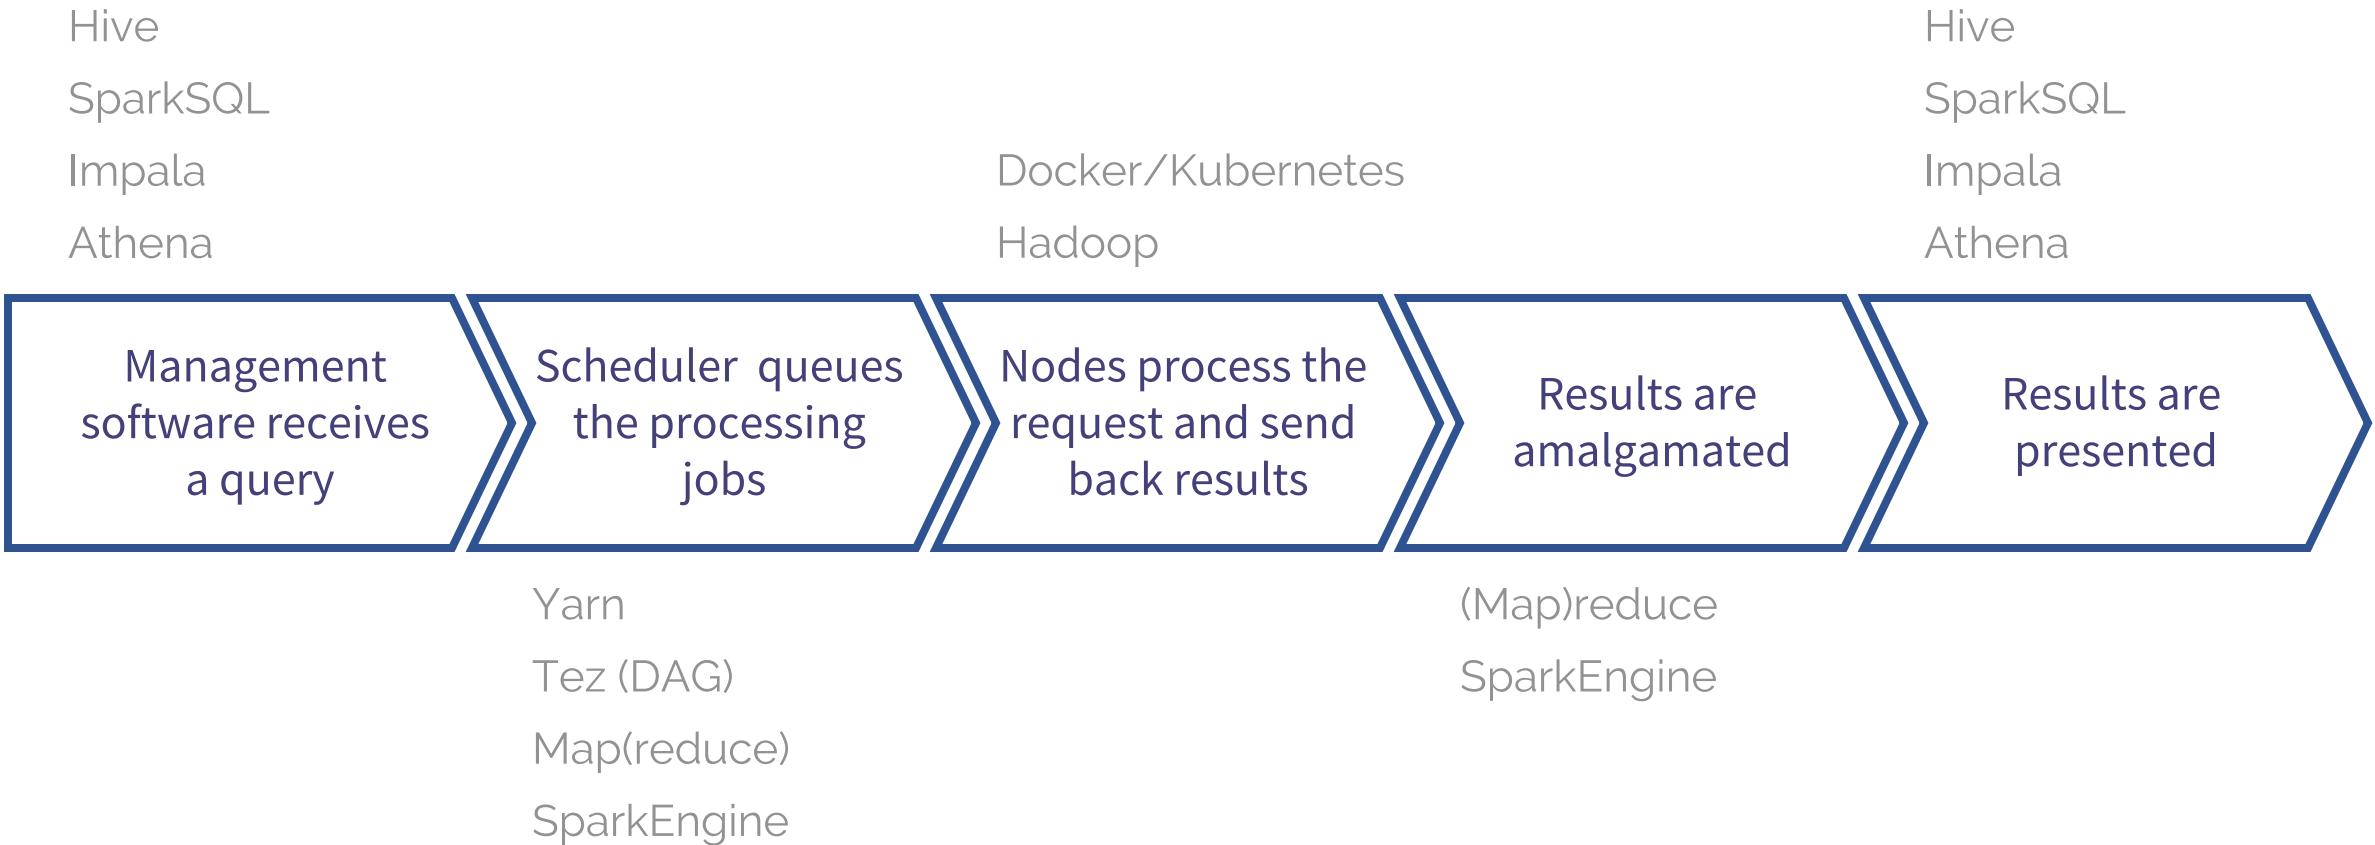
Results reduced down

Each node returns different results from their dataset

They need to be reduced to a single result

Beware! there be dragons here

Stack Flow



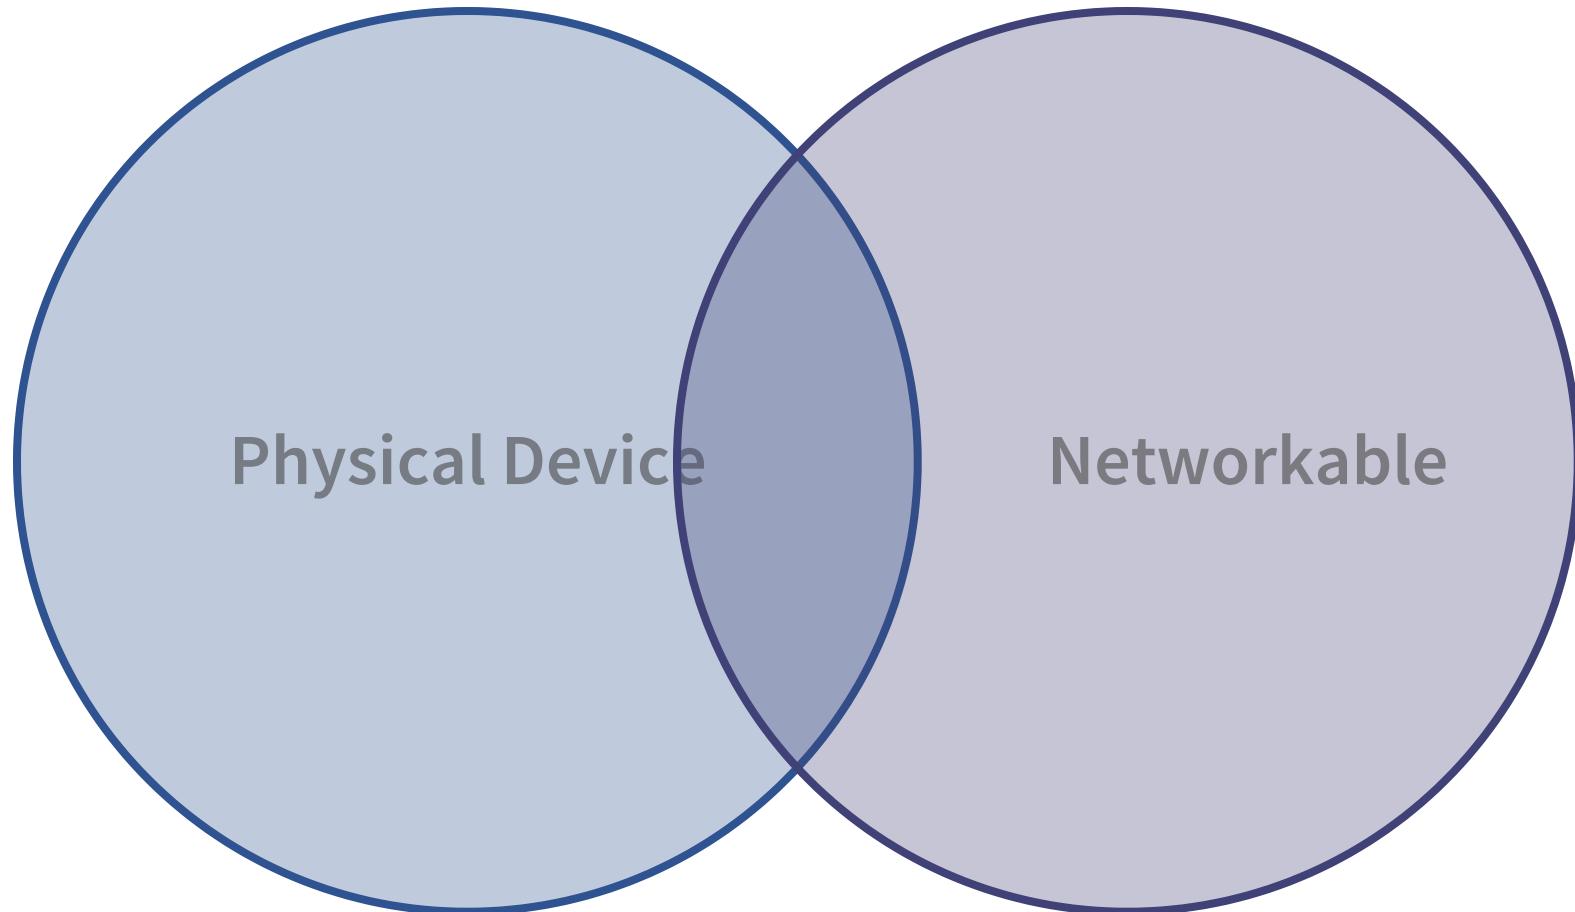
Section 5

Internet of Things

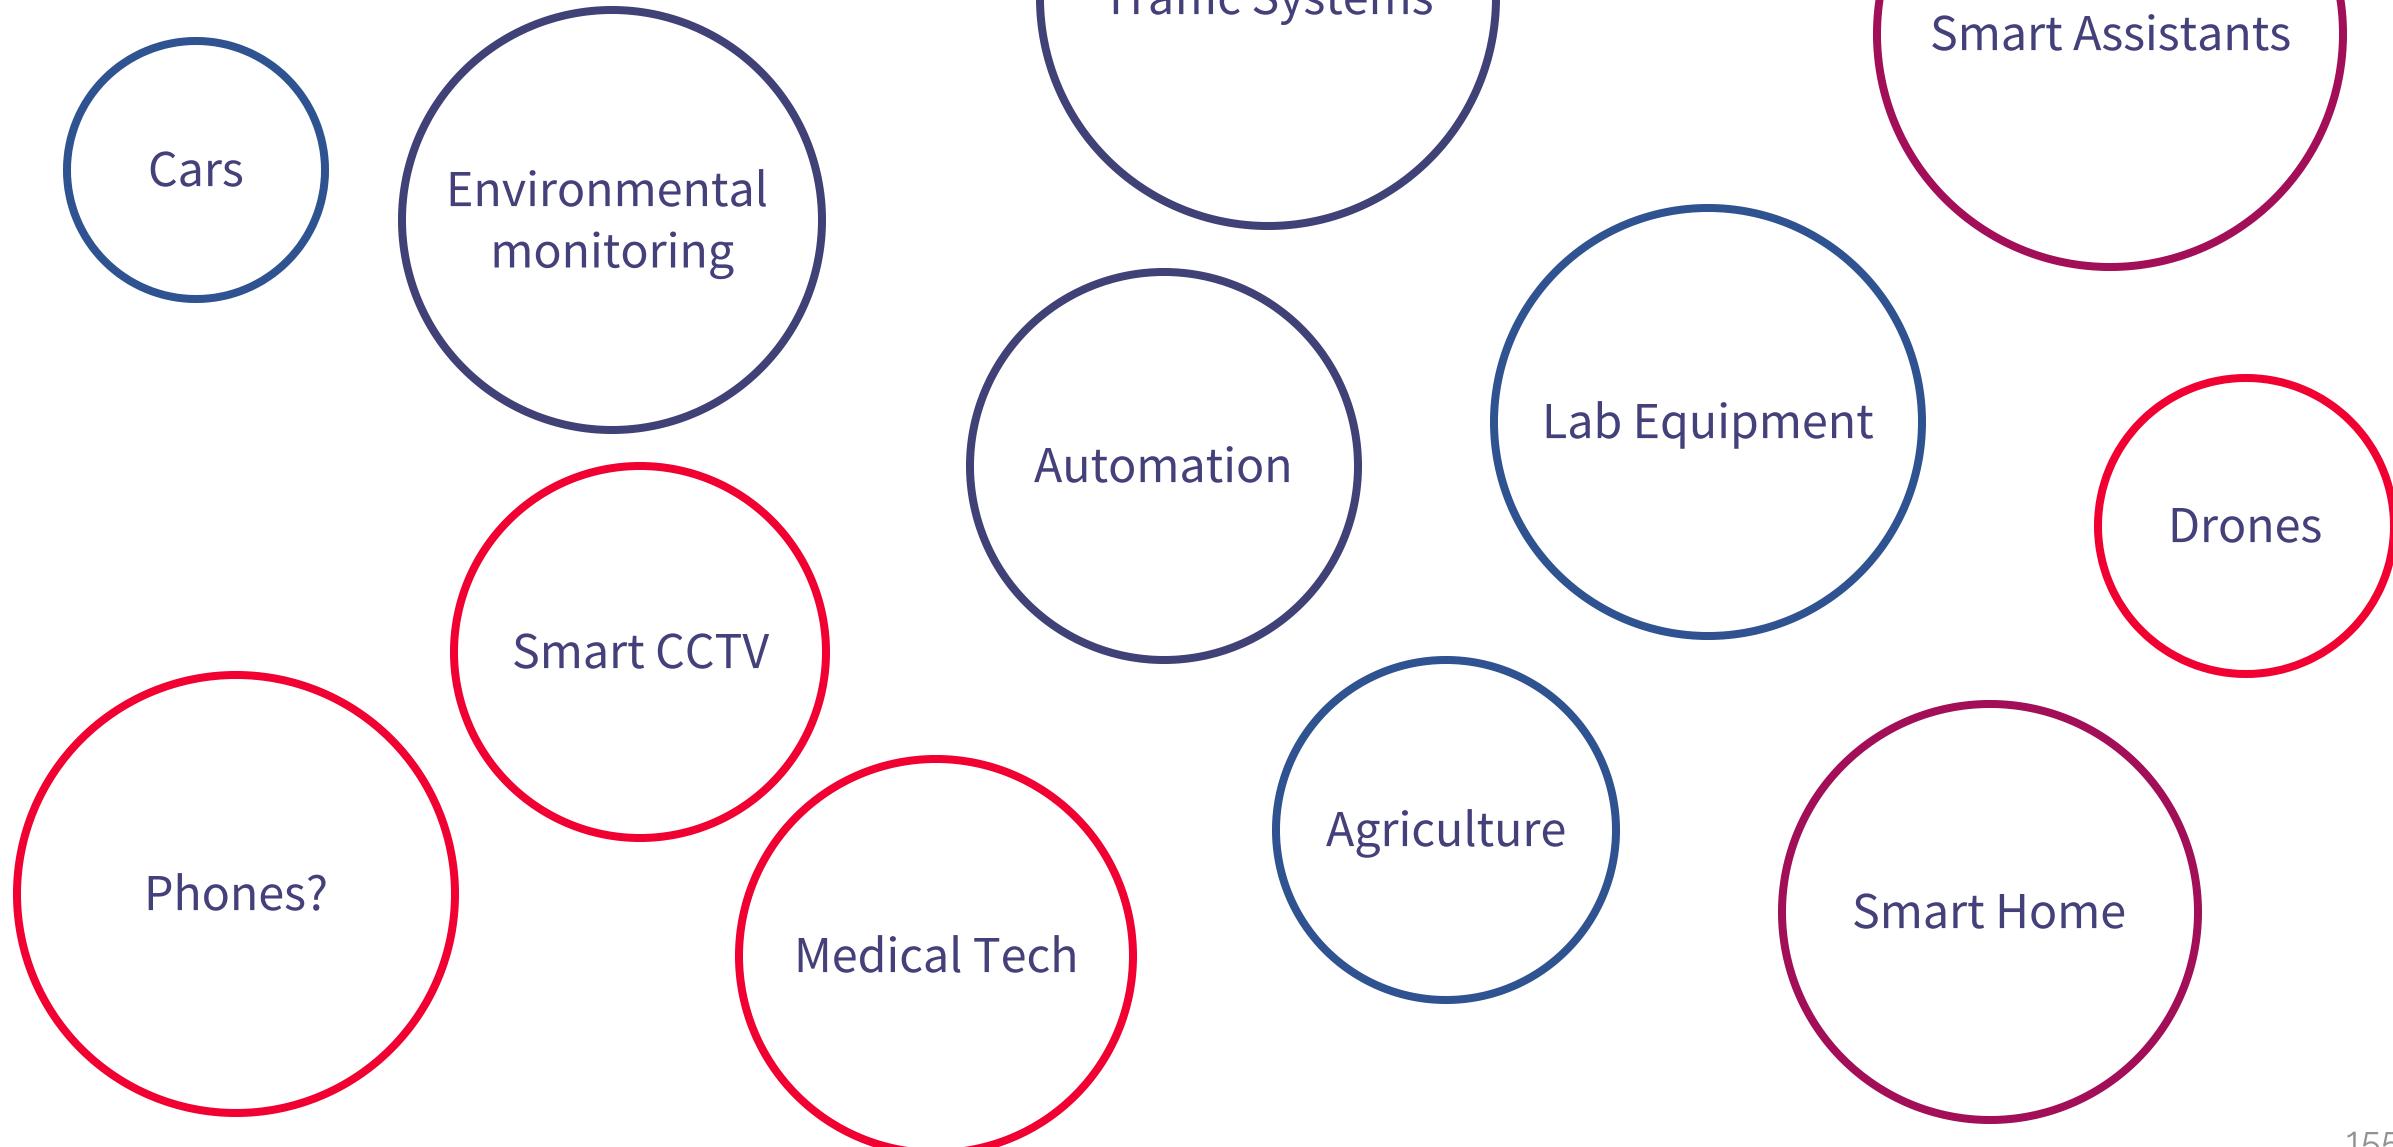


What is IoT?

IoT Devices



Device fields



Transmission



Collected



Batch



Stream

Received Data



Location Data

RTK

GPS/GLONASS/Galileo



Data Transmission

WiFi

GSM



User Input

Voice commands

Interaction

Break



Section 6

APIs

Application Program Interface

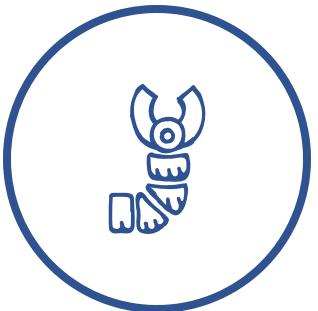
- Platform for providing and receiving information
- Other applications don't have to interact with a service on a code level
- Only exposes part of a system outwards
- can be used to transfer data or action commands

API examples



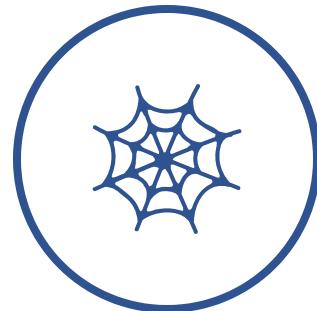
Phone Camera

Rather than make their app work with each phone's camera code, apps instead connect to an API that provides the image data.



Arduino

Much of the servo/stepped motor control runs through an API that accepts simple code requests that are converted into more assembly/complex machine code



Web service

To access data held on an online server, programs can call to a web API that provides the data they need in a structured way, and manages the service.

Web (Services) API

- RESTful (Representational State Transfer) APIs

The most common type of API for larger data providers

Requests can all be driven through HTTP web links

Often these providers have their own libraries for requests

Resources sit ready for being visited by the right URI (Unique Resource Identifier)

Often dumps data in JSON format

e.g. Met Office data

- SOAP (Simple Output Application Protocol)

Less common for data providers

More complex and less uniform than REST

The protocol allows operations to be driven, rather than just simple data recall

Older style of API

Can be ACID compliant - better for transactional data, and can communicate during data transfer

Supports better security measures

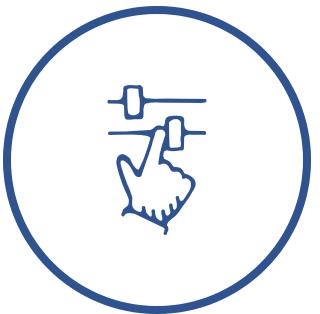
e.g. Mobile banking application talking to a banking server

Twitter API



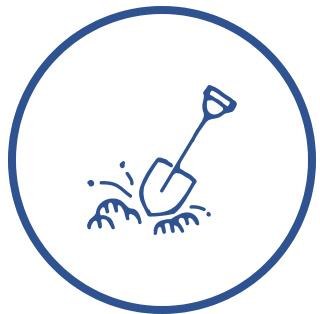
Query data

You can search by user, hashtag,
geolocation etc...



Control user accounts

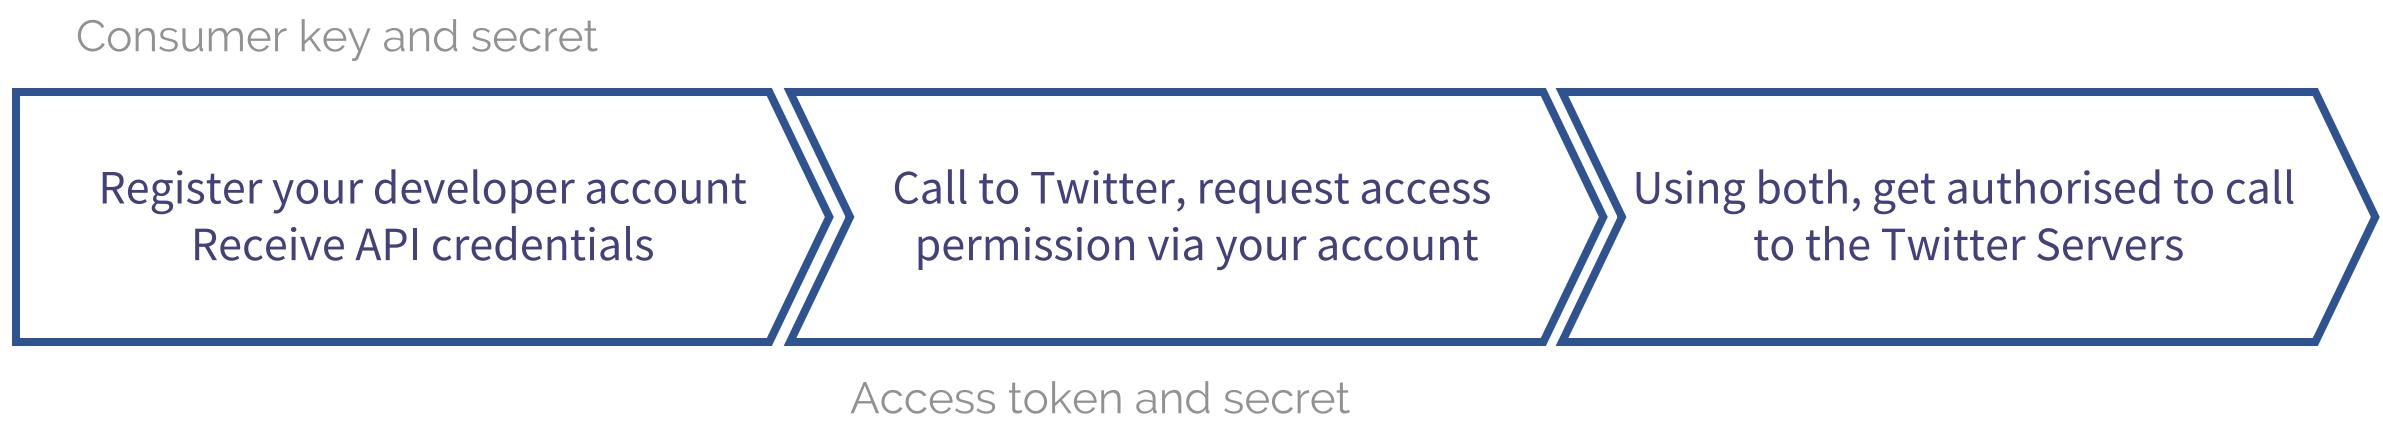
Interact with DMs, analyse follower
networks, post for a user...



Mine data

Pull data by the terabytes...

Twitter API - OAuth2



Your turn!

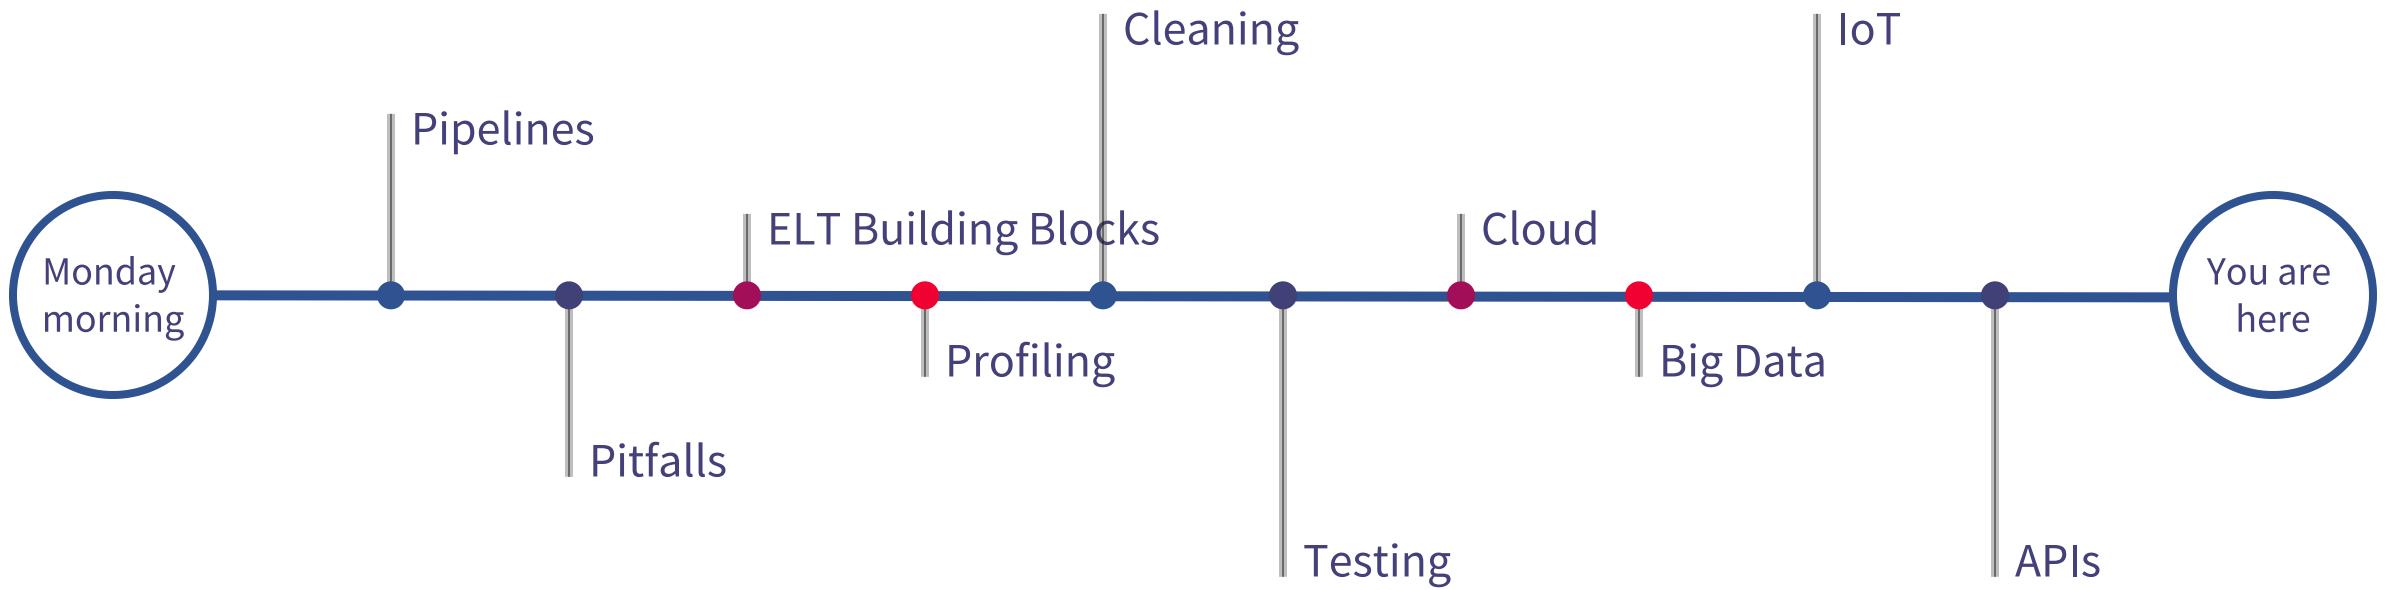
Twitter API lab

- Register with Twitter
- Establish token access
- Run queries
- Pull results
- Analyse!
- [https://github.com/R-Strange/
Data_Wrangling_Course_Oxford](https://github.com/R-Strange/Data_Wrangling_Course_Oxford)

Day 2 Summary

- Pipelines revisited
- Testing your Code
- Cloud Computing
- Cloud Storage
- Big Data
- MapReduce
- Internet of Things
- APIs
- REST
- SOAP

Course



What's next?

eBooks

Manning

O'Reilly

Pragmatic Programmers

Head First

eLearning

Coursera

Udemy

Cloud Guru (AWS Solutions Architect)

Testing

Behavioural Driven Development

Unit Testing

Experiment!

A year of free AWS micro instances

write a few tests in existing code

Build your skills outwards and upwards

You can make an informed
decision with your data

<https://bit.ly/2IgY3gw>



Richard Strange

Earth Sciences 30.17



richard.strange@linacre.ox.ac.uk



075684132854



<https://www.linkedin.com/in/richardpstrange/>



www.github.com/R-Strange

